Evaluating and Refining Arch-Stage

Supporting Architectural Hypotheses in Software Engineering Teams

José Gonçalves Lima Neto

Thesis presented to the
Institute of Mathematics,
Statistics, and Computer Science
of the University of São Paulo
in partial fulfillment
of the requirements
for the degree of
Master of Science

Program: Computer Science
Advisor: Prof. Dr. Paulo Roberto Miranda Meirelles

São Paulo
December, 2026

Evaluating and Refining Arch-Stage

Supporting Architectural Hypotheses in Software Engineering Teams

José Gonçalves Lima Neto

This is the original version of the thesis prepared by candidate José Gonçalves Lima Neto, as submitted to the Examining Committee.

The content of this work is published under the CC BY 4.0 license
(Creative Commons Attribution 4.0 International License)

Acknowledgments

Abstract

José Gonçalves Lima Neto. **Evaluating and Refining Arch-Stage: *Supporting Architectural Hypotheses in Software Engineering Teams***. Thesis (Master's). Institute of Mathematics, Statistics, and Computer Science, University of São Paulo, São Paulo, 2026.

Software engineering teams frequently face uncertainties related to their systems' architecture yet lack adequate methods and tools to manage them explicitly and collaboratively. The ArchHypo technique proposes using hypotheses engineering to make such uncertainties visible and manageable; however, its manual application imposes a significant learning curve. This work empirically evaluates and refines Arch-Stage, a support tool for the ArchHypo technique, by applying it with multiple software engineering teams in a real-world organization with a scaled software product. The research adopts a mixed-methods design combining analysis of artifacts produced in the tool, questionnaires, and interviews with team members. Preliminary results indicate how the tool supports the management of architectural uncertainties and inform its ongoing improvement.

Keywords: Software Architecture. Architectural Uncertainty. Hypotheses Engineering. ArchHypo. Arch-Stage. Empirical Evaluation.

Resumo

José Gonçalves Lima Neto. **Avaliação e Refinamento do Arch-Stage:** ***Apoio a Hipóteses Arquiteturais em Equipes de Engenharia de Software***. Dissertação (Mestrado). Instituto de Matemática, Estatística e Ciência da Computação, Universidade de São Paulo, São Paulo, 2026.

Equipes de engenharia de software frequentemente enfrentam incertezas relacionadas à arquitetura de seus sistemas, mas carecem de métodos e ferramentas adequadas para gerenciá-las de forma explícita e colaborativa. A técnica ArchHypo propõe o uso de engenharia de hipóteses para tornar essas incertezas visíveis e gerenciáveis, porém sua aplicação manual impõe uma curva de aprendizado significativa. Este trabalho avalia e refina empiricamente o Arch-Stage, uma ferramenta de apoio à técnica ArchHypo, por meio de sua aplicação com múltiplas equipes de engenharia de software em uma organização real com um produto em escala. A pesquisa adota um design de métodos mistos, combinando análise de artefatos produzidos na ferramenta, questionários e entrevistas com membros das equipes. Os resultados preliminares indicam como a ferramenta apoia o gerenciamento de incertezas arquiteturais e informam seu aprimoramento contínuo.

Palavras-chave: Arquitetura de Software. Incerteza Arquitetural. Engenharia de Hipóteses. ArchHypo. Arch-Stage. Avaliação Empírica.

List of Abbreviations

SA	Software Architecture
ArchHypo	Architecture Hypothesis (technique)
IDP	Internal Developer Portal
NFR	Non-Functional Requirements
QA	Quality Attribute
DSR	Design Science Research
HE	Hypotheses Engineering
IME	Institute of Mathematics and Statistics (Universidade de São Paulo)
USP	University of São Paulo
ADR	Architecture Decision Record

List of Figures

List of Tables

List of Programs

Contents

1 Introduction **1**

1.1 Narrowing the Scope 2

1.2 ArchHypo and Arch-Stage 2

1.3 Objective 4

1.4 Research Questions 4

1.5 Thesis Structure 4

2 Background **5**

2.1 Software Architecture and Architectural Uncertainty 5

2.2 Hypotheses Engineering and ArchHypo 6

2.2.1 Core Concepts 6

2.2.2 Assessment 7

2.2.3 Technical Plans 7

2.2.4 Lifecycle 8

2.2.5 Empirical Evaluations 8

2.3 Arch-Stage: The Initial Tool 9

2.4 Deriving Architecture from Requirements: The Tooling Gap 10

3 Related Work **13**

3.1 Centralization and Decentralization of Architectural Decisions 13

3.2 Tool Support for Architectural Decision-Making 14

3.3 Empirical Studies of Architectural Practices in Industry 15

4 Research Design **17**

4.1 Empirical Context 17

4.2 Research Approach 17

4.3 Data Sources 18

4.4 Analysis Plan 19

4.5 Ethical Considerations 19

4.6 Alignment with Qualification Feedback 19

Appendixes

A Research Instruments **21**

A.1 Questionnaire 21

A.2 Interview Guide 21

A.3 Informed Consent Form 21

References **23**

Chapter 1

Introduction

Software engineering organizations increasingly strive for fast value delivery through autonomous, empowered teams. Concepts such as Team Topologies and DevOps promote decentralized ownership and rapid feedback loops. Yet architecture decisions—those that constrain the structure and quality attributes of a system—often remain centralized in a handful of experienced practitioners. When requirements evolve continuously, as they do in scaled products, teams face a tension between speed and structural integrity: either they perform extensive upfront design that slows delivery or they allow the system to drift toward what Foote and Yoder termed a "Big Ball of Mud," sacrificing long-term sustainability for short-term velocity Foote and Yoder, 1999.

Understanding why this tension is so persistent requires acknowledging the organizational pressures that teams routinely face. Foote and Yoder identify several forces that push teams toward expedient, unstructured growth rather than deliberate design: tight delivery schedules that preclude architectural reflection, heterogeneous experience levels within teams, the compounding complexity introduced by rapid change, and cost constraints that make dedicated architectural effort difficult to justify in the short term. These pressures are not pathological—they are endemic to competitive software development—and any practical approach to managing architectural uncertainty must operate within them.

Software architecture itself has been defined in multiple ways, reflecting how the discipline has matured. Perry and Wolf's foundational 1992 account introduced a tripartite model in which architecture is constituted by elements (the processing, data, and connecting components), form (the weighted relationships among them), and rationale (the constraints that justify the chosen structure) Perry and Wolf, 1992. This model established that architecture is not merely a set of structural choices but an expression of reasoning under constraints. Decades later, Richards and Ford offered a more practitioner-oriented definition: software architecture encompasses the structure of a system together with the characteristics the system must support, the principles and decisions that govern its design, and the underlying design decisions themselves Richards and Ford, 2020. Both perspectives converge on the view that architectural decisions are inherently bound to uncertainty—uncertainty about which quality characteristics will matter most, which structural choices will best satisfy them, and how the constraints will evolve.

Architectural uncertainty is a natural consequence of operating in environments where business requirements, technology options, and quality-attribute trade-offs are only partially known at any point in time. Despite its pervasiveness, this uncertainty usually stays implicit. Existing methods for deriving and managing software architecture depend heavily on the tacit knowledge of experienced architects, offering limited guidance to less-experienced team members Souza *et al.*, 2019. The systematic mapping by Souza et al. analyzed 39 primary studies on architectural derivation methods and concluded that every method relied on architects' personal experience for its core decisions, with tool support being either absent or only partially effective. This reliance creates a bottleneck: when architectural expertise is scarce, teams either make uninformed decisions or wait for overloaded experts, both of which degrade delivery flow.

1.1 Narrowing the Scope

This thesis does not attempt to address software architecture management in general. Instead, it focuses specifically on **software engineering teams working on a scaled software product** within a single real-world organization. The target problem is the lack of explicit, tool-supported management of architecture-related uncertainties that could enable more shared and decentralized decision-making within and across these teams. By restricting scope to concrete teams and a tangible tool, the work responds directly to feedback received during the author's qualification examination, which recommended narrowing the original research agenda from broad surveys of centralization and decentralization in software architecture to a focused empirical contribution centered on a specific artifact.

1.2 ArchHypo and Arch-Stage

The ArchHypo technique, proposed by Silva et al., reframes software architecture as a collection of testable hypotheses rather than a fixed set of design decisions Silva, Melegati, Silveira, *et al.*, 2025; Silva, Melegati, Wang, *et al.*, 2024. Each architectural hypothesis captures an uncertainty—for instance, whether a chosen database technology will meet scalability demands—along with an assessment of its uncertainty level and potential impact. Teams then formulate technical plans, such as architectural spikes, tracer bullets, or triggers, to reduce uncertainty or defer decisions to a "responsible moment" when more evidence is available.

ArchHypo was developed using a Design Science Research (DSR) methodology and grew out of patterns for agile architecture that were initially disseminated at software engineering conferences and through corporate training programs Silva, Melegati, Wang, *et al.*, 2024. The technique thus carries both a theoretical grounding in DSR and a practical lineage in agile community practice.

ArchHypo has been evaluated in two primary empirical settings. The first evaluation, reported in an IEEE Software article Silva, Melegati, Wang, *et al.*, 2024, examined the Catch Solve startup—a company offering web testing and monitoring services—over a ten-month observation period. Four hypotheses were managed, addressing availability,

reusability of test templates, usability, and scalability of virtual machines. This constituted the first documented application of ArchHypo in a real operational environment. The study demonstrated that making architectural uncertainty explicit guided decisions on when to act versus when to defer, but relied entirely on manual tracking.

The second, more comprehensive evaluation, reported in the IEEE Transactions on Software Engineering Silva, Melegati, Silveira, *et al.*, 2025, took place within Jetsoft, a Brazilian software development company, and focused on PropSEP, a mission-critical budget proposal system for the Planning and Management Secretariat of the State of São Paulo. The system processed approved budget figures on the order of 317 billion BRL (approximately 60 billion USD), underscoring the high stakes of its architectural decisions. The evaluation ran from May to August 2022, spanning six sprints of seven business days each, and involved eight team members (P1–P8). Ten architectural hypotheses were identified, spanning six quality attributes: performance, security, reliability, flexibility, usability, and team productivity. This evaluation produced nine findings (detailed in Chapter 2) that collectively demonstrate ArchHypo's value while surfacing the adoption barriers that motivate the present thesis.

Of particular relevance to the present work are two findings from the TSE evaluation Silva, Melegati, Silveira, *et al.*, 2025. Finding 8 identified that changes in the team's way of working constituted a significant obstacle: mixing architectural and functional tasks, difficulty estimating hypothesis-related work, and the perception that the extra effort was not directly tied to product delivery all created friction. Finding 9, which was unanimous across all eight participants, was that the technique was hard to learn. Team members struggled to map architectural concerns to the hypothesis format, to define appropriate technical plan actions, and to articulate the assumptions underlying a hypothesis. After the initial workshop, no participant felt ready to apply ArchHypo independently without ongoing support.

Arch-Stage (originally named HypoStage) was conceived as a direct response to this tooling gap. Developed as a Backstage Internal Developer Portal plugin in an undergraduate capstone project supervised by the author of this thesis, Arch-Stage provides structured interfaces for creating and editing hypotheses, assessing uncertainty and impact on Likert scales, defining technical plans, and visualizing the portfolio of architectural hypotheses Corrêa, 2025.[1] By embedding these capabilities within an IDP, the tool integrates hypothesis management into the daily engineering workflow rather than keeping it as a separate, easily forgotten activity.

Neither the Catch Solve nor the Jetsoft evaluation involved any tool-mediated support for ArchHypo; both relied exclusively on manual means such as spreadsheets, meetings, and informal documentation. This thesis therefore represents a distinct contribution: an empirical evaluation of ArchHypo in its first tool-supported deployment with real teams.

[1] The initial design and implementation of Arch-Stage, up to commit `ceee509`, is the contribution of Pedro Henrique Mariano Corrêa's capstone project. Subsequent refinements reflect the author's empirical use with real teams.

1.3 Objective

The objective of this work is to empirically evaluate and refine the Arch-Stage tool, a support tool for the ArchHypo technique, by applying it with diverse software engineering teams in a real-world organization with a scaled software product, assessing how it supports managing architectural uncertainty and informing its improvement.

1.4 Research Questions

The following research questions guide this investigation:

RQ1 How do software engineering teams use Arch-Stage to identify and manage architectural hypotheses in practice?

RQ2 What refinements to Arch-Stage are needed to better support teams with varying levels of architectural experience?

RQ3 What types of architectural uncertainties do teams surface when using Arch-Stage, and how do these evolve over time?

1.5 Thesis Structure

The remainder of this thesis is organized as follows. Chapter 2 presents the background on software architecture uncertainty, hypotheses engineering, ArchHypo, the Arch-Stage tool, and the gap in tool support for architectural derivation. Chapter 3 provides a contextual synthesis of related work, drawing on the author's earlier mapping and multivocal review efforts. Chapter 4 describes the research design, including the empirical context, data sources, and analysis approach.

Chapter 2

Background

This chapter establishes the conceptual foundations for the empirical investigation carried out in this thesis. It begins with a concise overview of software architecture and architectural uncertainty in the context of scaled products, then introduces hypotheses engineering and the ArchHypo technique in full technical detail, describes the Arch-Stage tool, and concludes with a summary of findings from Souza et al.'s mapping study on the gap between architectural derivation methods and practical tool support.

2.1 Software Architecture and Architectural Uncertainty

Software architecture refers to the high-level structure of a software system, encompassing its components, their relationships, and the design principles that guide its evolution Bass *et al.*, 2021. In organizations that maintain scaled products—systems developed and operated by multiple teams—architectural decisions shape not only technical quality attributes such as performance, reliability, and maintainability, but also the ability of teams to work independently and deliver changes at speed Skelton and Pais, 2019.

Richards and Ford define software architecture as the structure of a system combined with the characteristics it must support, the principles and decisions guiding its design, and the design decisions themselves Richards and Ford, 2020. This definition emphasizes that architecture is a living body of reasoning, not a frozen blueprint. Perry and Wolf's earlier tripartite model Perry and Wolf, 1992 decomposed architecture into elements (the processing, data, and connecting constituents), form (the weighted relationships among those elements), and rationale (the set of constraints that justify the structure). Together, these accounts foreground the role of *uncertainty*: the rationale behind any architectural choice is valid only relative to the information available at the time of the decision, and that information is always incomplete.

Architectural uncertainty arises whenever teams must make structural decisions under incomplete information. This is a normal condition in continuous-delivery environments where business requirements, user behavior, and technology options are only partially

known. When uncertainty is left implicit, teams either defer decisions indefinitely (accumulating technical debt) or make premature commitments that constrain future options. Both outcomes slow delivery and degrade system quality Foote and Yoder, 1999.

The concept of "fast flow," popularized by the Team Topologies framework, emphasizes that organizations should design both their teams and their architecture to enable a continuous, rapid stream of software changes. Achieving fast flow presupposes that teams can make local architectural decisions without waiting for a central architect. This in turn demands that architectural uncertainties be made explicit and that lightweight mechanisms exist for teams to assess, prioritize, and revisit them as evidence accumulates.

2.2 Hypotheses Engineering and ArchHypo

Hypotheses engineering is an approach to managing technical uncertainty by formulating architectural decisions as testable hypotheses rather than settled facts. Instead of asserting that a particular technology or pattern *will* satisfy a quality requirement, teams state that it *might* do so and then design experiments to gather evidence.

The ArchHypo technique, proposed by Silva et al., operationalizes this idea for software architecture Silva, Melegati, Silveira, *et al.*, 2025. Grounded in Design Science Research (DSR) and evolved from a body of patterns for agile architecture first disseminated at conferences and through corporate training Silva, Melegati, Wang, *et al.*, 2024, ArchHypo provides a structured vocabulary and a set of practices that fit within iterative development workflows.

2.2.1 Core Concepts

An **architectural hypothesis** is a falsifiable statement that captures an uncertainty relevant to the design of a software architecture. It is important to distinguish the concept from adjacent notions:

- **Hypothesis vs. risk.** A risk involves a known event with estimable probability and impact. An architectural hypothesis, by contrast, stems from a *lack of information* sufficient to make a decision. The distinguishing feature is the epistemic gap, not the probability of an outcome Silva, Melegati, Silveira, *et al.*, 2025.
- **Hypothesis vs. requirement.** A requirement specifies what the system must do or how well it must do it. A hypothesis may be triggered by a requirement (e.g., a quality-attribute goal that cannot yet be confirmed as achievable), but the hypothesis captures the uncertainty about the path to satisfying that requirement, not the requirement itself.
- **Hypothesis vs. architectural decision.** An architectural decision resolves a design choice. A hypothesis precedes and informs decisions: it identifies what must be learned or validated before a decision can responsibly be made Silva, Melegati, Silveira, *et al.*, 2025.

Hypotheses arise from two principal sources: (a) *requirements*—when quality attribute targets are poorly specified or the team lacks information about whether a given approach

can satisfy them; and (b) *solutions*—when the consequences of an existing or candidate architectural solution are unknown.

2.2.2 Assessment

Each hypothesis is evaluated along two independent dimensions:

1. **Uncertainty level.** How far the team is from being able to refute or accept the hypothesis. Crucially, uncertainty level is *not* equivalent to probability. A hypothesis that is almost certainly true *and* a hypothesis that is almost certainly false both represent *low* uncertainty, because in either case the team already has enough information to act. High uncertainty means the team genuinely does not know which way the evidence will fall. Both dimensions can be assessed using a 3-point or 5-point qualitative scale; assessments are made independently and relatively, not on absolute probability Silva, Melegati, Silveira, *et al.*, 2025.

2. **Impact level.** The potential consequences—in terms of rework, schedule, cost, or quality degradation—if the hypothesis is invalidated. A high uncertainty combined with a high impact signals that a hypothesis demands urgent attention.

2.2.3 Technical Plans

For hypotheses with significant uncertainty or impact, teams formulate a **technical plan**: a set of concrete actions whose goal is to (a) accept or refute the hypothesis, (b) reduce its uncertainty level, (c) reduce its potential impact, or (d) define criteria under which the decision can responsibly be deferred Silva, Melegati, Silveira, *et al.*, 2025. ArchHypo identifies the following plan strategies:

- **Architectural Spike.** A time-boxed, isolated experiment to explore a technology or design option without coupling the result to the production codebase. The spike produces information rather than deliverable functionality.

- **Tracer Bullet.** A thin, end-to-end implementation within the actual system, exercising all architectural layers to validate an integration or performance characteristic under realistic conditions.

- **Software Analytics.** Using monitoring, profiling, or usage data already collected from the running system to gather evidence about a hypothesis without additional development effort.

- **Modularity / Flexibility.** Restructuring the affected subsystem to reduce coupling or increase replaceability, thereby lowering the *impact* of the hypothesis rather than resolving its uncertainty. This strategy is useful when full resolution is premature but protection against a bad outcome is warranted.

- **Guideline Implementation.** Tailoring the development process—for example, adopting a coding standard, a review protocol, or a testing policy—so that the team systematically generates evidence about the hypothesis as a by-product of normal work. This was the most frequently used strategy in the TSE evaluation Silva, Melegati, Silveira, *et al.*, 2025.

- **Architectural Trigger.** A condition-based deferral in which the team agrees to act on the hypothesis only when a specific observable event occurs (e.g., when the number of concurrent users exceeds a threshold, or when a performance regression appears in monitoring). Until the trigger fires, the hypothesis is deliberately deferred.
- **Agile Landing Zones.** Adopting architectural patterns that explicitly preserve flexibility, creating "landing zones" into which future decisions can be inserted without major disruption, regardless of which direction the evidence eventually points.

2.2.4 Lifecycle

Hypotheses progress through a defined lifecycle: they begin in an *open* state and transition to either *validated* or *invalidated* as evidence from technical plans accumulates. The responsible-moment principle—act when you have enough information, neither too early nor too late—guides when a hypothesis should be resolved rather than monitored further.

2.2.5 Empirical Evaluations

IEEE Software evaluation (Catch Solve startup). The first evaluation of ArchHypo in a real operating environment was reported in an IEEE Software article by Silva et al. Silva, Melegati, Wang, *et al.*, 2024. The setting was Catch Solve, a startup offering web testing and monitoring services. Four hypotheses were formulated and tracked over approximately ten months, covering availability, reusability of test templates, usability, and scalability of virtual machines. The study showed that naming and tracking uncertainties explicitly helped the technical lead decide when to invest in architectural improvements versus when to wait, and validated the overall structure of the ArchHypo approach.

TSE evaluation (PropSEP at Jetsoft). The larger and more controlled evaluation, published in the IEEE Transactions on Software Engineering Silva, Melegati, Silveira, *et al.*, 2025, examined the application of ArchHypo within Jetsoft, a Brazilian software development company, on the PropSEP project—the budget proposal system of the Planning and Management Secretariat of the State of São Paulo. PropSEP is mission-critical: the approved budget it processes amounts to approximately 317 billion BRL. The evaluation ran across six sprints of seven business days each (May–August 2022) and involved eight participants (P1–P8), all holding computer science degrees with certifications in Oracle, WebSphere, or GeneXus. Ten hypotheses were identified, spanning six quality attributes (performance, security, reliability, flexibility, usability, and team productivity), and seven distinct types of architectural practice were employed in the resulting technical plans.

The study produced nine key findings:

1. Some hypotheses carried uncertainty that was monitored throughout the entire project duration without reaching full resolution, demonstrating that the technique accommodates long-horizon uncertainties.
2. ArchHypo enabled the team to avoid upfront architectural design, distributing structural decisions across iterations as evidence became available.

3. Participants reported improvements to the development process in terms of safety, transparency, and predictability.
4. Hypotheses spanning different quality attributes were handled using qualitatively different technical strategies, illustrating the creativity the technique affords.
5. Guideline Implementation (process tailoring) emerged as the most frequently employed strategy across all technical plans.
6. The team collectively recognized that managing hypotheses improved the quality of their architectural decision-making.
7. Post-release evidence pointed to a positive impact on product quality: no issues related to the managed hypotheses were detected after deployment.
8. **Changes in the team's way of working constituted a significant adoption obstacle.** Mixing architectural and hypothesis-related tasks with functional development created estimation difficulties, and participants described the additional work as not feeling directly connected to product delivery.
9. **All eight participants unanimously considered the technique hard to learn.** Specific difficulties included: mapping architectural concerns to the hypothesis format, formulating appropriate technical plan actions, articulating the assumptions behind a hypothesis, and understanding hypotheses as the primary driver of architectural action. No participant felt prepared to apply ArchHypo without continued expert support after the initial workshop.

Findings 8 and 9 directly motivate the tooling contribution of this thesis: a tool that embeds ArchHypo scaffolding in the daily engineering workflow, reducing the cognitive overhead of manual application.

2.3 Arch-Stage: The Initial Tool

Arch-Stage (originally named HypoStage) is an open-source Backstage plugin developed as an undergraduate capstone project by Pedro Henrique Mariano Corrêa, supervised by the author of this thesis and co-supervised by Prof. Paulo Meirelles Corrêa, 2025. Its purpose is to operationalize ArchHypo within an Internal Developer Portal (IDP), embedding hypothesis management into the engineering workflow.

The initial version of Arch-Stage, corresponding to the code up to commit `ceee509` in the `hypo-stage` repository,[1] provides the following capabilities:

- **Hypothesis management**: creating, editing, listing, and deleting architectural hypotheses, each linked to a Backstage catalog entity (e.g., a service, component, or system). Integration with the Backstage catalog means that hypotheses are automatically associated with the software entities they affect, without requiring separate registration.

[1] https://github.com/ArchHypo/hypo-stage

- **Uncertainty and impact assessment**: rating each hypothesis on 5-point Likert scales for both uncertainty level and impact level, with free-text rationale fields. The use of a 5-point scale rather than a binary or 3-point scale provides finer-grained prioritization while remaining tractable for teams without formal risk management backgrounds.
- **Technical planning**: attaching one or more technical plan items to a hypothesis, specifying action type (spike, tracer bullet, trigger, etc.), description, target date, expected outcome, and documentation links.
- **Visualization**: a chart plotting hypotheses on an uncertainty-vs-impact plane to support prioritization, making the relative urgency of hypotheses visible at a glance.
- **Status lifecycle**: tracking each hypothesis through the defined lifecycle states (open → validated or invalidated), preserving the history of status transitions to support retrospective analysis.

By integrating with Backstage, Arch-Stage leverages the portal's service catalog, so that hypotheses are automatically associated with the software entities they affect. This design reduces context-switching: developers manage architectural hypotheses in the same environment where they consult documentation, inspect service health, and access CI/CD pipelines.

Known limitations. The initial tool design assumed familiarity with the ArchHypo vocabulary, which poses a barrier for teams new to hypothesis engineering. The case study in the capstone thesis reconstructed the Catch Solve startup scenario digitally but did not involve real-time use by a team. Furthermore, the tool lacked features for collaborative assessment (e.g., multiple raters per hypothesis) and for tracking the evolution of assessments over time. These gaps motivate the empirical evaluation and refinement effort in this thesis.

2.4 Deriving Architecture from Requirements: The Tooling Gap

The systematic mapping study by Souza, Moreira, and Goulão Souza *et al.*, 2019 examined the existing body of work on methods for deriving architectural models from requirements specifications. Starting from 2,575 candidate papers, the authors applied a PICOC (Population, Intervention, Comparison, Outcome, Context) search structure and a quality assessment checklist—composed of general criteria (weighted at 25%) and method-specific criteria (weighted at 75%)—to select 39 high-quality primary studies for detailed analysis. Their analysis employed the NIMSAD framework to characterize the structure of each reviewed method.

Among their key findings directly relevant to this thesis:

1. **Dependence on tacit knowledge.** Every reviewed method relied, to some degree, on the architect's experience and intuition. No method provided a fully automated or guided derivation path accessible to less-experienced practitioners.

2. **Limited support for less-experienced practitioners.** Methods assumed a skilled architect capable of navigating trade-offs among quality attributes. Teams without senior architectural expertise had little actionable guidance.
3. **Weak or absent tool support.** Although some tools were identified—most notably QuaDAI and Monarch—they addressed only fragments of the derivation process and offered only partial coverage of the architectural reasoning cycle. The mapping found no comprehensive tool integrating hypothesis formulation, assessment, and tracking within a team's daily workflow.
4. **Insufficient evaluation.** Many proposed methods lacked rigorous empirical validation, relying on toy examples or author-run case studies rather than independent evaluations with real teams in production environments.

Souza et al. also noted that 76.9% of the 39 primary studies addressed functional requirements, with a significant proportion also treating non-functional requirements; yet the methods for handling quality attributes—precisely the domain where architectural uncertainty is most consequential—were among the least well supported by tools or empirical evidence.

These findings reinforce the motivation for this thesis on two fronts. First, they confirm the existence of a "tooling void" in which architectural management remains disconnected from the daily engineering workflow. Second, they underscore the need for empirical evaluations with real teams—precisely the kind of evaluation this thesis aims to provide for Arch-Stage and ArchHypo.

Chapter 3

Related Work

This chapter provides a contextual synthesis of literature relevant to the empirical evaluation of Arch-Stage. Rather than presenting a standalone systematic review, it draws on the author's earlier qualification work—which included a tertiary systematic mapping study on software architecture for startups and a multivocal literature review on centralization and decentralization of architectural decision-making—and repositions those findings as narrative context for the present thesis.

3.1 Centralization and Decentralization of Architectural Decisions

A recurring theme in both academic and practitioner literature is the tension between centralizing architectural authority in a few experienced individuals and distributing it across development teams. Organizations pursuing fast flow and team autonomy increasingly favor decentralized models, yet several studies report that without explicit governance mechanisms, decentralization leads to inconsistent decisions and architectural drift.

The author's qualification examination Lima Neto, 2024 explored this tension through two complementary lenses. The first chapter of the qualification presented a tertiary systematic mapping study surveying the existing landscape of secondary studies on software architecture for startup and scale-up organizations, identifying recurring themes in how early-stage teams approach architecture governance under resource and time constraints. The second chapter conducted a multivocal literature review that synthesized both academic publications and grey literature (practitioner blogs, conference talks, and industry reports), covering three interrelated perspectives: deriving software architecture from requirements, the emerging socio-technical architecture trend (which situates architectural decisions within organizational and team structures), and the tension between centralized and decentralized decision-making.

The multivocal review surfaced a framing that the qualification termed the "paradox of software architecture decision-making": industry advocates strongly for decentralized decisions to enable team autonomy and fast flow, yet the methods and tools available

for architectural decision-making depend heavily on the tacit knowledge of experienced architects. Decentralization thus risks transferring decision responsibility to practitioners who lack the conceptual vocabulary or methodological support to exercise it effectively. The qualification committee, however, found this "paradox" framing ultimately unproductive as a primary research framing: the tension it described reflects a recognized design trade-off rather than a genuine logical contradiction, and building a thesis around resolving a paradox risked over-extending the scope of the contribution. This thesis consequently does not employ that framing; instead, it focuses on the concrete, tractable problem of providing tool support for one well-defined uncertainty management technique.

The third chapter of the qualification proposed a questionnaire-based research design to investigate how software engineering teams actually navigate architectural decisions in distributed settings. The proposed questionnaire targeted four dimensions: demographic and organizational context, knowledge of modern software architecture practices, collaboration and decision-making patterns within and across teams, and specific challenges experienced in decentralized architectural decision-making. While this research design would have yielded broad survey data, the qualification committee assessed the scope as too wide and the proposed SMS and multivocal review as insufficient as primary thesis contributions without a concrete artifact or evaluation at the center. The present thesis responds to that feedback by reorienting around the evaluation and refinement of Arch-Stage, a concrete tool artifact, rather than a survey of industry practice.

3.2 Tool Support for Architectural Decision-Making

Several tools and platforms have been proposed to support architectural decision-making:

- **ADR tools** (e.g., adr-tools, Log4brains) focus on recording and versioning Architecture Decision Records (ADRs). They capture the *what* and *why* of decisions but do not model uncertainty, impact, or technical plans for hypotheses that are not yet decisions. ADRs document settled choices; they are not designed to manage the pre-decision state of uncertainty that ArchHypo addresses.
- **Architecture analysis tools** (e.g., ATAM-based workshops, trade-off analysis methods) provide structured evaluation of architectural alternatives but are typically heavyweight, workshop-based activities rather than continuous, team-embedded processes. Their value lies in one-shot or periodic analysis sessions, not in the ongoing, iterative management of uncertainty throughout a product's lifecycle.
- **Internal Developer Portals** such as Backstage, Port, and Cortex offer a catalog-centric view of services and infrastructure. While they improve discoverability and documentation, none of these platforms natively supports hypothesis engineering or architectural uncertainty management. They provide the integration surface that Arch-Stage exploits, but leave the hypothesis engineering layer entirely absent.

Arch-Stage differentiates itself by being explicitly designed around the ArchHypo lifecycle: from hypothesis elicitation through assessment, technical planning, and resolution. Its integration with Backstage places this lifecycle within the developer's daily

tooling environment. No existing tool, to the best of the author's knowledge, provides this combination.

3.3 Empirical Studies of Architectural Practices in Industry

Industry-focused empirical studies on architectural practices remain relatively scarce compared with tool proposals and method descriptions. Among the most relevant:

- The ArchHypo evaluations by Silva et al. Silva, Melegati, Silveira, *et al.*, 2025; Silva, Melegati, Wang, *et al.*, 2024 constitute the closest antecedents to this work. They demonstrated the value of hypothesis engineering in two distinct settings—a web-testing startup and a mission-critical government budget system—but relied entirely on manual application with no tool support. This thesis introduces the first evaluation of ArchHypo in a tool-mediated setting.
- Studies on architectural governance in agile organizations (e.g., the work on SAFe and Spotify-model variants) document how companies structure roles and ceremonies for architecture, but rarely examine explicit uncertainty-management techniques.
- The multivocal review conducted during the author's qualification surfaced practitioner reports (blog posts, conference talks, white papers) describing ad-hoc approaches to architectural uncertainty. These reports corroborate the need for structured methods but, being grey literature, carry weaker evidential value.

This thesis contributes to the body of empirical evidence by evaluating a tool-supported application of ArchHypo with multiple teams, addressing a gap in both the ArchHypo literature (which has not yet studied tool-mediated use) and the broader architectural-practices literature (which lacks studies of hypothesis-engineering tools in scaled settings).

Chapter 4

Research Design

This chapter describes the empirical approach adopted to evaluate and refine the Arch-Stage tool. The design follows a mixed-methods strategy aligned with the thesis objective: understanding how software engineering teams use a tool for managing architectural hypotheses in a real-world setting, and identifying refinements needed to improve its practical utility.

4.1 Empirical Context

The study takes place within **CompanyX**, an organization that develops and operates **ProductY**, a scaled software product serving a significant user base. Multiple software engineering teams collaborate on different subsystems of ProductY, each responsible for distinct functional areas but sharing cross-cutting architectural concerns such as scalability and reliability.

Teams and participants. At least two teams are involved in the study, representing varying levels of architectural experience—from teams with a senior architect to teams where architectural responsibility is distributed informally among developers.

Tool deployment. Arch-Stage was deployed within CompanyX's Backstage instance, making it accessible alongside existing developer portal features. Teams were introduced to the ArchHypo concepts and the Arch-Stage interface through a short onboarding session before beginning their use of the tool.

4.2 Research Approach

The research adopts a **mixed-methods** approach combining qualitative and quantitative data sources. This choice is motivated by the exploratory nature of the study—there are no prior studies of tool-supported ArchHypo use—and by the need to capture both behavioral patterns (how teams use the tool) and perceptual data (how teams experience the tool).

The study can be characterized as an *embedded case study* with multiple units of analysis (the participating teams), following guidelines from Yin and Runeson et al. for empirical software engineering research.

The design is informed by the approach taken in the TSE evaluation of ArchHypo Silva, Melegati, Silveira, *et al.*, 2025, in which a participant-observer role was used to collect multi-source evidence: Kanban board history, user stories, incident reports, the ArchHypo tracking spreadsheet, relevant e-mails, and a questionnaire administered to all team members. That study demonstrated the value of combining artifact-based and perceptual evidence for understanding both how a technique is applied and how participants experience it. The present thesis extends this methodological approach to the tool-mediated setting, where Arch-Stage itself becomes both a data source (its stored artifacts) and an object of evaluation (its usability and fit).

Critically, the TSE evaluation used manual spreadsheets as its primary tracking medium. This thesis evaluates a qualitatively different mode of application: one in which the tool scaffolds hypothesis management, potentially reducing the cognitive overhead and process-change friction identified in Findings 8 and 9 of that study.

4.3 Data Sources

Three primary data sources are planned:

1. **Arch-Stage artifacts.** The hypotheses, assessments, and technical plans recorded by teams in Arch-Stage constitute the principal data set. For each hypothesis, the following attributes are available: statement, quality attribute, source type, uncertainty rating, impact rating, technical plan items (type, description, target date, status), and lifecycle status changes over time. Unlike the spreadsheet-based data of the TSE study, these artifacts are captured in a structured database, enabling systematic extraction and longitudinal analysis.

2. **Questionnaires.** Short questionnaires administered to team members after a defined usage period (e.g., one sprint cycle or one month) to capture perceptions of usability, perceived value, and difficulties encountered.

3. **Semi-structured interviews.** Interviews with a subset of participants (e.g., team leads, developers who created hypotheses, developers who did not) to explore themes emerging from the artifact and questionnaire data in greater depth.

Optionally, the study may include:

- **Observation of meetings.** Attending architecture discussions or sprint-planning sessions where Arch-Stage is referenced, to observe how the tool mediates conversation.

- **Usage logs.** If available, server-side logs of Arch-Stage interactions (page views, edits, time spent) could provide objective usage metrics.

4.4 Analysis Plan

Artifact analysis. Hypotheses and technical plans will be coded thematically to identify recurring types of architectural uncertainty (e.g., technology-fit uncertainties, integration uncertainties, performance uncertainties) and patterns in how teams formulate and resolve them. Uncertainty and impact ratings will be analyzed descriptively (distributions, medians) and, if sample size permits, compared across teams.

Questionnaire analysis. Likert-scale responses will be summarized with descriptive statistics. Open-ended responses will be coded using the same thematic framework applied to the artifact data.

Interview analysis. Interview transcripts will be analyzed using thematic coding, following the six-phase approach of Braun and Clarke. Codes will be organized into themes addressing each research question (tool usage patterns, needed refinements, types of uncertainty).

Triangulation. Findings from the three data sources will be triangulated to strengthen the validity of conclusions. Where artifact data and interview data converge, confidence in the finding increases; where they diverge, the divergence itself becomes a finding requiring interpretation.

4.5 Ethical Considerations

All participants will be asked for informed consent before data collection begins. Company and product names will be anonymized in all publications. Individual participant data (questionnaire responses, interview transcripts) will be stored securely and accessible only to the research team.

4.6 Alignment with Qualification Feedback

This research design responds to the main feedback received during the author's qualification examination Lima Neto, 2024:

- **Narrowed scope**: the study focuses on specific teams and a concrete tool rather than broad surveys of software architecture trends. The qualification committee found the earlier scope—encompassing both a tertiary SMS and a multivocal review as principal contributions—too wide for a master's thesis.
- **Rejection of the "paradox" framing**: the "paradox of SA decision-making" terminology used in the qualification was judged unproductive by the committee, as it overstated the nature of what is in practice a design trade-off. This thesis does not employ that framing.
- **Limited research questions**: three focused questions replace the earlier, broader set.

- **Concrete artifact**: Arch-Stage serves as the tangible object of evaluation, grounding the research in observable behavior and artifacts. The qualification's proposed questionnaire-based survey was repositioned as contextual motivation rather than a primary contribution.
- **Empirical rigor**: the mixed-methods design, triangulation plan, and explicit ethical considerations aim for a level of rigor appropriate for a master's thesis in empirical software engineering. By modeling the data collection strategy on the TSE study's multi-source approach while extending it to the tool-mediated context, the design offers a coherent methodological lineage.

Appendix A

Research Instruments

A.1 Questionnaire

A.2 Interview Guide

A.3 Informed Consent Form

References

[Bass *et al.* 2021] Len Bass, Paul Clements, and Rick Kazman. *Software Architecture in Practice*. 4th ed. Addison-Wesley, 2021 (cit. on p. 5).

[Corrêa 2025] Pedro Henrique Mariano Corrêa. *ArchHypo Tooling: Enabling Practical Hypothesis Engineering for Software Architecture*. Capstone Project Report (Bachelor), Institute of Mathematics and Statistics, University of São Paulo. Supervised by José G. Lima Neto; co-supervised by Paulo Meirelles. 2025 (cit. on pp. 3, 9).

[Foote and Yoder 1999] Brian Foote and Joseph Yoder. "Big ball of mud". *Pattern Languages of Program Design* 4 (1999), pp. 654–692 (cit. on pp. 1, 6).

[Lima Neto 2024] José G. Lima Neto. *Challenges Towards Decentralized Software Architecture Decision-Making for Fast Flow of Requirements*. Qualification Examination, Institute of Mathematics and Statistics, University of São Paulo. Dec. 2024 (cit. on pp. 13, 19).

[Perry and Wolf 1992] Dewayne E. Perry and Alexander L. Wolf. "Foundations for the study of software architecture". *ACM SIGSOFT Software Engineering Notes* 17.4 (1992), pp. 40–52. doi: 10.1145/141874.141884 (cit. on pp. 1, 5).

[Richards and Ford 2020] Mark Richards and Neal Ford. *Fundamentals of Software Architecture: An Engineering Approach*. O'Reilly Media, 2020. isbn: 9781492043454 (cit. on pp. 1, 5).

[Silva, Melegati, Silveira, *et al.* 2025] Eduardo Silva, Jorge Melegati, Moises Silveira, *et al.* "ArchHypo: managing software architecture uncertainty using hypotheses engineering". *IEEE Transactions on Software Engineering* 51.2 (2025), pp. 481–499. doi: 10.1109/TSE.2024.3520477 (cit. on pp. 2, 3, 6–8, 15, 18).

[Silva, Melegati, Wang, *et al.* 2024] Eduardo Silva, Jorge Melegati, Xiaofeng Wang, and Paulo Meirelles. "Using hypotheses to manage technical uncertainty and architecture evolution in a software start-up". *IEEE Software* 41.4 (2024), pp. 78–87. doi: 10.1109/MS.2024.3383628 (cit. on pp. 2, 6, 8, 15).

[Skelton and Pais 2019] Matthew Skelton and Manuel Pais. *Team Topologies: Organizing Business and Technology Teams for Fast Flow*. IT Revolution Press, 2019 (cit. on p. 5).

[Souza *et al.* 2019] Eric Souza, Ana Moreira, and Miguel Goulão. "Deriving architectural models from requirements specifications: a systematic mapping study". *Information and Software Technology* 109 (2019), pp. 26–39. doi: 10.1016/j.infsof.2019.01.004 (cit. on pp. 2, 10).