

**Hypo-Stage: Tool-Supported Architectural
Hypothesis Engineering
in Professional Software Development
Teams — A Case Study**

José Gonçalves Lima Neto

THESIS PRESENTED TO THE
INSTITUTE OF MATHEMATICS,
STATISTICS, AND COMPUTER SCIENCE
OF THE UNIVERSITY OF SÃO PAULO
IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS
FOR THE DEGREE OF
MASTER OF SCIENCE

Program: Computer Science

Advisor: Prof. Dr. Paulo Roberto Miranda Meirelles

São Paulo
June, 2026

**Hypo-Stage: Tool-Supported Architectural
Hypothesis Engineering
in Professional Software Development
Teams — A Case Study**

José Gonçalves Lima Neto

This version of the thesis includes the corrections and modifications suggested by the Examining Committee during the defense of the original version of the work, which took place on June 11, 2026.

A copy of the original version is available at the Institute of Mathematics, Statistics, and Computer Science of the University of São Paulo.

Examining Committee:

Prof. Dr. Eduardo Martins Guerra (Presidente) — UniBz

Prof. Dr. Jorge Augusto Melegati Gonçalves — U.Porto

Prof. Dr. Marco Tulio de Oliveira Valente — UFMG

Prof. Dr. Paulo Roberto Miranda Meirelles (Suplente) — USP

Prof^a. Dr^a. Carla Ilane Moreira Bezerra (Suplente) — UFC

Prof. Dr. Breno Bernard Nicolau de França (Suplente) — Unicamp

*The content of this work is published under the CC BY 4.0 license
(Creative Commons Attribution 4.0 International License)*

Ficha catalográfica elaborada com dados inseridos pelo(a) autor(a)
Biblioteca Carlos Benjamin de Lyra
Instituto de Matemática, Estatística e Ciência da Computação
Universidade de São Paulo

Lima Neto, José Gonçalves

Hypo-Stage: Tool-Supported Architectural Hypothesis
Engineering in Professional Software Development Teams – A
Case Study / José Gonçalves Lima Neto; orientador, Paulo
Roberto Miranda Meirelles. – São Paulo, 2026.

94 p.

Dissertação (Mestrado) – Programa de Pós-Graduação em
Ciência da Computação / Instituto de Matemática, Estatística
e Ciência da Computação / Universidade de São Paulo.

Bibliografia

Versão original

1. Software Architecture. 2. Architectural Uncertainty.
3. Hypotheses Engineering. 4. ArchHypo. 5. Hypo-Stage.
6. Internal Developer Portal. 7. Backstage. 8. Case
Study. 9. Empirical Evaluation. I. Miranda Meirelles,
Paulo Roberto. II. Prof. Dr.

Bibliotecárias do Serviço de Informação e Biblioteca
Carlos Benjamin de Lyra do IME-USP, responsáveis pela
estrutura de catalogação da publicação de acordo com a AACR2:
Maria Lúcia Ribeiro CRB-8/2766; Stela do Nascimento Madruga CRB 8/7534.

This is the original version of the thesis prepared by candidate José Gonçalves Lima Neto,
as submitted to the Examining Committee.

Prof. Dr. Paulo Roberto Miranda Meirelles

Advisor

Institute of Mathematics, Statistics, and Computer Science — University of São Paulo

São Paulo, _____ of _____ of 2026.

*To my wife and family,
whose unwavering support made this journey possible.*

Acknowledgments

*Good decisions come from experience.
Experience comes from making bad decisions.*
— Mark Twain

Special thanks to my advisor, Prof. Dr. Paulo Meirelles, for his exceptional guidance and support throughout this work. His mentorship shaped not only the direction of this research but also the way in which it was conducted. I also extend my gratitude to the staff of IME-USP for their assistance and institutional support, and to the teams who participated in this study for their openness, generosity with their time, and genuine engagement with the research. Finally, I thank my wife and family for their unwavering motivation and encouragement — without them, none of this would have been possible.

Abstract

José Gonçalves Lima Neto. **Hypo-Stage: Tool-Supported Architectural Hypothesis Engineering in Professional Software Development Teams — A Case Study**. Thesis (Master's). Institute of Mathematics, Statistics, and Computer Science, University of São Paulo, São Paulo, 2026.

Software engineering teams in continuous-delivery organizations often encounter architectural decisions under incomplete information, yet few lightweight tools exist to help them record, prioritize, and learn from those uncertainties collaboratively. The ArchHypo technique addresses this challenge by applying hypotheses engineering to software architecture context by enabling teams to formulate uncertainties as falsifiable hypotheses, assess each by uncertainty level and impact, and define a technical plan for handling them. However, prior works have reported a substantial learning curve and dependence on expert facilitation, both augmented by the absence of tool support. This work reports a case study that evaluates and iteratively refines a tool-mediated application of ArchHypo in a professional software engineering organization, centered on Hypo-Stage, an open-source Backstage plugin deployed on the organization's environment. The study adopts a flexible, multiple case-study design with formative evaluation and action-research elements in line with flexible-design guidance (ROBSON and McCARTAN, 2016), involving two teams — a feature-delivery (Product) team of twelve engineers and an internal-tooling (Platform) team of four — over five 2-week sprints, with data collected through participant observation, artifact analysis of thirteen registered hypotheses and their technical plans, and open-ended active listening sessions with the involved senior technical leaders. Four contributions emerge from this work: (1) a detailed observational report of how both teams adopted and used Hypo-Stage, characterizing their hypothesis registration dynamics, role stratification, and tool adoption patterns; (2) a characterization of the architectural uncertainties the teams made explicit, including their sources, dominant quality attributes, uncertainty-impact profiles, and technical plan strategies; (3) five Structural Hypothesis Patterns — Hypothesis Cluster, Regulatory Trigger Hypothesis, Service Decomposition Hypothesis, Monitored Quality-Attribute Hypothesis, and Shared-Resource Hypothesis; and (4) nine categories of empirically grounded refinements to Hypo-Stage, each traced to a concrete observation, including a hypothesis evolution chart, a prioritization dashboard, and inline formulation guidance that reduces the technique's learning barrier.

Keywords: Software Architecture; Architectural Uncertainty; Hypotheses Engineering; ArchHypo; Hypo-Stage; Internal Developer Portal; Backstage; Case Study; Empirical Evaluation.

Resumo

José Gonçalves Lima Neto. **Hypo-Stage: Engenharia de Hipóteses Arquiteturais Apoiada por Ferramenta em Equipes Profissionais de Desenvolvimento de Software — Um Estudo de Caso**. Dissertação (Mestrado). Instituto de Matemática, Estatística e Ciência da Computação, Universidade de São Paulo, São Paulo, 2026.

Equipes de engenharia de software em organizações de entrega contínua enfrentam recorrentes decisões arquiteturais diante de informações incompletas, porém poucas ferramentas existem para ajudá-las a registrar, priorizar e aprender com essas incertezas de forma colaborativa. A técnica ArchHypo aborda esse desafio ao aplicar a engenharia de hipóteses ao contexto da arquitetura de software, permitindo que as equipes formulem incertezas como hipóteses falseáveis, avaliem cada uma pelo nível de incerteza e impacto, e definam um plano técnico para tratá-las. No entanto, trabalhos anteriores relataram uma curva de aprendizado acentuada e dependência de facilitação especializada, ambas intensificadas pela ausência de suporte de ferramentas. O presente trabalho relata um estudo de caso que avalia e refina iterativamente uma aplicação mediada por ferramenta do ArchHypo em uma organização profissional de engenharia de software, centrado no Hypo-Stage, um plugin de código aberto para o Backstage implantado no ambiente da organização. O estudo adota um design de estudo de caso múltiplo e flexível, com avaliação formativa e elementos de pesquisa-ação alinhados à orientação sobre designs flexíveis (ROBSON e MCCARTAN, 2016), envolvendo duas equipes — uma equipe de entrega de funcionalidades (Produto) de doze engenheiros e uma equipe de ferramentas internas (Plataforma) de quatro — ao longo de cinco sprints de duas semanas, com dados coletados por meio de observação participante, análise de artefatos de treze hipóteses registradas e seus respectivos planos técnicos, e sessões abertas de escuta ativa com os líderes técnicos seniores envolvidos. Quatro contribuições emergem deste trabalho: (1) um relatório observacional detalhado de como ambas as equipes adotaram e utilizaram o Hypo-Stage, caracterizando suas dinâmicas de registro de hipóteses, estratificação de papéis e padrões de adoção da ferramenta; (2) uma caracterização das incertezas arquiteturais que as equipes tornaram explícitas, incluindo suas fontes, atributos de qualidade dominantes, perfis de incerteza-impacto e estratégias de planos técnicos; (3) cinco novos padrões emergentes de hipóteses arquiteturais técnicas — *Hypothesis Cluster*, *Regulatory Trigger Hypothesis*, *Service Decomposition Hypothesis*, *Monitored Quality-Attribute Hypothesis* e *Shared-Resource Hypothesis*; e (4) nove categorias de refinamentos fundamentados empiricamente para o Hypo-Stage, cada um rastreado a uma observação concreta, incluindo um gráfico de evolução de hipóteses, um painel de priorização e orientações de formulação *inline* que reduzem a barreira de aprendizado da técnica.

Palavras-chave: Arquitetura de Software; Incerteza Arquitetural; Engenharia de Hipóteses; ArchHypo; Hypo-Stage; Portal do Desenvolvedor Interno; Backstage; Estudo de Caso; Avaliação Empírica.

List of Abbreviations

SA	Software Architecture
ArchHypo	Architecture Hypothesis (technique)
IDP	Internal Developer Portal
NFR	Non-Functional Requirements
QA	Quality Attribute
DSR	Design Science Research
HE	Hypotheses Engineering
IME	Institute of Mathematics and Statistics (Universidade de São Paulo)
USP	University of São Paulo
ADR	Architecture Decision Record

List of Figures

2.1	Initial Hypo-Stage hypotheses dashboard (Backstage plugin), corresponding to the CORRÊA 2025 baseline at commit ceee509 . The listing exposes each hypothesis with uncertainty, impact, status, source type, and creation date.	19
4.1	Conceptual framework of the study.	32
4.2	Overview of the ArchHypo technique showing the iterative cycle through which engineers identify architectural hypotheses, assess their uncertainty and impact, define technical plans, and take or postpone architectural decisions. Source: adapted from SILVA, MELEGATI, SILVEIRA, et al. (2025)	36
5.1	EntityHypothesesTab: Hypo-Stage surfaces only the hypotheses linked to the Backstage catalog component the engineer is currently viewing, eliminating the need to navigate to a separate tool or filter by component.	48
5.2	Prioritization Dashboard in Hypo-Stage: the focus panel surfaces hypotheses according to the urgency signals derived from their uncertainty–impact combination, operationalizing the Postpone Uncertain Decisions pattern from the ArchHypo pattern literature (SILVA, ADOLFO, et al., 2024).	55
5.3	Evolution Chart in Hypo-Stage: tracking uncertainty and impact rating changes across successive assessments for a registered hypothesis. The screenshot is illustrative—a representative example of the chart rather than a depiction of one specific hypothesis from the study dataset. Each point corresponds to a saved assessment; the chart makes visible whether technical plan actions are reducing uncertainty or impact over time.	56
5.4	Hypothesis creation form in Hypo-Stage after the usability-oriented form refinements in this section: inline contextual guidance adjacent to the Statement field scaffolds the formulation of falsifiable hypotheses, reducing the cognitive barrier at the moment of registration.	57

List of Tables

1.1	Alignment between research questions and study objectives.	6
2.1	Distinguishing architectural hypotheses from adjacent concepts (SILVA, MELEGATI, SILVEIRA, <i>et al.</i> , 2025).	11
2.2	ArchHypo technical plan strategies and their primary effect (SILVA, MELEGATI, WANG, <i>et al.</i> , 2024; SILVA, MELEGATI, SILVEIRA, <i>et al.</i> , 2025).	13
2.3	Summary of prior empirical evaluations of ArchHypo.	16
2.4	Backstage architecture layers and their role (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024).	18
2.5	Hypo-Stage initial capabilities and known limitations identified prior to this study (CORRÊA, 2025).	20
3.1	Comparison of tools relevant to architectural decision-making support. .	27
4.1	Case comparison: Team 1 (Product) and Team 2 (Platform).	34
4.2	Data collection instruments, sources, and research questions addressed. .	38
4.3	Trustworthiness criteria and strategies applied in this study (ROBSON and McCARTAN, 2016).	41
5.1	From open codes to chapter sections: traceability of the thematic coding process.	46
5.2	Overview of the thirteen registered architectural hypotheses. Req. = Requirements-sourced; Sol. = Solution-sourced. U = Uncertainty level; I = Impact level. Ratings: VL = Very Low, M = Medium, H = High, VH = Very High.	49
5.3	Quality attribute frequency across the thirteen hypotheses.	50
5.4	Technical plan strategies observed across the hypothesis data.	51
5.5	Refinements to Hypo-Stage mapped to motivating observations and research questions. RQ columns indicate primary (P) or secondary (S) linkage. .	58

5.6	Emerging findings identified in the case study: the four Adoption-Condition Candidates (A1–A4) and the five Structural Hypothesis Patterns (B1–B5). Both families describe problem and context only; solution-shaped guidance for each item is deferred to the future research agenda (Section 6.3).	64
5.7	Correspondence between the research questions (Chapter 1) and the organization of evidence in this chapter.	65
B.1	Anonymized Backstage entity references and their descriptions.	80

List of Programs

Hypo-Stage — open-source Backstage plugin that operationalizes the ArchHypo technique for hypothesis management; empirical artifact of this thesis. Repository: <https://github.com/ArchHypo/hypo-stage>.

Backstage — open-source framework for Internal Developer Portals (CNCF); organizational platform in which Hypo-Stage is deployed in this study. Repository: <https://github.com/backstage/backstage>.

Contents

List of Programs	xviii
1 Introduction	1
1.1 Architectural Uncertainty and Hypotheses Engineering	2
1.2 From Requirements to Architecture and the Role of Tool Support	3
1.3 Problem Statement	4
1.4 Objectives	4
1.5 Research Questions	5
1.6 Research Contributions	5
1.7 Dissertation Structure	7
2 Background	9
2.1 Software Architecture and Architectural Uncertainty	9
2.2 Hypotheses Engineering and ArchHypo	10
2.2.1 Core Concepts	10
2.2.2 Assessment	12
2.2.3 Technical Plans	12
2.2.4 Lifecycle	14
2.2.5 Empirical Evaluations	14
2.3 Internal Developer Portals and Backstage	15
2.3.1 The Role of Internal Developer Portals (IDPs)	16
2.3.2 Backstage: Open Platform for Internal Developer Portals	17
2.4 Hypo-Stage: The Initial Tool	18
2.5 Deriving Architecture from Requirements: The Tooling Gap	20
2.6 Summary: State of the Literature for This Work	21
3 Literature Context and Positioning	23
3.1 Centralization and Decentralization of Architectural Decisions	23
3.2 Socio-Technical Architecture	24

3.2.1	Conway's Law and Its Implications	24
3.2.2	Socio-Technical Architecture as an Industry Trend	25
3.3	Tool Support for Architectural Decision-Making	26
3.4	Empirical Studies of Architectural Practices in Industry	26
4	Research Design	29
4.1	Research Strategy	29
4.1.1	Key Methodological Concepts	30
4.2	Conceptual Research Framework	32
4.3	Research Context	33
4.3.1	Case Selection	33
4.4	Researcher's Role	34
4.5	Procedure	35
4.5.1	Phase 1 – Training and Introduction	35
4.5.2	Phase 2 – Observation and Iterative Refinement	35
4.5.3	Phase 3 – Active Listening and Wrap-up	37
4.6	Data Collection	37
4.6.1	Participant Observation	37
4.6.2	Artifact Analysis	38
4.6.3	Open-Ended Active Listening Sessions	39
4.7	Data Analysis	39
4.8	Establishing Trustworthiness	40
4.9	Ethical Considerations	42
5	Results and Refinements	45
5.1	Research Methodology Traceability	45
5.2	RQ1 – How Teams Used Hypo-Stage in Practice	46
5.2.1	Hypothesis Registration Dynamics	46
5.2.2	Role-Based Engagement	47
5.2.3	Tool Adoption Patterns	47
5.3	RQ2 – Architectural Uncertainties Registered	48
5.3.1	Data Overview	48
5.3.2	Source Types	48
5.3.3	Quality Attribute Distribution	49
5.3.4	Uncertainty and Impact Profile	49
5.3.5	Technical Plans in Use	50
5.3.6	Lifecycle and Status	51
5.4	RQ3 – Perceived Benefits and Challenges	51

5.4.1	Perceived Benefits	51
5.4.2	Perceived Challenges	52
5.5	RQ4 — Refinements to Hypo-Stage	53
5.5.1	Backstage Integration and Catalog Surfacing	53
5.5.2	Backend API and Filtering	53
5.5.3	Prioritization Dashboard	54
5.5.4	Hypothesis Lifecycle Consistency	54
5.5.5	Technical-Planning-Driven Assessment Updates	55
5.5.6	Evolution Chart Improvements	55
5.5.7	Form Usability Improvements	56
5.5.8	Summary: Refinements Mapped to Observations	56
5.6	Emerging Findings from the Case Study	57
5.6.1	Adoption-Condition Candidates for Hypothesis Engineering	58
5.6.1.1	Senior-Curated Hypothesis Backlog	59
5.6.1.2	Long-Horizon Registration Priority	59
5.6.1.3	Registration as a Slack-Dependent Practice	59
5.6.1.4	Negative-Formulation Barrier	60
5.6.2	Structural Hypothesis Patterns for Hypothesis Engineers	60
5.6.2.1	Hypothesis Cluster	61
5.6.2.2	Regulatory Trigger Hypothesis	61
5.6.2.3	Service Decomposition Hypothesis	61
5.6.2.4	Monitored Quality-Attribute Hypothesis	62
5.6.2.5	Shared-Resource Hypothesis	63
5.7	Revisiting the Research Questions	63
5.7.1	Correspondence between research questions and chapter sections	63
5.7.2	Brief interpretive synthesis	63
6	Conclusion	67
6.1	Synthesis of Contributions	68
6.2	Implications for Practice and Research	69
6.3	Limitations and Future Work	70
6.3.1	Scope and Transferability	70
6.3.2	Dataset Scope and Hypothesis Lifecycle	71
6.3.3	Researcher Role and Potential Bias	71
6.3.4	Deviation from Prescribed Hypothesis Formulation	71
6.3.5	Future Work Agenda	72

Appendixes

A	Research Instruments	75
A.1	Open-Ended Active Listening Guide	75
A.2	Informed Consent Framework	76
A.3	Training Communication	77
B	Registered Hypotheses Data	79
B.1	Anonymization Legend	79
B.2	Hypothesis Records	80
B.3	Technical Plans Catalog	87

	References	91
--	-------------------	-----------

Chapter 1

Introduction

Software architecture (SA) plays a central role in the development and evolution of software-intensive systems. The SA concept structures a system into elements and relations that form the basis of its organization, conveying coherence, scalability, and maintainability for the evolving software over time (PERRY and WOLF, 1992). Despite the ever-changing practices of modern software development, SA still provides the basis for analyzing the quality attributes of modern systems, guiding technical decisions made throughout the development process via a diverse set of architectural methodologies (SOUZA *et al.*, 2019). In organizations that rely on software as a core part of their business, architecture is not a one-time design artifact, but rather a continuous decision-making process that must accommodate changing requirements, evolving technologies, and critical organizational constraints (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).

In recent years, agile and continuous development practices have intensified this fast-paced pressure. Teams are expected to deliver more features in shorter cycles while preserving scalability, performance, and other quality attributes of their systems that continue to grow in size and complexity, as observed in SILVA, MELEGATI, SILVEIRA, *et al.* (2025) and SILVA, MELEGATI, WANG, *et al.* (2024). In such environments, architecture decisions are frequently made under uncertainty: solutions or requirements may be incomplete or unstable, the impact of technology choices may not be fully understood, and the behavior of the system under realistic workloads may only become visible after deployment (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). Despite such challenges, architectural uncertainties often remain implicit and unmanaged, and decisions tend to rely on the tacit knowledge and intuition of a small group of experienced architects, as summarized in systematic evidence on architectural derivation methods (SOUZA *et al.*, 2019).

At the same time, many organizations seek to organize themselves around autonomous, cross-functional teams responsible for end-to-end delivery of features. Recent industry reports, such as the 2024 InfoQ Software Architecture and Design Trends Report (BETTS *et al.*, 2024), highlight socio-technical architecture as an emerging trend that emphasizes the co-design of technical and organizational structures, and the importance of considering how team boundaries and communication patterns shape system architectures (HERBSLEB and GRINTER, 1999). In such a scenario, architecture work is distributed across teams rather

than centralized in a separate function or role. Moreover, in these scenarios, the Conway's Law is treated as a design constraint to be intentionally leveraged rather than an accident to be ignored (BETTS *et al.*, 2024). However, when architecture decisions depend on implicit assumptions held by a few experts in the organization, this distribution becomes difficult in practice. Less experienced team members may struggle to participate meaningfully in architectural discussions, and teams may find it hard to reason about the risks and trade-offs of alternative solutions even when they are applying any one of the well-known and validated architectural methodologies mapped in prior work (SOUZA *et al.*, 2019). This challenge reveals the case of the architectural uncertainty management, which still is a critical aspect of the architecture work, but is often neglected or poorly supported in practice by the available tools and methodologies.

1.1 Architectural Uncertainty and Hypotheses Engineering

Uncertainty in architecture decisions proceeds from limited knowledge about requirements, future usage scenarios, and the properties of candidate solutions as reviewed in SILVA, MELEGATI, SILVEIRA, *et al.* (2025). Traditional risk management techniques assume that relevant events and their probabilities are known or can be estimated, which is not always the case when dealing with evolving systems and emerging technologies (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). A complementary perspective is to acknowledge uncertainty explicitly and as an inevitable aspect of the architecture work, and therefore, to manage it as a first-class concern during the architecture work as proposed by SILVA, MELEGATI, SILVEIRA, *et al.* (2025).

In this regard, Hypotheses Engineering has emerged as an approach to manage uncertainty in software development by treating assumptions as explicit, falsifiable hypotheses that can be prioritized and tested through experiments (MELEGATI, WANG, and ABRAHAMSON, 2019; OLSSON and BOSCH, 2014; FAGERHOLM *et al.*, 2017; SILVA, MELEGATI, SILVEIRA, *et al.*, 2025; SILVA, MELEGATI, WANG, *et al.*, 2024). In this context, ArchHypo is the technique that adapts hypotheses engineering to software architecture. It proposes that teams identify *architectural hypotheses* — statements capturing uncertainties that may affect architectural decisions — assess each hypothesis by its uncertainty level and impact, and define a technical plan for handling it (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).

The technical plan may include actions to reduce uncertainty, such as experiments, architectural spikes, or tracer bullets, and actions to reduce the impact of potential changes, such as modularization or flexible deployment strategies (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). By iteratively updating hypotheses and technical plans over time, ArchHypo aims to support a more deliberate and distributed approach to architectural decision-making, allowing teams to postpone irreversible decisions when justified and to focus their efforts on the most critical uncertainties (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025; SILVA, MELEGATI, WANG, *et al.*, 2024).

Empirical studies have started to explore how ArchHypo works in practice. An experience report in a software start-up describes how the technique helped the founder identify and prioritize architectural uncertainties related to availability, usability, and

scalability over a ten-month period, guiding the evolution of the product and avoiding premature investments in infrastructure (SILVA, MELEGATI, WANG, *et al.*, 2024). A design science study in a mission-critical public-sector system shows that integrating ArchHypo into the development process enabled the team to distribute architectural work across iterations and to make better-informed decisions in the face of changing requirements (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). At the same time, both studies report that the technique introduces a non-trivial learning curve and requires adjustments to existing processes, which might pose a challenge to its adoption by the team as reported in SILVA, MELEGATI, SILVEIRA, *et al.* (2025) and SILVA, MELEGATI, WANG, *et al.* (2024).

1.2 From Requirements to Architecture and the Role of Tool Support

The challenges of managing architectural uncertainty are closely related to the long-standing problem of deriving architectural models from requirements specifications. SOUZA *et al.* (2019) reviewed 39 methods for this derivation and concluded that existing approaches still rely heavily on the tacit knowledge of architects, provide limited support for less experienced practitioners, and often lack explicit evaluation mechanisms. Many methods describe high-level processes or frameworks, but offer little guidance on how to operationalize them in everyday practice as reported in SOUZA *et al.* (2019).

The same study highlighted gaps in tool support. A significant portion of the methods did not include any supporting tool, and when tools were present, they tended to address only parts of the process as also reported in SOUZA *et al.* (2019). The authors called for more automation, better mechanisms to capture and reuse architectural knowledge, and empirical validation of tools in realistic settings. These findings resonate with the observations from ArchHypo studies: while the technique provides a structured way to make architectural uncertainties explicit and to plan actions, applying it in practice is highly demanding without adequate supporting tool (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025; SILVA, MELEGATI, WANG, *et al.*, 2024).

Grounded in these studies' contributions, an initial version of Hypo-Stage — a support tool designed to help teams apply the ArchHypo technique in practice — was developed as prior work (CORRÊA, 2025); this present work takes that tool as its starting artifact rather than as one of its contributions. Up to a specific point in its history, the work focused on designing and implementing the core functionality for enabling ArchHypo capabilities, such as representing architectural hypotheses, recording their uncertainty and impact assessments, and managing associated technical plans.¹ The work also discussed initial design decisions and limitations, including usability aspects and open questions about how the tool would fit into different team workflows. Chapter 2 presents a more detailed description of Hypo-Stage and its associated contributions (CORRÊA, 2025).

¹The first version of Hypo-Stage and the corresponding source code history up to commit `ceee509776508081ebdd473c2c4f710b8ef55947` document the initial design and implementation of Hypo-Stage (CORRÊA, 2025).

1.3 Problem Statement

When reviewed in intersection, the reviewed studies in this present work outline a clear problem. On the one hand, modern software organizations need to manage architectural uncertainty in a way that supports continuous evolution and involves the broader development team. On the other hand, empirical evidence and systematic mapping studies indicate that:

- Methods for deriving and evolving architectures from requirements still depend strongly on architects' tacit knowledge and lack comprehensive tool support as reported in [SOUZA *et al.* \(2019\)](#).
- ArchHypo offers a promising structure for handling architectural uncertainty, but teams perceive the technique as hard to learn and requiring process changes, especially without specialized support as reported in [SILVA, MELEGATI, SILVEIRA, *et al.* \(2025\)](#) and [SILVA, MELEGATI, WANG, *et al.* \(2024\)](#).
- Hypo-Stage provides an initial implementation of tool support for ArchHypo, but its impact on real software engineering teams and the refinements needed for effective use in industry remain unknown as reported in [CORRÊA \(2025\)](#).

In this context, there is a need to understand how software engineering teams, with diverse roles, tenures, and responsibilities, interact with a tool like Hypo-Stage when working on their products and services. By addressing this problem, this work seeks to gain an understanding of what architectural uncertainties teams manage through such the tool, and which aspects of Hypo-Stage and its integration into the development process still require refinement. Investigating these concerns can contribute both to the practical adoption of ArchHypo via Hypo-Stage and to the broader discussion on how to support teams in sharing architectural responsibilities for accelerating adoption of the socio-technical architecture paradigm mapped in [BETTS *et al.* \(2024\)](#).

1.4 Objectives

The overall objective of this work is to empirically evaluate and refine Hypo-Stage, the tool-supported implementation of the ArchHypo technique introduced in prior Hypo-Stage work ([CORRÊA, 2025](#)), by applying it with diverse professional software development teams in a real-world organizational context that maintains a scaled software product in Latin America. More specifically, this work aims to:

- Investigate how software engineering teams appropriate and use Hypo-Stage and ArchHypo in their day-to-day activities when working on an existing, scaled product.
- Characterize the architectural hypotheses that teams register and manage through Hypo-Stage, identifying emerging patterns among them.
- Identify benefits and challenges perceived by team members regarding the use of Hypo-Stage, including their engagement with the tool, usability, and fit with existing practices.

- Derive and implement refinements to Hypo-Stage and its usage guidelines based on empirical findings, with the goal of improving its suitability for supporting architectural work in teams.

These objectives deliberately build on, rather than replicate, the previous studies' contributions in the ArchHypo research area. While CORRÊA (2025) focused on designing and implementing Hypo-Stage, this work treats the tool as an existing artifact and concentrates on its evaluation and evolution in a professional software development context.

1.5 Research Questions

To achieve the above objectives, this work is guided by four research questions, presented in the same order in which they are addressed in Chapter 5. Their mapping to the objectives is summarized in Table 1.1.

- RQ1:** How do software engineering teams use Hypo-Stage and the ArchHypo technique when working on a scaled software product in a real-world organization?
- RQ2:** What types and recurring structural shapes of architectural uncertainties do these teams identify, record, and manage using Hypo-Stage, and what patterns emerge from their content and organizational context?
- RQ3:** What benefits and challenges do team members perceive in using Hypo-Stage to support architectural decision-making and uncertainty management?
- RQ4:** Which refinements to Hypo-Stage and its usage guidelines emerge from the empirical study, and how do they address the identified challenges?

Chapter 4 operationalizes these questions through aligned data sources: participant observation and Hypo-Stage artifacts support RQ1 (tool use) and RQ2 (recorded uncertainties); active listening foregrounds RQ3 (perceived benefits and challenges) while also surfacing needs that motivate refinements under RQ4.

1.6 Research Contributions

By answering its set of research questions, this work makes the following contributions:

- An empirical account of how multiple professional software development teams use Hypo-Stage — the tool-supported implementation of ArchHypo (CORRÊA, 2025) — while working on a scaled software product, including observed adoption patterns, practices, and obstacles.
- A characterization of the architectural uncertainties these teams made explicit and managed through Hypo-Stage, describing their sources, quality attributes, structural archetypes, and evolution.
- A set of novel *Structural Hypothesis Patterns* derived from thematic analysis of the hypotheses registered during the study. These patterns — Hypothesis Cluster, Regulatory Trigger Hypothesis, Service Decomposition Hypothesis, Monitored

RQ	Research Question	Primary Objective Addressed
RQ1	How do software engineering teams use Hypo-Stage and ArchHypo when working on a scaled product?	Investigate appropriation and use patterns of Hypo-Stage in day-to-day team activities.
RQ2	What types and recurring structural shapes of architectural uncertainties do teams identify, record, and manage using Hypo-Stage, and what patterns emerge from their content and organizational context?	Characterize architectural uncertainties by source, quality attribute, and lifecycle; identify recurring structural archetypes of hypotheses that emerge from the empirical data as a novel contribution to the ArchHypo pattern literature.
RQ3	What benefits and challenges do team members perceive in using Hypo-Stage?	Identify perceived benefits and challenges related to engagement with the tool, usability, and workflow fit.
RQ4	Which refinements emerge from the empirical study and how do they address identified challenges?	Derive and implement empirically grounded refinements to Hypo-Stage and its usage guidelines.

Table 1.1: *Alignment between research questions and study objectives.*

Quality-Attribute Hypothesis, and Shared-Resource Hypothesis — constitute the first empirically grounded answer to the question of what recurring shapes architectural hypotheses take when ArchHypo is applied in practice, addressing a gap explicitly identified as future work in the prior ArchHypo pattern literature (SILVA, ADOLFO, *et al.*, 2024).

- A set of empirically grounded refinements to the Hypo-Stage tool and its usage guidelines, aimed at improving the tool’s usability and engagement, and better integrating the technique into team workflows as reported in CORRÊA (2025).
- Reflections on the implications of these findings for supporting distributed architectural work in agile and continuous development settings, connecting the ArchHypo experience to broader concerns about tacit knowledge and tool support in software architecture as reported in SOUZA *et al.* (2019) and BETTS *et al.* (2024).

Preview of key findings. This study produced three families of concrete results. First, hypothesis registration proved consistently senior-led in both teams: twelve of the thirteen hypotheses were authored by senior or staff engineers, and sprint delivery pressure was the dominant barrier to broader participation — a finding that directly motivates both the tool improvements of RQ4 and the adoption-condition candidate *Registration as a Slack-Dependent Practice* (A3, Section 5.6.1). Second, the thirteen hypothesis profiles reveal a strong concentration of compliance-driven and shared-infrastructure uncertainties, producing two of the five emerging patterns (*Regulatory Trigger Hypothesis* and *Shared-Resource Hypothesis*) that are novel contributions to the ArchHypo pattern literature. Third, the refinement cycle produced nine categories of empirically grounded improvements

to Hypo-Stage, including an evolution chart for tracking assessment changes over time. These findings are previewed here to give the reader an early overview. A complete result analysis and discussion appear in Chapter 5 and Chapter 6.

1.7 Dissertation Structure

The remainder of this work is organized as follows.

Chapter 2 provides the theoretical foundations for the study. It introduces key concepts in software architecture and architectural uncertainty; presents the ArchHypo technique in detail, including its core concepts, assessment model, technical plan strategies, and lifecycle; reviews the two prior empirical evaluations of ArchHypo (at Catch Solve and PropSEP/Jetsoft); describes Internal Developer Portals and the Backstage platform as the delivery vehicle for Hypo-Stage; presents the initial Hypo-Stage tool (CORRÊA, 2025); and surveys systematic evidence on the gap between architectural derivation methods and tool support (SOUZA *et al.*, 2019).

Chapter 3 synthesizes related literature and industry context relevant to tool-supported architectural decision-making and positions this work within that space.

Chapter 4 presents the research design and methods adopted in this work. It describes the research strategy (a flexible multiple case study with formative evaluation and action research elements), the conceptual research framework that organizes the key constructs of the study, the organizational context and case selection rationale for the two participating teams (a Product Team and a Platform Team), the researcher’s dual role as participant observer and facilitator, the data collection procedures (participant observation, artifact analysis, and open-ended active listening sessions), the thematic coding analysis approach, and the criteria used to establish trustworthiness and ethical conduct.

Chapter 5 presents the empirical results and refinements of the study. The chapter follows the research questions in order: how teams used Hypo-Stage and ArchHypo in practice (RQ1); the architectural uncertainties and hypotheses they recorded, including aggregate data characterization and the Structural Hypothesis Patterns — the recurring structural shapes that hypotheses take, developed in Chapter 5 (RQ2); the benefits and challenges they perceived (RQ3); and the refinements to Hypo-Stage and its usage guidelines that emerged from the study, each traced to observations and linked primarily to (RQ4). Presenting results and refinements together makes explicit how empirical findings have driven iterative tool evolution throughout the observation period.

Finally, Chapter 6 builds on the integrated empirical account in Chapter 5. It synthesizes findings in relation to the research questions and to prior ArchHypo evidence, discusses implications for practitioners adopting hypothesis-engineering tools and for researchers studying tool-mediated architectural practices, and outlines limitations together with directions for future work on Hypo-Stage and on broader support for distributed architectural decision-making in agile organizations.

Chapter 2

Background

This chapter establishes the theoretical foundations for the empirical investigation carried out in this work. It begins with an overview of software architecture and architectural uncertainty, then introduces hypotheses engineering and the ArchHypo framework in technical details, describes Internal Developer Portals and their role as platforms for enabling tools such as Hypo-Stage, presents the Hypo-Stage tool itself, and concludes with a summary of findings from the systematic mapping study that mapped the gaps between architectural derivation methods and practical tool support (SOUZA *et al.*, 2019).

2.1 Software Architecture and Architectural Uncertainty

Software architecture refers to the high-level structure of a software system, encompassing its components, their relationships, and the design principles that guide its evolution (L. BASS *et al.*, 2021). However, in organizations that maintain software at the core of their business developed and operated by multiple teams, the ever-happening architectural decisions shape not only technical quality attributes such as performance, reliability, and maintainability of systems, but also the ability of engineering teams to work independently and deliver changes quickly and efficiently at a speed that needs to match the speed of the business (SKELTON and PAIS, 2019).

More embracing to that rapidly changing business scenarios, RICHARDS and FORD (2020) added that the design decisions themselves are part of the software architecture itself, making it not only a static structure or frozen blueprint of a system, but a living body of reasoning that is constantly being refined and adapted to the changing business needs. This framing is adopted in this work to ground the concept of Architectural Uncertainty: the rationale behind any architectural choice is valid only relative to the information available at the time of the decision, and that information is always incomplete and subject to change.

In other words, Architectural Uncertainty arises whenever teams must make structural decisions under incomplete information. This is a normal condition in continuous-delivery environments where business requirements, user behavior, and technology options are

only partially known. When uncertainty is left implicit, teams either defer decisions indefinitely (accumulating technical debt) or make premature actions that may constrain the architecture evolution of their systems. Both outcomes can slow down delivery and degrade system quality over time, turning the software architecture into what Foote and Yoder (1999) famously characterized as a “Big Ball of Mud.”

Moreover, the concept of “fast flow”, popularized through Team Topologies (SKELTON and PAIS, 2019), emphasizes that organizations should design both their teams and their architecture to enable a continuous, rapid stream of software changes. Achieving fast flow presupposes that teams can make local architectural decisions without waiting for a central architect to approve them. This in turn demands that architectural uncertainties be made explicit and that lightweight mechanisms exist for teams to assess, prioritize, and revisit them as evidence accumulates. This is where Hypotheses Engineering may be used to play a key role in the architecture work, as proposed by SILVA, MELEGATI, SILVEIRA, *et al.* (2025).

2.2 Hypotheses Engineering and ArchHypo

Hypotheses engineering is a general approach to managing uncertainty in software development, first articulated by MELEGATI, WANG, and ABRAHAMSSON (2019). Its core idea is to treat uncertain decisions as testable hypotheses rather than settled facts: instead of asserting that a particular choice *will* produce a desired outcome, teams state that it *might* do so and then design experiments to gather evidence that validates or invalidates the hypothesis. Although it originated in the broader context of software startups and product development under uncertainty, the same discipline of making assumptions explicit and evidence-driven can be applied wherever teams must act despite incomplete information.

The ArchHypo technique, proposed by SILVA, MELEGATI, SILVEIRA, *et al.* (2025), is the architecture-specific adaptation of hypotheses engineering: it operationalizes this idea for software architecture, where a particular technology or pattern is treated as a hypothesis about whether it will satisfy a quality requirement rather than an assumed fact. Grounded in Design Science Research (DSR) and evolved from a body of patterns for agile architecture first disseminated at conferences and through corporate training (SILVA, ADOLFO, *et al.*, 2024; SILVA, MELEGATI, WANG, *et al.*, 2024), ArchHypo provides a structured vocabulary and a set of practices that fit within iterative development workflows where architecture uncertainty is a constant challenge for rapidly-evolving software systems.

2.2.1 Core Concepts

An *architectural hypothesis* is a *falsifiable* statement that captures an uncertainty relevant to the design of a software architecture. In this context, *falsifiable* means that the possibility of proving the statement false exists. This means that the hypothesis is refutable and testable through evidence, and *not* that the hypothesis must be phrased as something the team expects to be false. ArchHypo prescribes that hypotheses be written as simple, direct statements that express the team’s current beliefs, formulated freely to reduce the friction of adopting the technique (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). The practice descends from Hypotheses Engineering, first articulated by MELEGATI, WANG, and ABRAHAMSSON (2019); interrogative or questioning templates such as QUEST (MELEGATI

and WANG, 2019) are an optional, compatible way to represent hypotheses rather than a requirement of the technique. In this sense, an architectural hypothesis is fundamentally *a way of recording and registering an item of architectural uncertainty*: each hypothesis pairs one uncertainty with an explicit, testable statement about it, turning otherwise implicit assumptions into first-class artifacts that can be assessed, tracked, and revisited as evidence accumulates (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). It is this hypothesis and uncertainty pairing that distinguishes the concept from adjacent concepts:

- **Hypothesis versus Risk:** A risk involves a known event with estimable probability and impact. An architectural hypothesis, by contrast, stems from a *lack of information* sufficient to make a decision. The distinguishing feature is the epistemic gap, not the probability of an outcome as observed in SILVA, MELEGATI, SILVEIRA, *et al.* (2025).
- **Hypothesis versus Requirement:** A requirement specifies what the system must do or how well it must do it. A hypothesis may be triggered by a requirement (e.g., a quality-attribute goal that cannot yet be confirmed as achievable), but the hypothesis captures the uncertainty about the *path* to satisfying that requirement, not the requirement itself.
- **Hypothesis versus Architectural Decision.** An architectural decision resolves a design choice. A hypothesis *precedes and informs* decisions: it identifies what must be learned or validated before a decision can responsibly be made ((SILVA, MELEGATI, SILVEIRA, *et al.*, 2025)).

Concept	Definition	Relation to Architectural Hypothesis
Risk	A known event with estimable probability and impact.	A hypothesis stems from an <i>epistemic gap</i> — lack of information — not from a known probability.
Requirement	Specifies what the system must do or how well it must do it.	A requirement may <i>trigger</i> a hypothesis when the path to satisfying it is uncertain; they are not the same artifact.
Architectural Decision	Resolves a design choice.	A hypothesis <i>precedes and informs</i> a decision; it captures what must be learned before a decision can responsibly be made.

Table 2.1: Distinguishing architectural hypotheses from adjacent concepts (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).

Hypotheses arise from two principal sources:

- (a) *requirements*: when quality attribute targets are poorly specified or the team lacks information about whether a given approach can satisfy them; and

- (b) *solutions*: when the consequences of an existing or candidate architectural solution are unknown.

2.2.2 Assessment

Each hypothesis is evaluated along two independent dimensions:

1. **Uncertainty level.** How far the team is from being able to refute or accept the hypothesis. Crucially, uncertainty level is *not* equivalent to probability. A hypothesis that is almost certainly true and a hypothesis that is almost certainly false both represent *low* uncertainty, because in either case the team already has enough information to act. High uncertainty means the team genuinely does not know which way the evidence will fall.
2. **Impact level.** The potential consequences — in terms of rework, schedule, cost, or quality degradation — if the hypothesis is invalidated. A high uncertainty combined with a high impact signals that a hypothesis demands urgent attention.

Both dimensions can be assessed using a 3-point or 5-point qualitative scale; assessments are made independently and relatively, not on absolute probability as reported in [SILVA, MELEGATI, SILVEIRA, *et al.* \(2025\)](#).

2.2.3 Technical Plans

For hypotheses with significant uncertainty or impact, teams formulate a *technical plan*: a set of concrete actions whose goal is to (a) accept or refute the hypothesis, (b) reduce its uncertainty level, (c) reduce its potential impact, or (d) define criteria under which the decision can responsibly be deferred. ArchHypo identifies the following plan strategies as observed in [SILVA, MELEGATI, WANG, *et al.* \(2024\)](#):

- **Architectural Spike.** A time-boxed, isolated experiment to explore a technology or design option without coupling the result to the production codebase. The spike produces information rather than deliverable functionality.
- **Tracer Bullet.** A thin, end-to-end implementation within the actual system, exercising all architectural layers to validate an integration or performance characteristic under realistic conditions.
- **Software Analytics.** Using monitoring, profiling, or usage data already collected from the running system to gather evidence about a hypothesis without additional development effort.
- **Modularity / Flexibility.** Restructuring the affected subsystem to reduce coupling or increase replaceability, thereby lowering the *impact* of the hypothesis rather than resolving its uncertainty. This strategy is useful when full resolution is premature but protection against a bad outcome is warranted.
- **Guideline Implementation.** Tailoring the development process—for example, adopting a coding standard, a review protocol, or a testing policy—so that the team systematically generates evidence about the hypothesis as a by-product of normal

work. This was the most frequently used strategy in the evaluation reported by [SILVA, MELEGATI, SILVEIRA, et al. \(2025\)](#).

- **Architectural Trigger.** A condition-based deferral in which the team agrees to resolve the hypothesis only when a specific observable event occurs (e.g., when the number of concurrent users exceeds a threshold, or when a performance regression appears in monitoring). Until the trigger fires, the decision is deliberately deferred rather than made prematurely.
- **Agile Landing Zones.** Adopting architectural patterns that explicitly preserve flexibility, creating “landing zones” into which future decisions can be inserted without major disruption, regardless of which direction the evidence eventually points.

Strategy	Description	Primary Effect
Architectural Spike	Time-boxed isolated experiment; produces information, not deliverable functionality.	Reduces uncertainty.
Tracer Bullet	Thin end-to-end implementation in the real system to validate integration or performance.	Reduces uncertainty under realistic conditions.
Software Analytics	Uses existing monitoring or profiling data to gather evidence.	Reduces uncertainty without extra development effort.
Modularity / Flexibility	Restructures the subsystem to reduce coupling or increase replaceability.	Reduces impact (not uncertainty).
Guideline Implementation	Tailors the development process (coding standards, review protocols) to generate evidence as a by-product of normal work.	Reduces uncertainty incrementally; most frequently used (SILVA, MELEGATI, SILVEIRA, et al., 2025).
Architectural Trigger	Defers action until a specific observable condition is met.	Manages timing of decision; reduces premature cost.
Agile Landing Zones	Adopts patterns that preserve flexibility for future insertions.	Reduces impact proactively.

Table 2.2: ArchHypo technical plan strategies and their primary effect ([SILVA, MELEGATI, WANG, et al., 2024](#); [SILVA, MELEGATI, SILVEIRA, et al., 2025](#)).

2.2.4 Lifecycle

Hypotheses progress through a defined lifecycle: they begin in an *open* state and transition to either *validated* or *invalidated* as evidence from technical plans accumulates. The *responsible moment* is the key decision point in the hypothesis lifecycle: it is the point at which the team has accumulated enough evidence to act on the architectural decision — neither so early that the decision would be premature, nor so late that deferral would create unnecessary risk.

2.2.5 Empirical Evaluations

Start-up evaluation (Catch Solve). The first evaluation of ArchHypo in a real operating environment was reported in an IEEE Software article in [SILVA, MELEGATI, WANG, et al. \(2024\)](#). The setting was Catch Solve, a startup offering web testing and monitoring services based in Italy’s South Tyrol region. Four architectural hypotheses were formulated and tracked over approximately ten months, covering quality attributes of availability, reusability of test templates, usability, and scalability of virtual machines. For example, the hypothesis “*The lack of redundancy can cause problems with application availability for the customers*” was assessed with a low impact and very low uncertainty, which led the team to defer investment in redundant infrastructure and instead set up an Architectural Trigger based on the number of new customers. The study showed that naming and tracking uncertainties explicitly helped the technical lead decide when to invest in architectural improvements versus when to wait, and validated the overall structure of the ArchHypo approach. A key insight from this evaluation was that explicit hypothesis management changed how the team prioritized technical work: rather than treating architectural uncertainty as urgent and tacit, it became visible, assessed, and deliberately managed within a defined lifecycle.

Mission-critical evaluation (PropSEP at Jetsoft). The larger and more controlled evaluation, published in the IEEE Transactions on Software Engineering ([SILVA, MELEGATI, SILVEIRA, et al., 2025](#)), examined the application of ArchHypo within Jetsoft, a Brazilian software development company, on the PropSEP project—the budget proposal system of the Planning and Management Secretariat of the State of São Paulo. PropSEP is mission-critical: the approved budget it processes amounts to approximately 317 billion BRL ([SILVA, MELEGATI, SILVEIRA, et al., 2025](#)). The evaluation ran across six sprints of seven business days each (May–August 2022) and involved eight participants (P1–P8), all holding computer science degrees with certifications in Oracle, WebSphere, or GeneXus. Ten architectural hypotheses were identified, spanning six quality attributes (performance, security, reliability, flexibility, usability, and team productivity), and seven distinct types of architectural practice were employed in the resulting technical plans ([SILVA, MELEGATI, SILVEIRA, et al., 2025](#)).

The study produced nine key findings:

1. Some hypotheses carried uncertainty that was monitored throughout the entire project duration without reaching full resolution, demonstrating that the technique accommodates long-horizon uncertainties ([SILVA, MELEGATI, SILVEIRA, et al., 2025](#)).

2. ArchHypo enabled the team to avoid upfront architectural design, distributing structural decisions across iterations as evidence became available (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
3. Participants reported improvements to the development process in terms of safety, transparency, and predictability (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
4. Hypotheses spanning different quality attributes were handled using qualitatively different technical strategies, illustrating the creativity the technique affords (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
5. Guideline Implementation (process tailoring) emerged as the most frequently employed strategy across all technical plans (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
6. The team collectively recognized that managing hypotheses improved the quality of their architectural decision-making (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
7. Post-release evidence pointed to a positive impact on product quality: no issues related to the managed hypotheses were detected after deployment (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
8. **Changes in the team's way of working constituted a significant adoption obstacle.** Mixing architectural and hypothesis-related tasks with functional development created estimation difficulties, and participants described the additional work as not feeling directly connected to product delivery (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).
9. **All eight participants unanimously considered the technique hard to learn.** Specific difficulties included: mapping architectural concerns to the hypothesis format, formulating appropriate technical plan actions, articulating the assumptions behind a hypothesis, and understanding hypotheses as the primary driver of architectural action. No participant felt prepared to apply ArchHypo without continued expert support after the initial workshop (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). One participant explicitly stated that “*more training would make the team less dependent on the team of architects*” (P3) (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). The authors themselves acknowledge this structural dependency in their discussion. The absence of a supporting tool that encodes this knowledge into the workflow risks ArchHypo to remain a technique that only functions when an experienced architect is actively steering its application.

Findings 8 and 9 directly motivate the tooling contribution of this work: a tool that embeds ArchHypo scaffolding in the daily engineering workflow, reducing the cognitive overhead of manual application and knowledge transfer.

2.3 Internal Developer Portals and Backstage

The adoption of microservice architectures and the pursuit of team autonomy through practices such as Team Topologies have created significant complexity in the daily workflow of software engineering teams. In response, the industry has converged around the

Study	Setting	Scope	Key Finding	Limitation
(SILVA, MELEGATI, WANG, et al., 2024)	Catch Solve startup (Italy)	4 hypotheses, 10 months	Explicit hypothesis tracking changed how team prioritized architectural investments.	Single practitioner; no tool support.
(SILVA, MELEGATI, SILVEIRA, et al., 2025)	PropSEP / Jetsoft (Brazil)	10 hypotheses, 6 sprints, 8 participants	ArchHypo distributed architectural work across iterations; all participants found technique hard to learn.	Manual application; expert steering required; no supporting tool.
This work	Large-scale tech organization (Latin America)	Multiple teams, scaled product	To be reported — first evaluation in a tool-mediated setting via Hypo-Stage.	Single organizational context; Backstage dependency.

Table 2.3: Summary of prior empirical evaluations of ArchHypo.

concept of *Internal Developer Portals* (IDPs) as a platform-engineering discipline aimed at reducing this complexity while preserving team independence ([ANJUM, 2026](#)).

2.3.1 The Role of Internal Developer Portals (IDPs)

An Internal Developer Portal serves as the primary interface between developers and the underlying platform infrastructure. Unlike a simple wiki or a project-management tool, an IDP is an operational hub that aggregates tools, services, documentation, and data into a unified “pane of glass” ([ANJUM, 2026](#)). Its primary goal is to reduce *cognitive load*: by abstracting the complexity of infrastructure and standardizing workflows, IDPs enable what practitioners call *Golden Paths* – recommended, supported, and automated paths for building and deploying software ([ANJUM, 2026](#)).

This goal aligns directly with the theory of autonomous stream-aligned teams summarized in Team Topologies guidance ([SKELTON and PAIS, 2019](#)): if teams are to own full slices of a product without depending on a central gatekeeper, the platform must lower the barriers to acting on that ownership. An IDP operationalizes this by ensuring that the information and tooling a team needs—service catalog, documentation, CI/CD pipelines, monitoring dashboards, and domain-specific plugins—are accessible from a single environment, without context-switching.

Conceptually, IDPs sit at the intersection of two ideas that are highly relevant to this

work. First, they embody the principle of *decentralized architecture*: rather than requiring teams to escalate architectural concerns to a central architecture group, a well-designed IDP makes it possible to encode architectural guidance, constraints, and tools directly into the developer workflow. Second, they are the natural delivery vehicle for hypothesis-engineering support: because ArchHypo must become part of the daily engineering cadence – not a separate off-sprint activity – embedding it in an IDP portal where developers already spend their working time is strategically motivating.

2.3.2 Backstage: Open Platform for Internal Developer Portals

Backstage¹ is an open-source framework for building Internal Developer Portals, originally developed at Spotify to manage the complexity arising from rapid growth and an expanding microservice estate. Spotify’s internal version of Backstage serves as the central nervous system for its infrastructure, used by over 2,700 engineers to manage more than 14,000 software components (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024). Spotify open-sourced the framework in March 2020, and it was subsequently donated to the Cloud Native Computing Foundation (CNCF) as an Incubating project (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024). As of 2025, more than 3,000 companies have adopted Backstage for their own portals (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024).

Backstage is not a rigid application but rather a modular ecosystem driven by configuration. Its architecture rests on three layers: (1) the *Core Framework*, which manages essential utilities including the plugin loader, routing, authentication, and error handling; (2) the *Plugin Ecosystem*, built around a split-plugin model where frontend plugins handle views and UI components while backend plugins manage API services, data processing, and storage; and (3) *External Integrations*, through which backend plugins serve as gateways to external systems and APIs, allowing the portal to aggregate data from cloud providers, CI/CD tools, and other infrastructure (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024).

Crucially, the plugin architecture allows organizations to extend the portal’s functionality by integrating custom tools directly into the developer’s daily workflow, without forcing context-switching. It is precisely this extensibility that makes Backstage the appropriate delivery vehicle for the Hypo-Stage tool: a plugin can introduce a new class of first-class object – the architectural hypothesis – into the same environment where a developer already inspects service health, reviews documentation, and triggers CI/CD runs (SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024).

CORRÊA (2025) demonstrated the technical feasibility of building such a plugin and validated the integration approach. That work established the initial version of the Hypo-Stage plugin for Backstage, which constitutes the artifact under empirical evaluation in this work.

¹ <https://github.com/backstage/backstage>

Layer	Responsibilities	Relevance to Hypo-Stage
Core Framework	Plugin loader, routing, authentication, error handling.	Provides the lifecycle hooks that Hypo-Stage’s frontend and backend plugins register into.
Plugin Ecosystem	Frontend plugins (views, UI); backend plugins (API services, data, storage).	Hypo-Stage is implemented as a split plugin: the frontend renders hypotheses and assessments; the backend persists data and exposes the API.
External Integrations	Backend plugins as gateways to external systems (cloud, CI/CD, APIs).	Hypo-Stage integrates with the Backstage catalog to auto-associate hypotheses with software entities.

Table 2.4: Backstage architecture layers and their role (*SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024*).

2.4 Hypo-Stage: The Initial Tool

Hypo-Stage is an open-source Backstage plugin developed as an undergraduate capstone project by Pedro Henrique Mariano Corrêa, supervised by the author of this work and co-supervised by Prof. Paulo Meirelles (CORRÊA, 2025). Its purpose is to operationalize ArchHypo within the Spotify’s IDP Backstage, embedding hypothesis management into the engineering platform workflow. Figure 2.1 shows the main hypotheses dashboard of that initial version.

The initial version of Hypo-Stage, corresponding to the code up to commit `cee509776508081ebdd473c2c4f710b8ef55947` in the `hypo-stage` repository,² provides the following capabilities:

- **Hypothesis management:** creating, editing, listing, and deleting architectural hypotheses, each linked to a Backstage catalog entity (e.g., a service, component, or system). Integration with the Backstage catalog means that hypotheses are automatically associated with the software entities they affect, without requiring separate registration.
- **Uncertainty and impact assessment:** rating each hypothesis on 5-point Likert scales for both uncertainty level and impact level, with free-text rationale fields. The use of a 5-point scale rather than a binary or 3-point scale provides finer-grained prioritization while remaining tractable for teams without formal risk management backgrounds.

² Source repository: <https://github.com/ArchHypo/hypo-stage>. The baseline revision for the initial tool design referenced throughout this thesis is commit `cee509776508081ebdd473c2c4f710b8ef55947`.




Figure 2.1: Initial Hypo-Stage hypotheses dashboard (Backstage plugin), corresponding to the *CORRÊA 2025* baseline at commit *ceee509*. The listing exposes each hypothesis with uncertainty, impact, status, source type, and creation date.

- **Technical planning:** attaching one or more technical plan items to a hypothesis, specifying action type (spike, tracer bullet, trigger, etc.), description, target date, expected outcome, and documentation links.
- **Visualization:** a chart plotting hypotheses on an uncertainty-vs-impact plane to support prioritization, making the relative urgency of hypotheses visible at a glance.
- **Status lifecycle:** tracking each hypothesis through the defined lifecycle states (open → validated or invalidated), preserving the history of status transitions to support retrospective analysis.

By integrating with Backstage, Hypo-Stage leverages the portal’s service catalog so that hypotheses are automatically associated with the software entities they affect. This design reduces context-switching: developers manage architectural hypotheses in the same environment where they consult documentation, inspect service health, and access CI/CD pipelines.

Known limitations: The initial tool design assumed familiarity with the ArchHypo vocabulary, which poses a barrier for teams new to hypothesis engineering. The case study (*CORRÊA, 2025*) reconstructed the Catch Solve startup scenario digitally but did not involve real-time use by a team. Furthermore, the tool lacked features for collaborative assessment (e.g., multiple raters per hypothesis) and for tracking the evolution of assessments over time. These gaps also motivated the empirical evaluation and refinement effort promoted in this work.

Capability	Description	Known Limitation
Hypothesis management	Create, edit, list, delete hypotheses linked to Backstage catalog entities.	Assumed prior ArchHypo vocabulary familiarity; no onboarding guidance.
Uncertainty & impact assessment	5-point Likert scales per dimension with free-text rationale.	No collaborative multi-rater assessment; evolution of assessments not tracked over time.
Technical planning	Attach plan items with action type, description, target date, and outcome.	Action types not enforced or explained in-tool; required expert facilitation to use correctly.
Visualization	Uncertainty-vs-impact scatter plot for prioritization.	Static; not updated reactively as assessments changed.
Status lifecycle	Tracks hypotheses through open → validated / invalidated states.	No history of intermediate state transitions preserved.

Table 2.5: *Hypo-Stage initial capabilities and known limitations identified prior to this study (CORRÊA, 2025).*

2.5 Deriving Architecture from Requirements: The Tooling Gap

The systematic mapping study (SOUZA *et al.*, 2019) examined methods for deriving architectural models from requirements specifications, analyzing 39 high-quality primary studies selected from a broader literature search.

Among their key findings, these were directly relevant to this work:

1. **Dependence on tacit knowledge.** Every reviewed method relied, to some degree, on the architect’s experience and intuition. No method provided a fully automated or guided derivation path accessible to less-experienced practitioners (SOUZA *et al.*, 2019).
2. **Limited support for less-experienced practitioners.** Methods assumed a skilled architect capable of navigating trade-offs among quality attributes. Teams without senior architectural expertise had little actionable guidance (SOUZA *et al.*, 2019).
3. **Weak or absent tool support.** Although some tools were identified—most notably QuaDAI and Monarch—they addressed only fragments of the derivation process and offered only partial coverage of the architectural reasoning cycle. The mapping found no comprehensive tool integrating hypothesis formulation, assessment, and tracking within a team’s daily workflow (SOUZA *et al.*, 2019).

4. **Insufficient evaluation.** Many proposed methods lacked rigorous empirical validation, relying on toy examples or author-run case studies rather than independent evaluations with real teams in production environments (SOUZA *et al.*, 2019).

The mapping study also reported that 76.9% of the 39 primary studies addressed functional requirements, with a significant proportion also treating non-functional requirements; yet the methods for handling quality attributes — precisely the domain where architectural uncertainty is most consequential — were among the least well supported by tools or empirical evidence (SOUZA *et al.*, 2019).

These findings characterize the state of architectural derivation methods at the time of the mapping: reliant on tacit expertise, weakly tooled, and thinly evaluated in production settings.

2.6 Summary: State of the Literature for This Work

Taken together, the background surveyed in this chapter frame architectural work as an exercise in managing uncertainty. Hypotheses Engineering (MELEGATI, WANG, and ABRAHAMSSON, 2019) and its architecture-specific realization, ArchHypo (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025), provide the conceptual basis: each architectural hypothesis records one item of uncertainty as an explicit, testable statement, which is then assessed for epistemic gap and priority and worked through a technical plan until enough evidence accumulates to resolve it. This turns otherwise implicit assumptions into first-class artifacts that a team can track and revisit as a project evolves.

However, the empirical picture exposes a persistent limitation. The evaluations of ArchHypo reported to date (SILVA, MELEGATI, WANG, *et al.*, 2024; SILVA, MELEGATI, SILVEIRA, *et al.*, 2025) showed real value in distributing architectural reasoning across iterations, but they were conducted through manual application, depended on expert team members, and offered no supporting tooling, leading participants to find the technique hard to learn and its practice disconnected from daily delivery. Additionally, the concept of Internal Developer Portals, and Backstage in particular (ANJUM, 2026; SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION, 2024), suggested a single platform for engineering workflows through which the techniques, such as ArchHypo, could be operationalized. The Hypo-Stage prototype (CORRÊA, 2025) took that first step, but it remained a feasibility demonstration rather than a workflow-ready tool.

The systematic mapping study (SOUZA *et al.*, 2019) also confirms a broader “tooling void” in which methods for deriving and managing architecture remain disconnected from the daily engineering workflow, weakly supported by tools, and rarely validated with real teams in production.

Grounded in these findings, this work takes Hypo-Stage as a starting artifact and asks how real engineering teams appropriate ArchHypo in a production setting. Chapter 3 connects this agenda to the adjacent research on architectural decision-making, tool support, and hypotheses-driven practice.

Chapter 3

Literature Context and Positioning

This chapter positions the present work within the existing literature rather than performing a classic approach-versus-approach comparison of competing techniques. Because the contribution of this thesis is the empirical, tool-mediated evaluation of a single well-defined technique rather than a new method competing with alternatives, the surrounding literature is best treated as context that situates and motivates the study. A feature-level comparison of existing tools against Hypo-Stage is consolidated in Table 3.1 in Section 3.3.

This chapter synthesizes the literature relevant to the empirical evaluation of Hypo-Stage. Rather than presenting a standalone systematic review, it exposes the author's earlier exploratory work, which included a tertiary systematic mapping study on software architecture for startups and a multivocal literature review on centralization and decentralization of architectural decision-making. The full tertiary mapping study and multivocal review are intentionally not reproduced as standalone chapters. This is a deliberate scope decision, so this work is not broadened beyond its focus on the tool-mediated evaluation of ArchHypo. Therefore, this literature context is summarized here as it motivated and framed that evaluation, keeping the work centered on how software engineering teams appropriate and refine Hypo-Stage in practice.

This chapter also incorporates references from the gray literature. Such sources are explicitly identified as relevant, despite having weaker evidential weight than peer-reviewed publications, because they often capture industry trends and real organizational experience more rapidly than academic publication cycles allow.

3.1 Centralization and Decentralization of Architectural Decisions

A recurring theme in both academic and gray literature is the tension between centralizing architectural authority in a few experienced individuals versus distributing it across development teams. Organizations pursuing fast flow and team autonomy increasingly favor decentralized models as observed by SALAMEH and J. M. BASS (2021) and OLIVEIRA DE

DIANA *et al.* (2011), yet several studies report that without explicit governance mechanisms, decentralization leads to inconsistent decisions and architectural drift as seen in the work of DA SILVEIRA NETO *et al.* (2025), DASANAYAKE *et al.* (2016), and OLIVEIRA DE DIANA *et al.* (2011).

The author’s preliminary exploratory work examined this tension through two complementary lenses. Firstly, the author conducted a tertiary systematic mapping study surveying the existing landscape of secondary studies on software architecture for startup and scale-up organizations, identifying recurring themes in how early-stage teams approach architecture governance under resource and time constraints. For the second lens, the author performed a multivocal literature review that synthesized both academic publications and gray literature (practitioner blogs, conference talks, and industry reports), covering three interrelated perspectives: deriving software architecture from requirements, the emerging socio-technical architecture trend (which situates architectural decisions within organizational and team structures), and the tension between centralized and decentralized decision-making.

The multivocal review surfaced a framing where the industry advocates strongly for decentralized decisions to enable team autonomy and fast flow (SKELTON and PAIS, 2019), yet the methods and tools available for architectural decision-making depend heavily on the tacit knowledge of experienced architects (SOUZA *et al.*, 2019). Decentralization thus risks transferring decision responsibility to practitioners who lack the conceptual vocabulary or methodological support to exercise it effectively. However, the author recognized that this tension reflects a recognized design trade-off rather than a genuine logical contradiction, and framing this work around resolving that tension would risk over-extending its scope. For this reason, this work consequently does not employ that framing; instead, it focuses on the concrete, tractable problem of providing tool support for one well-defined uncertainty management technique.

3.2 Socio-Technical Architecture

A significant body of work, both academic and practitioner-oriented, frames architectural decision-making not merely as a technical concern but as a *socio-technical* one — shaped by the organizational structure, communication patterns, and cognitive constraints of the teams involved, as observed in TANG *et al.* (2017), RAZAVIAN *et al.* (2019), and HERBSLEB and GRINTER (1999). This framing is directly relevant to the motivation for the Hypo-Stage tool, because tool-embedded scaffolding does not only lower the technical cost of applying the ArchHypo technique, but it also contributes to changing the social dynamics of how architectural uncertainty is surfaced and addressed within a team, as studied by BOROWA *et al.* (2023) and RAZAVIAN *et al.* (2019).

3.2.1 Conway’s Law and Its Implications

The foundational observation underlying the socio-technical framing is Conway’s Law (CONWAY, 1968): “Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization’s communication structure.” Conway’s Law implies that architectural decisions cannot be analyzed in isolation from the

organizational and team structure in which they are made. Considering that perspective, this work also supports the idea that team boundaries become architectural boundaries.

SKELTON and PAIS (2019) extended this insight into an actionable framework with the concept of the *Inverse Conway Maneuver*: rather than allowing the organization’s structure to passively determine the architecture, teams should deliberately design team topologies to produce the target architecture. In practical terms, this means that architectural governance — including how uncertainty is managed and who is empowered to make decisions — must be considered a design decision with organizational as well as technical consequences.

3.2.2 Socio-Technical Architecture as an Industry Trend

The 2024 InfoQ Software Architecture and Design Trends Report (BETTS *et al.*, 2024) identified *socio-technical architecture* as a distinct “Early Adopter” trend, explicitly describing it as replacing the earlier framing of “architecture as a team sport.” The report notes that the trend encompasses two related ideas: on one side, the evolving role of architects as mentors and technical leaders rather than sole decision-makers; on the other, the growing recognition — driven in part by the dissemination of Team Topologies (SKELTON and PAIS, 2019) — that the design of a system must consider all the people who build, support, and maintain it (BETTS *et al.*, 2024).

It is important to note that InfoQ is a practitioner-oriented publication, not a peer-reviewed academic venue. The Trends Report reflects the editorial judgment of an experienced editorial board drawing on community submissions, conference talks, and practitioner case studies. As gray literature, its findings carry weaker evidential weight than systematic empirical studies; nonetheless, they are valuable as an indicator of what concepts are gaining traction in industrial practice and what challenges practitioners are actively facing, as observed in BETTS *et al.* (2024).

In a related InfoQ article, HUNTER and STOJMILOVA (2025) described a transformation journey at a software organization adopting Team Topologies and decentralized architectural decision-making. Their account illustrates several socio-technical dynamics that are directly relevant to this work: the initial disbelief of teams that they are actually empowered to make architectural decisions, the flood of decisions that follows once empowerment is accepted, the consequences teams face when decisions have unforeseen cross-boundary impacts, and the eventual self-correction as teams develop peer accountability mechanisms. All four phases are driven by social factors at least as much as technical ones.

A further gray-literature source relevant to this work is an InfoQ article that argues that “distributed architecture is more social than technical” (TIBBITTS, 2025). Tibbitts draws on Andrew Harmel-Law’s concept of the Architecture Advice Process — in which any team can make any architectural decision as long as it has looked for advice from the affected parties and documented the decision in an ADR. This approach is presented as a mechanism for replacing permission-seeking with advice-seeking. The social shift from “who approves this?” to “who should I talk to before we do this?” is presented as a cultural unlock that enables decentralized governance without sacrificing coherence.

3.3 Tool Support for Architectural Decision-Making

Along with the author’s exploratory research, several tools and platforms have been proposed to support architectural decision-making:

- **ADR tools** (e.g., adr-tools, Log4brains), documented in practitioner and gray-literature sources (NYGARD, 2011; PRYCE, 2024; VAILLANT, 2024), focus on recording and versioning Architecture Decision Records (ADRs). They capture the *what* and *why* of decisions but do not model uncertainty, impact, or technical plans for hypotheses that are not yet decisions. ADRs document settled choices; they are not designed to manage the pre-decision state of uncertainty that ArchHypo addresses.
- **Architecture analysis tools** (e.g., ATAM-based workshops, trade-off analysis methods) (KAZMAN *et al.*, 1998) provide structured evaluation of architectural alternatives but are typically heavyweight, workshop-based activities rather than continuous, team-embedded processes. Their value lies in one-shot or periodic analysis sessions, not in the ongoing, iterative management of uncertainty throughout a product’s lifecycle.
- **Internal Developer Portals** such as Backstage, Port, and Cortex offer a catalog-centric view of services and infrastructure. While they improve discoverability and documentation, none of these platforms natively supports hypothesis engineering or architectural uncertainty management. They provide the integration surface that Hypo-Stage exploits, but leave the hypothesis engineering layer entirely absent.

Hypo-Stage differentiates itself by being explicitly designed around the ArchHypo lifecycle: from hypothesis elicitation through assessment, technical planning, and resolution. Its integration with Backstage places this lifecycle within the developer’s daily tooling environment. No existing tool, to the best of the author’s knowledge, provides this combination.

3.4 Empirical Studies of Architectural Practices in Industry

Industry-focused empirical studies on architectural practices remain relatively scarce compared with tool proposals and method descriptions. Among the most relevant:

- The ArchHypo evaluations by SILVA, MELEGATI, SILVEIRA, *et al.* (2025) and SILVA, MELEGATI, WANG, *et al.* (2024) constitute the closest antecedents to this work. They demonstrated the value of hypothesis engineering in two distinct settings — a web-testing startup and a mission-critical government budget system — but relied entirely on manual application with no tool support. This work introduces the first evaluation of ArchHypo in a tool-mediated setting.
- Studies on architectural governance in agile organizations, including tailored Spotify-model variants (SALAMEH and J. M. BASS, 2021) and architecture groups in multi-team agile settings (OLIVEIRA DE DIANA *et al.*, 2011), document how companies struc-

Tool / Platform	Primary Purpose	Models Uncertainty	Technical Plans	Lifecycle Tracking	Team-Embedded
ADR tools (adr-tools, Log4brains)	Record settled decisions.	No	No	Partial	No
ATAM / Trade-off workshops	Evaluate architectural alternatives.	Partial	No	No	No
IDPs (Backstage, Port, Cortex)	Developer portal; catalog; documentation.	No	No	No	Yes
Hypo-Stage (this work)	Support ArchHypo hypothesis engineering.	Yes	Yes	Yes	Yes (via Backstage)

Table 3.1: Comparison of tools relevant to architectural decision-making support.

ture roles and ceremonies for architecture, but rarely examine explicit uncertainty-management techniques.

- The multivocal review conducted during the author’s exploratory research surfaced practitioner reports (blog posts, conference talks, white papers) describing ad-hoc approaches to architectural uncertainty. These reports corroborate the need for structured methods but, being gray literature, carry weaker evidential value.

Based on the findings of these related studies, this work contributes to the body of empirical evidence by evaluating a tool-supported application of ArchHypo with multiple teams, addressing a gap in both the ArchHypo literature (which has not yet studied tool-mediated use) and the broader architectural-practices literature (which lacks studies of hypothesis-engineering tools in scaled settings).

Chapter 4

Research Design

This chapter describes and justifies the research design adopted in this work. It covers the adopted research strategy along with the conceptual framework created to fulfill the research questions. Moreover, it covers the study context and case selection, the researcher’s role, the data collection procedures, and the criteria adopted to establish the trustworthiness of the findings. The chapter concludes with notes on the ethical considerations that governed the study.

4.1 Research Strategy

The objective of this work was to empirically evaluate and refine Hypo-Stage, a support tool for the ArchHypo technique, by applying it with multiple software engineering teams in a real-world organizational setting. Answering the research questions posed in Chapter 1 requires understanding not only *what* outcomes the tool produces, but also *how* teams interact with it, *why* certain difficulties arise, and *which* adjustments improve its fit with existing practices. Following the design framework for flexible qualitative inquiry (ROBSON and McCARTAN, 2016), the purpose of this work combined with its research questions clearly pointed out the need for a **flexible research design**.

Within flexible designs, the study follows a **multiple case study strategy** as defined in methodological treatments of case-study research (ROBSON and McCARTAN, 2016; YIN, 2009). A case study is defined by them as an empirical investigation of a particular contemporary phenomenon within its real-life context, using multiple sources of evidence and where the boundary between phenomenon and context is not clearly distinct. This definition fits the present work precisely: the phenomenon under study is the use of Hypo-Stage by a software engineering team, and its context — the organizational, technical, and social environment in which the team operates — is both relevant and inseparable from the phenomenon itself. A purely experimental design, which would require controlled manipulation of variables and prespecification of the design before data collection, is not appropriate here. The number of teams is small, the settings are naturalistic and non-manipulable, and the research purpose calls for rich, process-oriented accounts rather than statistical inference (ROBSON and McCARTAN, 2016).

The study involves **two cases**, each consisting of a distinct software engineering team. Using multiple cases rather than a single case follows the logic of *analytic (or theoretical) generalization* from prior methodological accounts (ROBSON and McCARTAN, 2016; YIN, 2009): the two cases are not selected to constitute a statistical sample of teams, but to provide complementary and contrasting instances through which patterns of tool use, difficulties, and refinements can be examined and cross-validated. Multiple case studies are analogous to multiple experiments where each case either replicates findings or provides theoretically predicted contrast — both strengthening the analytic basis for generalization (ROBSON and McCARTAN, 2016).

Beyond its descriptive and explanatory purposes, the study also has an **evaluative purpose**: it seeks to assess the worth and fitness of Hypo-Stage as a practical tool for managing architectural uncertainty in software engineering teams. Flexible-design accounts distinguish *formative evaluation*, which aims at understanding and improving an ongoing program or intervention, from *summative evaluation*, which renders a final verdict on overall effectiveness (ROBSON and McCARTAN, 2016). This study is explicitly formative: data collected during the observation period are used to identify difficulties and trigger tool refinements in near real-time, rather than to produce a retrospective pass-or-fail judgment.

Finally, the iterative loop through which the researcher observes difficulties, implements refinements to the tool, deploys an updated version, and resumes observation incorporates a core element of **action research** as described in the methodological literature (ROBSON and McCARTAN, 2016): a deliberate cycle of plan, act, observe, and reflect, with the explicit intent of producing practical change alongside research knowledge. This does not transform the study into a full action research project — the primary goal remains the production of knowledge about tool use and architectural uncertainty, not organizational change more broadly — but the iterative refinement component follows the same logic and must be acknowledged as part of the design.

In summary, the research strategy can be formally described as a **flexible, multiple case study design with formative evaluation purpose and action research elements**, following the taxonomy from ROBSON and McCARTAN (2016).

4.1.1 Key Methodological Concepts

Because the subsequent sections, and the results chapter, rely repeatedly on a compact set of methodological terms drawn from flexible-design guidance (ROBSON and McCARTAN, 2016), the definitions below clarify their meaning for the remainder of this work. They are collected here so that every later mention resolves to an unambiguous referent.

Flexible design: A research design in which the detailed plan (data sources, instruments, and categories of analysis) is allowed to evolve during the study in response to what the researcher learns in the field, in contrast to a *fixed design*, where all elements are pre-specified before data collection begins (ROBSON and McCARTAN, 2016). Flexible designs are the standard choice for case studies, ethnographies, and grounded-theory work, where the phenomenon under study is only partially understood at the outset and the researcher expects to refine the analytical lens as evidence accumulates.

Multiple case study: An empirical investigation that examines two or more contemporary phenomena within their real-life contexts, using multiple sources of evidence, where the boundary between phenomenon and context is not clearly distinct (ROBSON and McCARTAN, 2016; YIN, 2009). The cases are selected not as a statistical sample but as analytical units that either replicate findings (*literal replication*) or produce theoretically predicted contrasts (*theoretical replication*).

Formative evaluation: An evaluative purpose aimed at understanding and improving an ongoing program, intervention, or artifact while it is still being shaped, as distinguished from *summative evaluation*, which delivers a retrospective pass-or-fail verdict on overall effectiveness (ROBSON and McCARTAN, 2016). In the present study, the evaluative object is the Hypo-Stage tool, and the formative stance justifies the iterative refinement cycle described in Section 4.5.

Action research: A tradition of inquiry in which the researcher deliberately alternates between observing a setting, intervening in it, and reflecting on the effects of that intervention, following a recurrent *plan–act–observe–reflect* cycle (ROBSON and McCARTAN, 2016). The present study borrows the cycle as a design element without adopting the emancipatory or organizational-change commitments that define full action research.

Purposive sampling: A non-probabilistic sampling strategy in which participants or cases are chosen because they are expected to be informative about the phenomenon of interest, rather than because they are representative of a population (ROBSON and McCARTAN, 2016). The two teams in this study were selected purposively to provide complementary contrast (product vs. platform scope), not to constitute a statistical sample of engineering teams.

Analytic generalization: The form of generalization proper to case study research, in which findings are extended to theoretical propositions rather than to populations (ROBSON and McCARTAN, 2016; YIN, 2009). It is contrasted with *statistical generalization*, which requires a representative sample and probabilistic inference. The claims in Chapter 5 are analytic: they describe conditions under which the observed patterns are theoretically expected to recur, not the frequency with which they would appear in a wider population.

Member checking: A trustworthiness procedure in which interpretations derived by the researcher are returned to the participants whose behavior or statements produced them, so that the participants may verify, correct, or elaborate on the researcher’s reading (ROBSON and McCARTAN, 2016). Section 4.6.3 describes how member checking was operationalized in this study; Section 4.8 addresses it under credibility.

These concepts recur throughout Sections 4.3, 4.6, and 4.7, and frame the claims made in Chapters 5 and 6.

4.2 Conceptual Research Framework

The conceptual research framework organizes the key constructs of this work and their relationships in a way that is consistent with the research questions and guides the data collection and analysis as described for flexible designs (ROBSON and McCARTAN, 2016). Figure 4.1 describes the framework.




Figure 4.1: Conceptual framework of the study.

The central piece is the *Hypo-Stage Tool Usage*, understood as the practices through which team members appropriate and practice the ArchHypo technique via Hypo-Stage: creating and editing architectural hypotheses, assessing uncertainty levels and impact, defining and tracking technical plans, and revisiting hypotheses over successive sprint cycles (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). The tool usage is conditioned by two sets of contextual factors: the *team context* (composition, experience profile, scope of work, and familiarity with architectural concerns) and the *organizational context* (the company’s development platform, the role of Backstage as the delivery portal, and the scaled product being developed). The tool usage produces two observable streams: the *artifacts* recorded in Hypo-Stage (hypotheses, assessments, technical plans, and their evolution over time) and the *experiential observations* of team members (benefits, difficulties, and adaptation strategies). These streams feed the *Hypo-Stage Tool Refinements* — changes made to Hypo-Stage and its usage guidelines — which in turn modify subsequent tool usage, forming an iterative cycle. This framework is consistent with the formative evaluation and action research elements of the design, and directly links each element to the research questions (RQ1–RQ4) stated in Chapter 1.

4.3 Research Context

The study was conducted within a technology organization that operates large-scale software products across Latin America and employs over 500 software engineers. The organization is an intensive user of Backstage,¹ an open-source developer portal, which serves as the central platform for internal tooling, documentation, and engineering workflows. The availability of this platform was an enabling condition for the study: Hypo-Stage was deployed as a Backstage plugin, so its adoption is conditioned by the availability of that platform. The Hypo-Stage tool was deployed directly in the company's production environment, integrated with the same catalog and workflow interface that engineers use daily, reducing adoption friction considerably. This allowed the researcher to observe the tool usage in a naturalistic setting, with minimal interference from the researcher.

4.3.1 Case Selection

Two teams were selected as cases. Following methodological guidance on purposive sampling in flexible designs (ROBSON and McCARTAN, 2016), the cases were not chosen to be statistically representative of a population of teams, but to provide *complementary contrast*: teams with different compositions, scopes, and relationships to the broader system, so that differences and similarities in tool use can be analytically compared. There are exactly two cases in this study and each is treated as a single, holistic unit of analysis. The design does not embed sub-units (such as individual engineers) as separate units within a case; observations at the level of individuals are read as evidence about the team as a whole.

A second, pragmatic criterion also governed the selection: both teams were chosen because the researcher, as an employee of the organization, had insider access to them. This access was not merely an enabling condition but an explicit selection factor that made it possible to conduct close-range observation of architectural decision-making that this study requires. The accepted trade-off is that such access facilitates the research while introducing a risk of insider bias. Such risk is acknowledged and addressed through the trustworthiness measures described in Section 4.8.

Team 1 — Product Team. A feature-delivery (Product) team of 12 software engineers (5 senior, 7 mid-level) responsible for developing new product features with direct impact on the company's business objectives and end-user experience. This team works on a bounded product domain, with close proximity to business requirements and user feedback. Its architectural concerns are primarily about supporting new features reliably and at scale while preserving existing quality attributes.

Team 2 — Platform Team. An internal-tooling (Platform) team of 4 software engineers (2 senior, 2 mid-level) responsible for developing and maintaining broad-scope internal tooling that supports software development across the entire organization. This team's architectural decisions affect how all engineers in the company build, deploy, and operate software. Its architectural concerns are therefore more horizontal, with higher impact on organizational engineering practices and greater uncertainty about usage patterns and requirements.

¹ <https://github.com/backstage/backstage>

Literal and theoretical replication. In the multiple-case replication logic for case studies (YIN, 2009), two or more cases are deliberately selected to perform one of two complementary analytical roles. *Literal replication* chooses cases that share the relevant conditions and are therefore predicted to produce similar findings; convergence across literal replications strengthens the claim that the finding is genuinely tied to the intervention rather than to idiosyncratic case features. *Theoretical replication* chooses cases that differ on a theoretically meaningful dimension and are therefore predicted to produce contrasting findings; divergence along the predicted direction strengthens the theoretical account of *why* the intervention behaves differently across conditions.

Although the two teams of this study are exposed to the same intervention (use of Hypo-Stage and ArchHypo), their contrasting scopes and compositions are expected to produce theoretically predicted differences in how architectural uncertainty is experienced and in the challenges of adopting a hypothesis-driven approach. ROBSON and McCARTAN (2016) suggests that such divergence is analytically informative in its own right. Therefore, as this study’s design predicts divergence, any convergence observed across the two cases despite their configurational differences was still considered analytically informative for the research questions.

Dimension		Team 1 — Product Team	Team 2 — Platform Team
Size		12 engineers (5 senior, 7 mid-level).	4 engineers (2 senior, 2 mid-level).
Scope		Bounded product domain; new features with direct business and user impact.	Broad-scope internal tooling supporting all engineers in the organization.
Architectural concerns		Supporting new features reliably at scale; preserving existing quality attributes.	Horizontal decisions affecting build, deploy, and operate practices organization-wide.
Uncertainty (expected)	profile	Moderate uncertainty, high proximity to user feedback.	Higher uncertainty about usage patterns; high cross-team impact of decisions.

Replication logic (case pair): the two cases together form a *theoretical replication* — under the same intervention, their contrasting scopes and compositions are expected to surface different challenges. Replication is a property of the case pair, not of either team individually.

Table 4.1: Case comparison: Team 1 (Product) and Team 2 (Platform).

4.4 Researcher’s Role

Throughout the study, the researcher occupied the dual role of **participant observer and facilitator**, in the terms used in flexible-design methodology (ROBSON and McCARTAN, 2016). The researcher is an employee of the organization in which the study was conducted, which makes this an instance of *insider research*. This position provided privileged access to the teams, their work context, and their daily practices, which would have been substan-

tially harder to obtain as an external investigator. It also allowed the researcher to introduce Hypo-Stage into the company’s infrastructure in a credible and practically feasible way.

As discussed in accounts of insider research (ROBSON and McCARTAN, 2016), such studies raise specific concerns regarding bias and the potential conflation of the researcher’s own perspective with those of participants. These concerns are addressed in Section 4.8. It is important to note that, while the researcher served as facilitator during training sessions and provided assistance on request throughout the observation period, he did not direct the teams in *which* architectural uncertainties to address or *how* to formulate specific hypotheses. The content of the work produced by the teams in Hypo-Stage reflects their own judgments and priorities.

4.5 Procedure

The study procedure unfolded in three sequential phases. In **Phase 1 (Training and Introduction)**, each team was introduced to the ArchHypo technique and to Hypo-Stage through a dedicated training session. In **Phase 2 (Observation and Iterative Refinement)**, the teams used Hypo-Stage in their regular development work while the researcher observed tool use, supported the teams on request, and refined the tool in response to the difficulties that emerged — the plan–act–observe–reflect loop introduced in Section 4.1. Finally, in **Phase 3 (Active Listening and Wrap-up)**, the researcher held brief reflective sessions with team members to capture their overall experience. The remainder of this section describes each phase in turn; the specific instruments used to capture data during them are detailed in Section 4.6.

4.5.1 Phase 1 — Training and Introduction

Before the observation period began, each team participated in a one-hour training session led by the researcher. The session covered:

- (i) The conceptual foundations of hypotheses engineering and the ArchHypo technique, including the distinction between architectural hypotheses, uncertainty levels, impact assessment, and technical plans.
- (ii) The context and motivation of the research, including the broader challenge of making architectural uncertainties explicit in teams working on scaled products.
- (iii) A live demonstration of Hypo-Stage running in the company’s Backstage instance, illustrating how to create and manage hypotheses within the tool. Teams were given written communication prior to the meeting (Appendix A.3) — including an illustration of the ArchHypo big picture (Figure 4.2) — so that they could arrive at the session with some prior exposure to the key concepts.

4.5.2 Phase 2 — Observation and Iterative Refinement

Following the training session, each team entered an observation period during which they were invited to use Hypo-Stage to register and manage any architectural uncertainties that arose in the course of their regular development work. The term *observation* is used




Figure 4.2: Overview of the ArchHypo technique showing the iterative cycle through which engineers identify architectural hypotheses, assess their uncertainty and impact, define technical plans, and take or postpone architectural decisions. Source: adapted from SILVA, MELEGATI, SILVEIRA, *et al.* (2025).

here in the action-research sense of the plan–act–observe–reflect cycle (Section 4.1) rather than in the sense of passive watching: the researcher actively supported the teams and iteratively refined the tool throughout this phase, as detailed below. The organization operates on two-week agile sprints. The initial participant communication (Appendix A.3) described a four-week horizon; in practice, teams continued to use the tool voluntarily beyond that window.

The empirical window for hypothesis creation is anchored in Hypo-Stage registration timestamps. As documented in Appendix B, the thirteen hypotheses were created between 11 February 2026 and 12 April 2026: an interval of 61 calendar days (approximately nine weeks), spanning five consecutive two-week sprint periods from the earliest to the latest registration. No constraints were imposed on the type, number, or formulation of hypotheses, nor on when to create them during the sprints. It is intrinsic to the nature of the organization in which the experience was conducted that the teams are free to adjust their agile practices to their own needs. Therefore, in contrast to the PropSEP evaluation setting (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025), the teams were not required to create hypotheses during the sprints, but rather to use the tools whenever they deemed necessary. The objective was to observe tool use in naturalistic conditions, with minimal researcher interference in the content of the teams’ work.

During this period, the researcher remained available to both teams for support on request. Support was of two distinct types:

1. **Conceptual support.** When a team member found it difficult to formulate a hypothesis or to assess its uncertainty and impact within the ArchHypo framework, the researcher provided clarification or facilitated a brief discussion. Interactions of this type were noted as observations about the conceptual difficulty of the technique.

2. **Tool support.** When a team member encountered a difficulty that originated from the tool itself — a usability problem, a missing feature, or a mismatch between the tool’s behavior and the expected workflow — the researcher logged this as a tool improvement candidate and, as a matter of priority, implemented and deployed an updated version of Hypo-Stage as quickly as practicable. This was done to ensure that tool limitations would not persistently obstruct teams’ ability to practice ArchHypo, and to allow subsequent observations to reflect a more refined version of the tool.

This second type of support operationalizes the *action research cycle* within the study: difficulty observed → refinement designed → updated tool deployed → use and observation resumed. Refinements were deployed in near real time: whenever tool friction was logged, the researcher implemented and deployed an updated Hypo-Stage version as quickly as practicable, then resumed observation. The resulting changes span commit ceee509 (the baseline artifact boundary established by CORRÊA 2025) through commit 3a53d4e over the observation window; each refinement cycle is documented in Section 5.5. Empirical findings for RQ1–RQ3 reported in Chapter 5 describe that window cumulatively rather than the pre-refinement tool state alone; see Section 5.2.

4.5.3 Phase 3 — Active Listening and Wrap-up

At the end of the observation period, the procedure concluded with a wrap-up phase in which the researcher held open-ended active-listening sessions with the senior technical leader of each team, inviting them to reflect on their overall experience of using Hypo-Stage and ArchHypo. These sessions served as the primary reflective moment of the study; their format and the data they produced are described as a data collection instrument in Section 4.6.3.

4.6 Data Collection

Following the multiple data source approach recommended for case studies (ROBSON and MCCARTAN, 2016; YIN, 2009), data were collected through three complementary channels:

4.6.1 Participant Observation

The researcher’s continuous presence in the organization and involvement in both teams’ environments provided an ongoing stream of observational data. Observations covered:

- (i) Team members’ reactions and difficulties during and after the training session;
- (ii) When and by whom Hypo-Stage was opened, had artifacts created in it, or was left unused during the sprints;
- (iii) Interactions in which team members requested conceptual or tool support;
- (iv) Informal commentary about the fit of ArchHypo with existing practices, gathered during the observed team rituals—daily standups, sprint planning, backlog refine-

Instrument	Data Produced	Collection Timing	RQs Addressed
Participant observation (field notes)	Reactions, when and by whom Hypo-Stage was opened, created in, or left unused, informal commentary from team rituals, support requests.	Continuous during observation period.	RQ1, RQ3
Artifact analysis (Hypo-Stage records)	Hypotheses, assessments, technical plans, lifecycle transitions.	End of observation period.	RQ1, RQ2
Open-ended active listening sessions	Perceived usefulness, difficulties, suggestions, overall assessment.	End of observation period.	RQ3, RQ4

Table 4.2: Data collection instruments, sources, and research questions addressed.

ment, retrospectives, and dedicated architecture and technical discussions (e.g. the Kafka session)—as well as through informal hallway, Slack, and asynchronous interaction; these observations were subsequently corroborated via team feedback.

Following established guidance on participant observation in flexible designs (ROBSON and McCARTAN, 2016), field notes were recorded as close as possible in time to the observed events, and the researcher maintained a reflexive awareness of his dual role as insider and investigator.

4.6.2 Artifact Analysis

The hypotheses, assessments, and technical plans created by each team in Hypo-Stage constitute a documentary record of how they used the tool over the observation period. These artifacts were collected at the end of the observation period and analyzed to characterize:

- (i) The number and types of hypotheses created;
- (ii) The quality attributes and sources of uncertainty they addressed;
- (iii) How hypotheses were assessed and what technical plan actions were defined;
- (iv) How hypotheses evolved (created, refined, or closed) across the observation period (five two-week sprints spanned by registration timestamps; Appendix B).

The content analysis of documentary record is a recognized data collection method in flexible case study designs (ROBSON and McCARTAN, 2016). In particular, the objective attributes captured in items (i) and (ii) — counts of hypotheses, quality attributes, sources of uncertainty, and technical-plan strategies — are treated *quantitatively* through descriptive statistics, complementing the qualitative thematic coding applied to the remaining data (see Section 4.7).

4.6.3 Open-Ended Active Listening Sessions

In addition to informal interactions arising from support requests, the researcher conducted open-ended active-listening sessions with the senior technical leader of each team at the end of the observation period—one session per team, two in total—inviting them to reflect on their experience of using Hypo-Stage and ArchHypo. Each session lasted approximately 20–30 minutes, consistent with the duration communicated in the participant invitation (Appendix A.3). These conversations followed a semi-structured format guided by the thematic scaffold in Appendix A.1, covering the perceived usefulness of the technique, the difficulties encountered, the changes they would suggest to the tool or its integration into their workflow, and their overall assessment of the value of making architectural uncertainties explicit in this way. Semi-structured or focused interviews are a central method for flexible design research and are particularly appropriate when the researcher seeks to capture participants’ perspectives and interpretations rather than confirming predetermined hypotheses as described for flexible qualitative designs (ROBSON and McCARTAN, 2016).

During each session, the researcher also surfaced interpretations formed during the observation period—for instance, inferences about when and why hypotheses were registered—and invited the team leader to confirm, correct, or elaborate on those readings. This *member-checking* step strengthened the credibility of subsequent thematic analysis (Section 4.8).

The sessions were genuinely open-ended and were *not* audio-recorded or transcribed verbatim. The researcher reconstructed participants’ perspectives from dated field notes and post-session written recollection. This was a deliberate, low-friction choice appropriate to insider research in a live delivery setting. It avoided the formality and consent overhead of recording colleagues during ordinary work conversations, but it carries a methodological cost: the RQ3 account rests on the researcher’s own notes and memory recall rather than on a verbatim transcript that an independent auditor could re-inspect.

4.7 Data Analysis

Data from the three collection sources described in Section 4.6 were analyzed using *thematic coding analysis* as specified for qualitative data gathered under a flexible design (ROBSON and McCARTAN, 2016). Thematic coding is a family of qualitative analytic techniques that produces a structured account of what is present textual or observational data by iteratively assigning short labels — *codes* — to meaningful segments, aggregating those codes into higher-order *categories* or *themes*, and then comparing the resulting structure across units of analysis (ROBSON and McCARTAN, 2016). It is distinguished from fully inductive *grounded theory* by its willingness to accept a priori analytical constructs (in this study, those of the conceptual framework in Figure 4.1) and from fully deductive *content analysis* by its commitment to letting the codebook emerge from the data rather than being fixed in advance.

In line with established guidance on flexible designs (ROBSON and McCARTAN, 2016), the analysis proceeded through three stages, each answering a distinct analytical question.

Stage 1 – Inductive open coding: All field notes, extracted Hypo-Stage artifact records, and the reconstructed active-listening notes (Section 4.6.3) were reviewed line-by-line. A short descriptive *code* was assigned to each segment that carried analytical weight, capturing a recurring observation, a reported experience, or a notable silence. Codes at this stage are descriptive, not interpretive: the aim is faithful, low-inference labeling of what is present in the data. No a priori categories were imposed. Representative codes and their traceability to later sections appear in Table 5.1.

Stage 2 – Axial grouping into conceptual categories: The descriptive codes produced in Stage 1 were grouped into higher-order *axial categories* that correspond to the constructs of the conceptual framework. The term *axial* follows the Strauss–Corbin tradition adopted in accounts of thematic coding under flexible designs (ROBSON and McCARTAN, 2016): it denotes the relational re-organization of open codes around a central analytical axis, in this study the constructs of tool usage, architectural uncertainty, perceived benefits, perceived challenges, candidate refinements, and organizational-adoption conditions. Six axial categories emerged, each corresponding to one construct in Figure 4.1 or to an axis along which the framework’s observables were stratified.

Stage 3 – Cross-case comparison: The axial categories were then examined across the two cases (Product Team and Platform Team) to identify findings common to both teams and findings that were specific to one. Similarities are reported as shared observations in Chapter 5; differences are attributed to the distinct scopes and compositions of product versus platform work, following the theoretical replication logic set out in Section 4.3.1.

The complete mapping from representative open codes to axial categories, research questions, and chapter sections appears in Table 5.1 at the start of Chapter 5. That table is the primary audit trail linking raw observations to every empirical claim made in this thesis.

Consistent with the observation that flexible designs may include modest quantitative summaries without becoming multi-strategy designs (ROBSON and McCARTAN, 2016), the artifact data extracted from Hypo-Stage were also summarized descriptively (counts of hypotheses per team, frequency of quality attributes, technical-plan strategy distribution). These summaries complement and illustrate the qualitative findings; they do not constitute a separate quantitative strand of analysis.

4.8 Establishing Trustworthiness

The trustworthiness of findings from flexible design research is discussed in terms of *credibility*, *transferability*, *dependability*, and *confirmability* (ROBSON and McCARTAN, 2016). Each is addressed below.

Credibility — the internal validity of the findings — was supported through:

- (i) Prolonged engagement with both teams over approximately nine calendar weeks of real work in their natural environment (Hypo-Stage registrations span five consecutive two-week sprints; Appendix B);

Criterion	Meaning	Strategy Applied
Credibility	Internal validity; do findings accurately represent participants' realities?	Prolonged engagement (≈ 9 weeks, five two-week sprints covered by registrations); triangulation across three data sources; member checking.
Transferability	Degree to which findings may inform other settings.	Thick description of context; analytic (not statistical) generalization claimed.
Dependability	Reliability and consistency of the research process.	Audit trail: dated written field notes (not recordings); Hypo-Stage artifact exports at end of window; each refinement linked to commit <code>cee509-3a53d4e</code> .
Confirmability	Degree to which findings reflect data rather than researcher bias.	Role separation maintained via typed support logs; disconfirming evidence coded alongside difficulties; claims traceable to dated notes, artifacts, or reconstructions (Table 5.1).

Table 4.3: *Trustworthiness criteria and strategies applied in this study (ROBSON and McCARTAN, 2016).*

- (ii) Multiple data sources (observation, artifact analysis, active listening sessions), providing triangulation of evidence;
- (iii) **Member checking** (Section 4.6.3): during the end-of-period active-listening sessions, interpretations that the researcher formed during the observation period were surfaced back to each team's senior technical leader for confirmation, correction, or elaboration, operationalizing the member-checking procedure recommended as a core credibility control for flexible qualitative designs by ROBSON and McCARTAN (2016).

Transferability — the degree to which the findings may inform other settings — is constrained by the specifics of the organizational context (a large technology company in Latin America with an established Backstage platform) and the tool (Hypo-Stage). The study does not claim statistical generalization. Rather, it offers analytic generalization (ROBSON and McCARTAN, 2016; YIN, 2009): the patterns of use, difficulties, and refinements documented here provide a theoretically informed account that researchers and practitioners can assess for relevance to their own contexts.

Dependability — the reliability of the research process — was addressed by maintaining an audit trail of three linked record types, consistent with the transparency choices documented in Section 4.6.3. First, **written field notes** were kept in a private research journal as close as practicable to each observed event (Section 4.6.1); each entry was timestamped and tagged to a specific team, sprint, or ritual. Second, at the end of the observation

window, **artifact records were exported** from the Hypo-Stage database (hypothesis bodies, assessments, technical plans, and lifecycle timestamps) and archived alongside the notes for independent re-inspection (Appendix B). Third, each tool refinement deployed during the study was **linked to its triggering observation** and to the corresponding commit in the public Hypo-Stage repository, spanning revision ceee509 through 3a53d4e (Section 5.5). Together, these records constitute the dependability audit trail that Table 5.1 extends into the thematic structure of Chapter 5.

Confirmability — the degree to which findings reflect the data rather than researcher bias — was addressed by:

- (i) **Explicit role separation in practice.** Support interactions were classified at the time of logging as either *conceptual support* (clarifying ArchHypo) or *tool support* (addressing Hypo-Stage friction), matching the two support types in Phase 2 (Section 4.5). Facilitation entries were labeled accordingly in the field notes; during team rituals and unprompted usage, the researcher recorded what teams did without steering subsequent behavior, keeping investigator observations separable from facilitation actions.
- (ii) **Active logging of disconfirming evidence.** The researcher applied a standing analytic rule during note-taking and open coding: for each logged difficulty or barrier, the notes also captured counter-instances where teams derived value from the technique or tool. Positive instances received descriptive codes on the same footing as challenge codes—for example, *tool-positive-catalog-integration* alongside *registration-barrier-time-pressure* in Table 5.1.
- (iii) **Claims grounded in specific observations.** Every empirical claim in Chapter 5—including Section 5.5—is traceable to a Hypo-Stage artifact identifier, a dated field-note entry, or a paraphrased reconstruction from active-listening notes (not a verbatim quotation; Section 4.6.3). Table 5.1 maps each reported theme to its source codes and chapter section.

Construct validity. One deviation from prescribed practice warrants explicit disclosure. During the study, the teams and the researcher adopted a local convention of registering hypotheses only in a *negative formulation*, phrased to describe the evidence under which each hypothesis would be considered false. ArchHypo does not require this: it prescribes that hypotheses be *falsifiable* (refutable and testable) and recommends writing them as simple, direct statements of the team’s current belief (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). This convention is therefore a plausible cause of the Negative-Formulation finding reported in Chapter 5 (Section 5.4.2 and pattern A4, Section 5.6.1.4), and the associated cognitive cost may reflect the study’s operationalization rather than the technique itself. Future studies should follow ArchHypo’s prescribed free-form falsifiable statements to test whether the barrier persists.

4.9 Ethical Considerations

The study raises ethical considerations associated with insider research (ROBSON and McCARTAN, 2016). Participation was voluntary: team members were informed of

the research objectives, were free to use Hypo-Stage as much or as little as they chose, and were invited but not required to engage in active listening sessions. No personal performance data were collected or reported; all analysis pertains to team-level patterns and tool-level observations. The organization in which the study was conducted and the teams involved are anonymized in reporting. As no sensitive personal data were collected, no ethical review was required for this study.

The dual role of the researcher as organizational employee and academic investigator is a potential source of role conflict. This was managed by maintaining a clear separation between the researcher's responsibilities to the organization (deploying and supporting a useful tool) and his responsibilities as a researcher (reporting findings honestly, including negative ones). The research was conducted with the knowledge and informal consent of the relevant engineering leadership in the organization.

Chapter 5

Results and Refinements

This chapter presents the empirical findings of this study and the refinements applied to Hypo-Stage that emerged from them. The results are organized around the four research questions defined in Chapter 1:

- RQ1 addresses how teams used the tool and the technique in practice;
- RQ2 characterizes the architectural uncertainties they registered;
- RQ3 reports the benefits and challenges perceived by participants;
- RQ4 documents each refinement, tracing it to a concrete observation and to the question it addresses.

The section 5.6 goes beyond a direct answer to those four questions by drawing two families of emerging findings: a set of *Adoption-Condition Candidates* — empirically grounded candidate conditions (not validated patterns) observed to shape whether hypothesis engineering becomes a sustainable practice — and a set of novel *Structural Hypothesis Patterns* for hypothesis engineers. The Structural Hypothesis Patterns constitute a distinct contribution of this work, addressing a gap left open by prior research in the ArchHypo line, as discussed in Section 5.6.2. The Adoption-Condition Candidates describe problem and context only; the derivation of their solutions is deferred to future work.

5.1 Research Methodology Traceability

The evidence presented in this chapter was collected and analyzed strictly in accordance with the research design described in Chapter 4. This section makes that correspondence explicit so that compliance with the research methodology can be assessed at a glance; it does not restate the methodology itself, which is defined in Chapter 4.

The evidence draws on the three data sources specified in Section 4.6: participant observation, artifact analysis, and open-ended active-listening sessions. The quantitative summaries in Sections 5.3 and 5.5.8 are derived exclusively from the Hypo-Stage artifact records. As explained in Section 4.6.3, the active-listening sessions were not recorded;

paraphrased, role-attributed observations drawn from the researcher’s reconstructed notes are reported only when they add evidence beyond the recorded artifacts.

These data were analyzed with the three-stage thematic coding procedure defined in Section 4.7. Table 5.1 provides the local traceability that this analysis adds: it lists representative open codes and traces each one to its axial category, its research question (or emerging-pattern group), and the chapter section where the corresponding finding is developed.

Open code	Axial category	RQ / Group	Discussed in Section
<i>kafka-cluster-burst-registration</i>	Patterns of tool use; Structural archetypes	RQ1; RQ2	5.2, 5.6.2 (B1)
<i>senior-author-only</i>	Organizational adoption conditions	Adoption candidate	5.6.1 (A1)
<i>registration-barrier-time-pressure</i>	Perceived challenges (registration time cost)	RQ3	5.4.2
<i>tool-positive-catalog-integration</i>	Perceived benefits (catalog / Backstage integration)	RQ3	5.4.1
<i>falsifiability-confusion</i>	Perceived challenges (negative-formulation learning curve); Candidate refinements	RQ3; RQ4	5.4.2, 5.5

Table 5.1: *From open codes to chapter sections: traceability of the thematic coding process.*

The artifact data from Hypo-Stage were summarized descriptively (hypothesis counts, quality-attribute frequency, technical plan distribution) to complement and illustrate the qualitative findings. Following methodological guidance for flexible qualitative designs (ROBSON and MCCARTAN, 2016), the inclusion of these descriptive statistics is consistent with a flexible qualitative design and does not constitute a multistrategy approach.

5.2 RQ1 — How Teams Used Hypo-Stage in Practice

Unless noted otherwise, the empirical findings for RQ1–RQ3 in Sections 5.2–5.4 describe the full observation window cumulatively; the refinements documented under RQ4 were deployed iteratively during that same window, in the observe–refine–deploy cycle described in Section 4.5, so the results below are not restricted to the pre-refinement tool state.

5.2.1 Hypothesis Registration Dynamics

Over the observation period (Hypo-Stage registrations span 11 February 2026–12 April 2026, approximately nine calendar weeks or five two-week sprints), both teams

together registered thirteen architectural hypotheses in Hypo-Stage. Six hypotheses originated from the Product Team and seven from the Platform Team. This distribution is noteworthy because the Platform Team is significantly smaller (four engineers) yet produced more registered hypotheses than the twelve-engineer Product Team. This asymmetry is partly explained by the registration of a dense cluster of five Platform hypotheses on a single day (12 April 2026), all addressing complementary aspects of the shared Kafka streaming infrastructure — a burst pattern discussed further in Section 5.6.2.

Across both teams, twelve of the thirteen hypotheses were authored by senior technical leaders — senior or staff engineers who held explicit architectural responsibilities for their respective domains. Table 4.1 reports the underlying team composition: five senior and seven mid-level engineers on the Product Team, and two senior and two mid-level engineers on the Platform Team. One hypothesis was registered by a mid-level engineer on the Product Team. Both technical leaders in each team independently reported that a much larger number of architectural uncertainties existed in their daily work than the number they managed to register. The principal constraint they named was time: formulating a well-structured hypothesis statement requires focused, uninterrupted attention that sprint delivery pressure rarely afforded. This observation resonates with findings from prior ArchHypo evaluations, where the learning curve and process-integration demands were described as significant barriers to adoption (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025; SILVA, MELEGATI, WANG, *et al.*, 2024).

5.2.2 Role-Based Engagement

Participant observation and the active listening sessions revealed a clear stratification in how different roles engaged with Hypo-Stage during the study period. Senior engineers treated the tool as a decision-support instrument: they used it to externalize concerns that had until then existed only as tacit knowledge, and they referenced the hypothesis backlog when discussing architectural trade-offs in planning sessions. Mid-level and junior engineers were aware of the tool’s existence and attended the training session, but neither team saw them register hypotheses independently beyond the initial introduction. This does not indicate disengagement; rather, both technical leaders explained that these engineers typically operate at a scope where architectural uncertainties are resolved quickly — within a sprint or less — leaving little perceived need for formal registration.

5.2.3 Tool Adoption Patterns

Usage of Hypo-Stage was non-uniform across the observation window. Registration activity was concentrated in two periods: a first wave during and immediately after the training session, and a second, more intensive wave near the end of the period when the Platform Team used the tool for a structured architectural discussion on Kafka failure-handling strategies. Between these waves, both technical leaders reported using the hypothesis detail view to review existing records and update technical plans, but the tool was not used in team-visible ceremonies such as sprint retrospectives or planning meetings. When asked, participants noted that embedding the tool into these ceremonies would require an explicit cultural agreement within the team — something that had not yet occurred within the observation window.

The integration of Hypo-Stage into the Backstage Internal Developer Portal was cited by both teams as a meaningful enabler of adoption. Because Backstage is already the daily entry point for service catalog, documentation, and CI/CD status, engineers encountered Hypo-Stage without navigating to a separate tool. The EntityHypothesesTab, shown in Figure 5.1, which surfaces only the hypotheses linked to the component an engineer is already viewing, was specifically mentioned as reducing the cognitive cost of context-switching into hypothesis management.




Figure 5.1: EntityHypothesesTab: Hypo-Stage surfaces only the hypotheses linked to the Backstage catalog component the engineer is currently viewing, eliminating the need to navigate to a separate tool or filter by component.

5.3 RQ2 — Architectural Uncertainties Registered

5.3.1 Data Overview

Table 5.2 presents the thirteen hypotheses registered during the study, anonymized in accordance with the ethical agreement described in Section 4.9. Full exported records (statements, ratings, entity links, and author notes) appear in Appendix B (Section B.2, *Hypothesis Records*). Each hypothesis is identified by a short label, its source type (Requirements or Solution), its lifecycle status at the end of the observation period, and its uncertainty and impact ratings.

5.3.2 Source Types

Eight of the thirteen hypotheses were classified as *Requirements*-sourced: the uncertainty arose from an incompletely specified quality-attribute goal or an external obligation whose implementation path is unknown. Five were classified as *Solution*-sourced: the team has an existing or candidate architectural decision and is uncertain about its consequences.

ID	Short label	Team	Source	U	I
H01	Kotlin/JOOQ/Spring template readiness	Product	Req.	VH	VH
H02	Kafka direct exposure via platform framework	Platform	Req.	VL	M
H03	Decouple eng. logs from security/compliance audit	Platform	Req.	M	VH
H04	Blob storage GDPR right-to-erasure	Product	Req.	M	VL
H05	Credit-bureau context-agnostic restructuring	Product	Sol.	H	VH
H06	Billing aggregation chunked keyset pagination	Product	Sol.	VL	VH
H07	Blob ownership transfer across service domains	Product	Req.	VL	VH
H08	Blob cross-client sharing governance	Product	Req.	VL	VH
H09	Dead Letter Topic (DLT) for audit vs. offset-based replay	Platform	Sol.	M	VH
H10	Synchronous retry for transient consumer failures	Platform	Sol.	M	H
H11	Failure state inside vs. outside Kafka flow	Platform	Sol.	M	H
H12	Error classification before retry	Platform	Sol.	M	H
H13	Payload immutability before re-enqueueing	Platform	Sol.	VL	VH

Table 5.2: Overview of the thirteen registered architectural hypotheses. *Req.* = Requirements-sourced; *Sol.* = Solution-sourced. *U* = Uncertainty level; *I* = Impact level. Ratings: VL = Very Low, M = Medium, H = High, VH = Very High.

The predominance of requirements-sourced hypotheses on the Product Team (five out of six) and the predominance of solution-sourced hypotheses on the Platform Team’s Kafka cluster (four out of five Kafka hypotheses) reflects the different nature of each team’s work: the Product Team faces open questions about whether architectural goals can be met at all, while the Platform Team is assessing the trade-offs of specific implementation choices for a shared infrastructure component.

5.3.3 Quality Attribute Distribution

Table 5.3 summarizes the frequency of quality attributes across the thirteen hypotheses. Attributes were assigned by the hypothesis authors using the multi-select form in Hypo-Stage; a hypothesis may carry more than one attribute. *Because a single hypothesis was frequently classified under more than one quality attribute (NFR), the counts in Table 5.3 sum to well more than thirteen:* the fourteen attributes together account for thirty-seven classifications across the thirteen hypotheses.

Reliability and Auditability emerge as the two dominant concerns. The prominence of Auditability is especially notable and directly traceable to two contextual forces: the organization’s internal security audit requirements, which drove H03, and the applicable privacy regulation governing document storage (H04), as well as the Kafka immutability and DLT governance questions (H09, H13). This concentration reveals a recurring archetype in which external compliance obligations materialized into architectural hypotheses — a pattern discussed in Section 5.6.2 as the *Regulatory Trigger Hypothesis*.

5.3.4 Uncertainty and Impact Profile

Nine of the thirteen hypotheses carry a *Very High* or *High* impact rating, signalling that the teams were not registering peripheral concerns but uncertainties with potentially

Quality Attribute	Count	% of hypotheses
Reliability	6	46%
Auditability	5	38%
Performance	4	31%
Interoperability	4	31%
Extensibility	3	23%
Usability	3	23%
Scalability	2	15%
Modifiability	2	15%
Flexibility	2	15%
Configurability	2	15%
Maintainability	1	8%
Security	1	8%
Observability	1	8%
Reusability	1	8%

Table 5.3: *Quality attribute frequency across the thirteen hypotheses.*

significant consequences for schedule, rework, or system quality. Uncertainty ratings are more evenly distributed: five hypotheses are rated *Very Low* (the team has strong evidence but the hypothesis is tracked for governance reasons), four are *Medium*, one is *High*, and one *Very High*.

The combination of Very Low uncertainty and Very High impact (H06, H07, H08, H13) is particularly instructive. In ArchHypo’s model, this quadrant maps to hypotheses that *can be postponed* — the team already has sufficient evidence to act, so the immediate risk is low — yet the downstream consequences of acting incorrectly are severe. Hypo-Stage’s prioritization dashboard surfaced these hypotheses in the "Can Postpone" segment of its focus panel, enabling teams to consciously defer them without losing visibility. This represents a practical instance of the *Postpone Uncertain Decisions* pattern described in the ArchHypo pattern literature (SILVA, ADOLFO, *et al.*, 2024).

5.3.5 Technical Plans in Use

Seven of the thirteen hypotheses had at least one associated technical plan at the end of the observation period. Together these seven hypotheses carried thirteen individual technical planning entries. The complete catalog of plan titles and strategies is reproduced in Section B.3 of Appendix B. Table 5.4 summarizes the strategies used.

Guideline Implementation and Architectural Spikes were the most frequently chosen technical plans, consistent with findings from the PropSEP/Jetsoft evaluation (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). The Tracer Bullet plan appeared once for hypothesis H10, where the Platform Team planned a controlled injection of transient failures into a real Kafka consumer path to measure ordering and latency characteristics. This application illustrates how the Tracer Bullet differs from an Architectural Spike. Rather than an isolated experiment, the Tracer Bullet plan exercises the full production consumer path under representative conditions (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025).

Strategy	Plans	Hypotheses using it
Guideline Implementation	3	3
Spike (Architectural Spike)	3	2
Modularization / Flexibility	2	1
Experiment	2	2
Tracer Bullet	1	1
Software Analytics (Trigger)	1	1
Other (implementation)	1	1
Total	13	—

Table 5.4: *Technical plan strategies observed across the hypothesis data.*

Six hypotheses had no technical plan created within the observation window. Three of these (H01, H04, H02) were registered but not yet acted upon; two (H07, H08) were validated by planning actions taken *before* the observation window began, and their technical plans were entered into Hypo-Stage only afterward, as a record of decisions already made rather than plans that guided work during the study; and one (H05) was under active discovery, with its plan document still being prepared.

5.3.6 Lifecycle and Status

At the end of the observation period, ten hypotheses had *Open* status, one was *In Review* (H05), and two were *Validated* (H07 and H08, both in the document storage domain). No hypothesis was invalidated or discarded during the study window. The predominance of the Open status is consistent with the short duration of the study over the observation window (Section 5.2.1) and with the nature of the uncertainties: most concern mid- to long-horizon architectural concerns that teams do not expect to resolve within a single sprint cycle.

5.4 RQ3 — Perceived Benefits and Challenges

5.4.1 Perceived Benefits

Active listening sessions with all four interviewed participants (the two technical leaders from each team) converged on four recurring themes regarding the perceived value of using Hypo-Stage. Because the sessions were not recorded (Section 4.6.3), the attributions below are the researcher’s paraphrased reconstructions from dated field notes and post-session written recollection—not verbatim quotations. The coding basis for the perceived-benefit theme appears in Table 5.1 (*tool-positive-catalog-integration*); challenge themes are developed in Section 5.4.2.

Externalization of tacit architectural knowledge: Both teams described the act of registering a hypothesis as valuable in itself. Participants noted that seeing a hypothesis written down with a clear statement, an explicit uncertainty rating, an impact rating, and a related quality attribute made it easier to share the concern with other colleagues who had not been involved in the original discussion.

Prioritization support: The *Where to Focus* panel of the Hypo-Stage dashboard, distinguishing hypotheses that *Need Attention* (high uncertainty and high impact) from those that *Can Postpone* (low impact), was described by both technical leaders as a useful input in architectural conversations. One technical leader recalled that the “Can Postpone” category in particular gave them a way to explain to other engineers outside of the team why a known concern was being deliberately deferred rather than being ignored.

Collective memory for architectural decisions: Both teams observed that the hypothesis backlog served a documentation function that was distinct from Architecture Decision Records (ADRs) or design documents. Whereas ADRs capture decisions after they are made, an architectural hypothesis captures an open question before a decision is available. Participants described this as filling a gap in their current toolchain: they had good mechanisms for recording what was decided but poor mechanisms for recording what was still unknown. Chapter 6 revisits this distinction when discussing implications for practice.

Catalog and Backstage workflow integration: Both teams described surfacing Hypo-Stage within the Backstage Internal Developer Portal as a perceived benefit rather than merely an adoption convenience. Because Backstage already served as the daily entry point for service catalog, documentation, and CI/CD status, engineers could review and register hypotheses without navigating to a separate tool or breaking out of the component context they were already working in. Both technical leaders recalled that the `EntityHypothesesTab`, which lists only hypotheses linked to the catalog component currently being viewed, reduced the cognitive cost of context-switching into hypothesis management. The open code *tool-positive-catalog-integration* labels this reconstructed theme; Section 5.5.1 later describes how the study operationalized it as a refinement.

5.4.2 Perceived Challenges

Three challenges were consistently raised across participants. As with the benefits above, these are reconstructed, role-attributed observations rather than verbatim quotations; their open-code anchors are *falsifiability-confusion* and *registration-barrier-time-pressure* (Table 5.1).

Negative-formulation learning curve: In this study, the teams and the researcher adopted a local convention of registering only hypotheses phrased to specify the evidence under which the team would conclude the hypothesis is incorrect—an inversion of the team’s belief into a “what would prove this false” form. All four participants described this inversion as non-intuitive: engineers are accustomed to stating what they believe to be true, and rephrasing that belief into its refutation condition requires a deliberate cognitive effort that is easy to skip under time pressure. This convention is stricter than what ArchHypo actually prescribes: the technique asks only that a hypothesis be a *falsifiable statement* (a refutable, testable claim in which the possibility of proving it false exists, formulated as a simple, direct statement of the team’s current belief; Section 2.2.1) (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). The cognitive cost observed here is therefore largely a property of the study’s negative-formulation convention rather than of ArchHypo’s guidance. The

underlying *learning curve*, however, is consistent with the difficulty reported in both prior ArchHypo evaluations (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025; SILVA, MELEGATI, WANG, *et al.*, 2024); Section 4.8 discusses this deviation as a threat to construct validity. The open code *falsifiability-confusion* labels this reconstructed theme.

Registration time cost: Both technical leaders described the difficulty of dedicating the focused time required to register a hypothesis properly. Fast-moving architectural micro-decisions, usually resolved within a day or less, were rarely registered because, by the time a hypothesis could be written, the question had already been answered informally. This creates a systematic under-representation of short-lifecycle uncertainties in the hypothesis backlog, which may be consequential when the same question re-emerges in a future context. The open code *registration-barrier-time-pressure* labels this reconstructed theme.

Absence of team-visible ceremonies: Neither team had established a recurring ceremony, such as an architectural review or a hypothesis triage meeting, in which Hypo-Stage records were reviewed as a team. Tool usage remained largely individual during the observation period.

5.5 RQ4 — Refinements to Hypo-Stage

This section presents the refinements implemented in Hypo-Stage from commit ceee509 (the artifact boundary established by CORRÊA 2025) to commit 3a53d4e. Each refinement was traced to a concrete observation from the study and linked to the research question it primarily addresses. All refinements address RQ4 by definition; selected refinements also related to RQ1 (adoption ease) and RQ3 (usability and learning curve).

5.5.1 Backstage Integration and Catalog Surfacing

Observation: During the training session and the first week of observation, both teams navigated to Hypo-Stage via a standalone entry point that was separate from their regular Backstage workflow. Engineers who were not technical leaders rarely revisited the tool after the initial session.

Refinement: Hypo-Stage was wired as a first-class Backstage experience, with routable pages for home, hypothesis creation, detail view, and edit flows. Crucially, a catalog component tab (`EntityHypothesesTab`), shown in Figure 5.1, was introduced, surfacing all hypotheses linked to the component an engineer is currently viewing in the Backstage catalog.

Impact: Reducing navigation friction lowered the perceived cost of checking the hypothesis backlog while already working with a service’s catalog entry.

5.5.2 Backend API and Filtering

Observation: Engineers with cross-team or platform-wide responsibilities needed to filter hypotheses by team and by specific service. The initial version of Hypo-Stage

returned all hypotheses without server-side filtering, which degraded usability as the hypothesis data grew.

Refinements: Three new API capabilities were implemented:

- GET /hypotheses extended with entityRef and team query parameters;
- GET /hypotheses/referenced-entity-refs to populate the component filter without listing the entire catalog;
- GET /hypotheses/teams deriving team names from catalog metadata for components referenced by hypotheses.

Impact: On the Hypo-Stage home page, these capabilities enabled filtering of the hypothesis list by team and by service (the Backstage catalog entity linked to a hypothesis), so engineers with cross-team or platform-wide responsibilities could narrow the backlog without loading the full unfiltered dataset.

5.5.3 Prioritization Dashboard

Observation: The technical leaders expressed the desire to quickly identify which hypotheses deserved immediate attention and which could safely be deferred. Manually scanning the full hypothesis list imposed an unnecessary cognitive burden, especially when the hypothesis data grew.

Refinement: A HypothesisListDashboard component was introduced above the hypothesis table (Figure 5.2), presenting:

- Overview cards with total count, validated count, and hypotheses created in the last thirty days;
- A *Where to Focus* section distinguishing *Need Attention* (high uncertainty and high impact) from *Can Postpone* (low impact), with explanatory subtext aligned with ArchHypo's prioritization model (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025);
- Distribution panels for uncertainty and impact Likert values.

The dashboard respects entity, team, and component scoping.

Impact (RQ3). The "Can Postpone" segment was specifically cited by participants as resolving a communication challenge: it provided a principled vocabulary for explaining deliberate deferral to non-technical stakeholders.

5.5.4 Hypothesis Lifecycle Consistency

Observation: In the initial implementation, create and update operations were not wrapped in database transactions, which created a risk of partial writes (hypothesis row updated without the corresponding event entry, or vice versa). Artifact analysis revealed two records with missing event entries, suggesting that this race condition had occurred in practice.

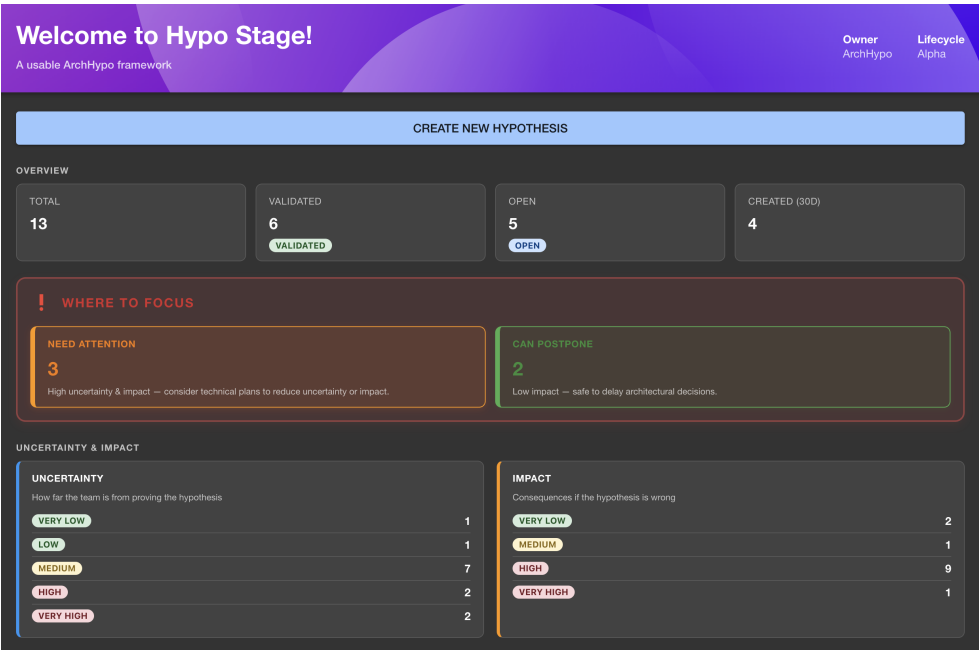


Figure 5.2: Prioritization Dashboard in Hypo-Stage: the focus panel surfaces hypotheses according to the urgency signals derived from their uncertainty–impact combination, operationalizing the Postpone Uncertain Decisions pattern from the ArchHypo pattern literature (SILVA, ADOLFO, et al., 2024).

Refinements: Create and update paths were refactored to use explicit database transactions. Delete was extended to cascade to technical planning rows and events, with clear API semantics and a confirmation dialog that requires the exact hypothesis statement before proceeding (paste-enabled to avoid penalizing users with long statements). Success and error toast notifications were added for destructive operations.

5.5.5 Technical-Planning-Driven Assessment Updates

Observation: Engineers reported that after completing a technical plan action, such as an architectural spike, they updated the hypothesis’s uncertainty rating separately from the plan record. This two-step process was error-prone and led to hypotheses whose assessment was stale relative to their technical plans.

Refinements: Technical planning create and update operations were extended to accept optional uncertainty and impact fields. When present, the service updates the parent hypothesis’s Likert values in the same transaction. The event model was refined to distinguish TECHNICAL_PLANNING_CREATE and TECHNICAL_PLANNING_UPDATE events from generic hypothesis updates, preserving a traceable audit trail of which assessment changes were driven by plan actions.

5.5.6 Evolution Chart Improvements

Observation: The initial uncertainty/impact evolution chart used opaque numeric y-axis ticks that participants could not interpret without consulting the Likert scale documentation separately. Multiple events on the same calendar day produced ambiguous

x-axis labels.

Refinements: The chart was updated to render verbal y-axis ticks (Very Low through Very High). Event source was made visually distinguishable: circles for manual or creation-driven points, squares for technical-planning-driven points. A legend explaining the marker shapes was added. When multiple events occur on the same calendar day, the x-axis uses date and time to avoid ambiguity. Tooltips were improved to label technical-planning-driven assessment moves with the plan number when identifiable.




Figure 5.3: Evolution Chart in Hypo-Stage: tracking uncertainty and impact rating changes across successive assessments for a registered hypothesis. The screenshot is illustrative—a representative example of the chart rather than a depiction of one specific hypothesis from the study dataset. Each point corresponds to a saved assessment; the chart makes visible whether technical plan actions are reducing uncertainty or impact over time.

5.5.7 Form Usability Improvements

Observation: Participants found the hypothesis creation form cognitively demanding, particularly the uncertainty and impact selectors, which displayed only numeric values, and the quality attribute field, which accepted free text without suggesting valid values.

Refinements: Uncertainty and impact selectors were redesigned as Likert star pickers showing both a numeric button (1–5) and an adjacent text label (Very Low through Very High) with disambiguating icons. Quality attributes were replaced by a multi-select autocomplete showing attribute name and description in the dropdown. Entity references for linking hypotheses to catalog components were given a catalog-backed autocomplete. Documentation link fields were given an improved list editor with add/remove and basic URL validation. Figure 5.4 shows the improved hypothesis creation form.

5.5.8 Summary: Refinements Mapped to Observations

Table 5.5 summarizes the nine categories of refinements, the primary observation that motivated each, and the research question each addresses most directly.

+ Create New Hypothesis
Define your hypothesis and assess its uncertainty and potential impact

Entity References
Please select at least one entity reference. Type to search catalog components.

Hypothesis Statement*
0/500 characters

Source Type*

Uncertainty Level
How far is the team from validating/falsifying this hypothesis?

1 2 3 4 5 **Not rated**

Impact Level
What is the potential impact on the system if the hypothesis proves false?

1 2 3 4 5 **Not rated**

Quality Attributes

Quality Attributes (select one or more)
Please select at least one quality attribute that this hypothesis affects. Click the field to open the list; each option shows a short concept in parentheses.

Figure 5.4: Hypothesis creation form in Hypo-Stage after the usability-oriented form refinements in this section: inline contextual guidance adjacent to the Statement field scaffolds the formulation of falsifiable hypotheses, reducing the cognitive barrier at the moment of registration.

5.6 Emerging Findings from the Case Study

Beyond the direct answers to RQ1–RQ4, thematic analysis of observation notes, reconstructed active-listening notes, and the combined hypothesis data produced two families of emerging findings. The **Adoption-Condition Candidates** describe team-behavior conditions that shape whether hypothesis engineering becomes a sustainable practice; they are candidate conditions rather than validated patterns. The **Structural Hypothesis Patterns** describe emerging patterns in the content of the architectural hypotheses themselves. These two families address orthogonal dimensions: the Adoption-Condition Candidates ask *when* hypothesis engineering works; the Structural Hypothesis Patterns ask *what shape* hypotheses take in practice. Table 5.6 at the end of this section summarizes the nine items across both families.

Scope and statistical caveat: The findings described in this section were identified through thematic analysis of thirteen architectural hypotheses produced by two software engineering teams within a single organizational context over the observation window (Section 5.2.1). This dataset is sufficient to motivate conceptual characterization and to propose initial vocabulary for the registered hypothesis, but it does not support statistical claims about frequency, generalizability, or prevalence. The Adoption-Condition Candidates and Structural Hypothesis Patterns should therefore be read as *emerging propositions* — grounded in real empirical data and theoretically coherent — but requiring replication across additional teams, organizations, and domains before they can be treated as consolidated findings. Future work directions for each item are discussed in Section 6.3.

Refinement	Motivating observation	RQ1	RQ4	RQ3
Backstage catalog tab	Infrequent revisit by non-leaders outside native workflow	S	P	S
Backend API filtering	Cross-team users needed scoped views as hypothesis data grew	S	P	—
Prioritization dashboard	Manual scanning of full backlog was cognitively costly	—	P	S
Lifecycle transaction consistency	Partial writes produced stale event records	—	P	—
TP-driven assessment updates	Assessment and plan records drifted apart post-spike	—	P	S
Evolution chart improvements	Numeric y-axis and same-day collisions caused confusion	—	P	S
Form usability improvements	Numeric-only selectors; free-text QA field caused errors	—	P	P
Safer deletion UX	Long statements made confirmation hostile without paste	—	P	—
Toast notifications	No feedback after destructive operations caused re-submission	—	P	—

Table 5.5: Refinements to Hypo-Stage mapped to motivating observations and research questions. RQ columns indicate primary (P) or secondary (S) linkage.

5.6.1 Adoption-Condition Candidates for Hypothesis Engineering

The four Adoption-Condition Candidates emerge by cross-referencing RQ1 and RQ3 findings. They describe organizational and individual-level conditions that the study observed to facilitate or obstruct the adoption of hypothesis engineering as a team practice. They extend the existing ArchHypo patterns observed by [SILVA, ADOLFO, et al. \(2024\)](#) and [SILVA, ADOLFO, et al. \(2024\)](#) by addressing the social and operational conditions in which ArchHypo mechanics operate in practice.

What these items are (and are not). The four items described below (A1–A4) are *candidate conditions* under which hypothesis engineering can be sustained as a team practice. These are *not* validated patterns. Unlike a pattern, which pairs a recurring problem with a proven solution, each candidate here characterizes only the *problem and context* it was observed in. They are emergent rather than confirmatory findings: each is supported by convergent evidence from both teams, but they have not been tested in the broader population of software engineering organizations. Accordingly, the *solution component is deliberately absent*: the derivation of actionable solutions from each candidate is deferred to the future research agenda (Section 6.3), which outlines the empirical validation required before any candidate could be consolidated into a pattern or adopted as a guideline.

5.6.1.1 Senior-Curated Hypothesis Backlog

Problem: In a team whose engineers operate at different scopes and timescales, who is positioned to surface and sustain the registration of architectural hypotheses in mid- and long-terms?

Observed evidence: Twelve of the thirteen registered hypotheses were authored by senior technical leaders — engineers who carried explicit architectural responsibility and whose work extended beyond the current sprint. Non-senior engineers did not sustain independent hypothesis registration after the initial training episode. The hypotheses that were registered shared a common trait: they concerned architectural uncertainties whose consequences would materialize weeks or months in the future, making them invisible to typical sprint-scope review practices.

5.6.1.2 Long-Horizon Registration Priority

Problem: Many architectural micro-decisions are made and resolved within hours or a single day. Rather than treating their speed as a blind spot to be closed, the actionable observation is the inverse: fast-resolved micro-decisions rarely repay the cost of registration, whereas the uncertainties genuinely worth registering are the slow, high-consequence ones whose architectural effects only materialize weeks or months later. Registration effort is therefore best concentrated on long-horizon uncertainties rather than spread evenly across every transient decision.

Observed evidence: Both technical leaders reported that the majority of architectural decisions encountered day-to-day were resolved before a hypothesis could be written, and that registration felt retrospectively unnecessary for them. Notably, these fast-resolved decisions were not captured as Architecture Decision Records (ADRs) either, so their absence from Hypo-Stage reflects a genuine registration threshold rather than a tool-specific gap. Yet both also described recurring situations where the same question re-surfaced in a new context, requiring the team to reconstruct reasoning that had previously been resolved informally and in a rapid pace — a residual cost the leaders judged acceptable relative to the effort of registering every micro-decision, which reinforces that the payoff of registration concentrates in the slow, high-consequence uncertainties.

5.6.1.3 Registration as a Slack-Dependent Practice

Problem: Translating a tacit architectural concern into a properly formulated Arch-Hypo hypothesis requires focused, uninterrupted attention. Sprint delivery pressure makes this kind of attention scarce and unreliable as a basis for sustained practice.

Observed evidence: Senior leaders explicitly described the difficulty of "dedicating proper time to register all possible hypotheses they encounter day-to-day." The observed set of thirteen hypotheses is known to under-represent the total uncertainties that leaders were aware of. Registration activity clustered around protected and restricted discussion time rather than being distributed across normal sprint work.

5.6.1.4 Negative-Formulation Barrier

Problem: Under the negative-formulation convention adopted in this study — registering hypotheses phrased to describe what would prove them false—writing a hypothesis required a cognitive inversion of how engineers typically articulate their beliefs. This inversion was a persistent adoption barrier even for practitioners who understood the technique conceptually. The barrier is conditioned on that convention: it is not an inherent requirement of ArchHypo, which prescribes free-form falsifiable statements of the team’s current belief (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025) and admits optional interrogative templates such as QUEST (MELEGATI and WANG, 2019).

Observed evidence: All four interviewed participants mentioned that formulating hypotheses under this convention was non-intuitive and difficult to apply under time pressure. The associated *learning curve* replicates findings reported in the PropSEP/Jetsoft evaluation (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025); the *negative-formulation* framing, however, is specific to this study’s convention. Whether ArchHypo’s prescribed free-form falsifiable statements avoid the barrier is a hypothesis for future work rather than an established property of the technique.

5.6.2 Structural Hypothesis Patterns for Hypothesis Engineers

The prior work in the ArchHypo research line has established the technique’s core concepts (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025), documented patterns for hypothesis assessment, uncertainty management, and decision deferral (SILVA, ADOLFO, *et al.*, 2024; SILVA, ADOLFO, *et al.*, 2024), and reported on the technique’s learning curve and process-integration challenges (SILVA, MELEGATI, WANG, *et al.*, 2024). What this body of work has not yet addressed is a question that only becomes answerable through empirical data on hypotheses registered in practice:

What recurring shapes do architectural hypotheses take when teams apply ArchHypo to real systems? And what structural properties of a hypothesis, its source type, quality-attribute profile, and relation to its organizational context, should an engineer recognize in order to write and manage it more effectively?

This gap was also identified in calls for future work within the ArchHypo pattern literature, which explicitly noted the need for "identifying and documenting patterns for various types of hypotheses" and for "exploration of common uncertainties related to specific quality attributes" (SILVA, ADOLFO, *et al.*, 2024). The thirteen-hypothesis dataset produced by this study constitutes the first empirical dataset that enables such identification. The five patterns described below represent a **novel contribution** of this work: they are grounded in real hypothesis data, not derived from theoretical analysis, and they provide a vocabulary that engineers can use both to recognize familiar situations more quickly and to communicate about hypothesis types across teams.

Statistical scope: All five patterns were identified from the same thirteen-hypothesis dataset. While each pattern maps to two or more hypotheses in the dataset, the small dataset size means that no frequency claims are made. The contribution lies in naming and structuring recurring shapes so that future practitioners and researchers have a vocabulary for recognition, not in establishing that any pattern will appear with a particular frequency

in other organizational contexts, nor in prescribing solutions for each one. Solution-shaped guidance and consolidation into the broader ArchHypo pattern language are delegated to the future-work research agenda in (Section 6.3).

5.6.2.1 Hypothesis Cluster

Context: A team is responsible for a shared platform component on which multiple other service teams depend. An architectural question about that component cannot be decomposed into a single hypothesis because the design space has several interrelated dimensions, each of which involves a distinct trade-off.

Problem: Formulating a single hypothesis for a multi-dimensional architectural concern either oversimplifies the problem (losing important trade-off structure) or produces an unmanageably broad statement that is hard to falsify.

Observed evidence: Five hypotheses (H09–H13) all concern the shared Kafka streaming infrastructure and were registered on the same day, as the Platform Team’s technical leader used Hypo-Stage to structure a design discussion. The five hypotheses address: (i) DLT audit semantics versus offset-based replay; (ii) synchronous versus asynchronous retry for transient consumer failures; (iii) failure-state placement inside versus outside the Kafka flow; (iv) error classification before retry; and (v) payload immutability before re-enqueueing. Each is independently falsifiable yet together they map the full decision space for consumer-side failure handling in the shared streaming platform.

5.6.2.2 Regulatory Trigger Hypothesis

Context: A software system operates in a regulated domain such as financial services, healthcare, or data privacy, where external compliance obligations create architectural constraints whose implementation path is not fully known.

Problem: A compliance requirement imposes an architectural obligation before the team has determined how to satisfy it. It is unclear whether the existing architecture already satisfies the obligation, what changes would be required if it does not, and what the cost of non-compliance would be.

Observed evidence: Two hypotheses in the dataset arise directly from regulatory obligations: H04 (the GDPR, the General Data Protection Regulation, right-to-erasure requirement for binary documents in the blob storage service) and H03 (organization-wide security audit requirements for database log correlation across the platform framework). In both cases the uncertainty is not about whether to comply — that is a legal requirement that must be met — but about whether the current architecture *can* comply and at what technical cost.

5.6.2.3 Service Decomposition Hypothesis

Context: A service has grown to handle multiple responsibilities such as domain logic, external integrations, and identity management, and the team is considering decomposing it into smaller services aligned with distinct domain boundaries — the recurring software-architecture-evolution challenge of splitting a coarse-grained service into microservices

along business-capability or subdomain lines (RICHARDSON, 2018). What the decomposition would entail — the correct contract boundaries and the migration path — remains unknown.

Problem: Deciding *whether* and *how* to decompose the service is itself uncertain: the team does not know where the correct service boundary should fall, which existing flows can be migrated without business-logic changes, and what the downstream impact on other teams would be. As RICHARDSON (2018) argues, decomposition can proceed by business capability or by subdomain, but committing to a boundary without this information risks a disruptive refactor that delivers less reuse than expected.

Observed evidence: H05 (credit-bureau service decomposition) captures this precisely. The service currently handles identity, bureau integration, and confidential document flows using domain-specific entities tied to a single business vertical. The hypothesis asks whether these integrations can be decoupled into a context-agnostic middleware service accepting only each provider's minimum required identifiers, reusable across domains without duplicated cost. At the end of the observation period, H05 was the only hypothesis in *In Review* status, with a discovery technical plan mapping (i) the dependency matrix of current versus real dependencies per integration, (ii) an agnostic interface contract proposal, and (iii) a migration plan identifying which flows can isolate immediately.

5.6.2.4 Monitored Quality-Attribute Hypothesis

Context: A hypothesis targets a quality attribute — such as performance or scalability under load — that is affected by many interacting factors (data volume, access patterns, resource limits, runtime behavior). The team has proposed a specific algorithmic or implementation approach and has already conducted a capacity or volume test at scale, producing evidence that the approach satisfies the target. However, the hypothesis has not been formally closed because the implementation is still in progress.

Problem: The hypothesis occupies the lowest uncertainty quadrant (Very Low) while retaining the highest impact rating (Very High), reflecting a situation where the team *knows* the solution works but the architectural commitment is still consequential. Because the quality attribute depends on many factors that can shift as the system evolves, it must be *continuously monitored* even after the uncertainty is validated: the hypothesis must be tracked until implementation is complete and the attribute confirmed in production, yet its characterization as "Open" may suggest unresolved doubt that no longer exists.

Observed evidence: H06 (billing aggregation chunked keyset pagination) exemplifies this. The team validated the approach — chunked keyset pagination with isolated repository-level transactions — through a capacity test at five times production scale (45 million rows total, 1.5 million rows per month). Profiling confirmed that JPA stream processing caused first-level cache growth, and that chunking scopes transactions to approximately ten thousand rows so the garbage collector can reclaim memory, producing a healthy sawtooth heap pattern. The hypothesis remains Open because the implementation is ongoing: the uncertainty has been effectively resolved, but the performance quality attribute continues to be monitored while the solution is rolled out.

5.6.2.5 Shared-Resource Hypothesis

Context: A platform service manages a shared resource such as binary objects, credentials, and configuration data, on which multiple product domains depend. New business requirements demand new ways of accessing that resource across domain boundaries, raising interacting consistency, security, and usability concerns whose resolution is not yet known — of which governance (ownership, cross-domain read/write, and transfer between owning services) is one prominent facet.

Problem: The team must decide how to model cross-domain access to the shared resource in a way that keeps it consistent, secure, and usable across domains. One central facet is governance — specifically, which service "owns" the resource at any given point, who may read or write it, and under what conditions ownership can be transferred — but these decisions interact with broader consistency, security, and usability concerns. Without explicit decisions, different product teams will adopt incompatible access patterns, creating consistency violations that are difficult to detect and remedy.

Observed evidence: H07 (blob ownership transfer across service domains) and H08 (blob cross-client sharing governance) are two instances of this pattern in the secure blob storage domain. H08 was driven by a concrete usage scenario in which a customer-facing channel receives documents that are later routed to support tooling; the support team needs read access without breaking the blob's ownership invariants. Both hypotheses were validated through a Guideline technical plan composed of an alignment session held before the observation window and registered in Hypo-Stage retrospectively, producing a governance decision enabling prioritization of the transfer and sharing implementation work.

5.7 Revisiting the Research Questions

This section concludes the analytical observations of the chapter by stating explicitly how the preceding empirical and design material responds to the four research questions formulated in Chapter 1. The intent is *consolidative* where no additional empirical claims are introduced. The aim is to make the logical organization of the chapter legible relative to the research questions.

5.7.1 Correspondence between research questions and chapter sections

Table 5.7 summarizes where each research question is principally addressed in this chapter, and where integrated interpretation (the Adoption-Condition Candidates and Structural Hypothesis Patterns) extends the principal sections.

5.7.2 Brief interpretive synthesis

Through the preceding analysis presented in this chapter, the four research questions tell a single story about how far a tool-mediated instantiation of ArchHypo can carry a team, and where the remaining barriers lie.

ID	Family	Name	Signature	Primary evidence
A1	Adoption candidate	Senior-Curated Hypothesis Backlog	Hy-Mid-/long-horizon uncertainties stay visible because senior engineers author them	12 of 13 hypotheses senior-authored
A2	Adoption candidate	Long-Horizon Registration Priority	Registration value concentrates in slow, high-consequence uncertainties; fast-resolved micro-decisions rarely repay the cost (and are absent from ADRs too)	Both technical leaders' testimony
A3	Adoption candidate	Registration as a Slack-Dependent Practice	Well-formulated hypotheses require protected time, not sprint slack	Registration clustering at training and Kafka session
A4	Adoption candidate	Negative-Formulation Barrier	The study's negative-formulation convention (not ArchHypo itself) makes hypothesis writing cognitively costly	All four interviewed participants
B1	Structural pattern	Hypothesis Cluster	Multi-dimensional platform concern decomposes into linked falsifiable hypotheses	H09–H13 (Kafka failure handling)
B2	Structural pattern	Regulatory Trigger Hypothesis	External compliance obligation surfaces as tracked architectural uncertainty	H03, H04
B3	Structural pattern	Service Decomposition Hypothesis	Uncertainty about whether and how to decompose a service into microservices along domain boundaries	H05 (credit-bureau service decomposition)
B4	Structural pattern	Monitored Quality-Attribute Hypothesis	Very-low uncertainty, very-high impact: quality attribute empirically validated but continuously monitored while implementation is open	H06 (billing aggregation)
B5	Structural pattern	Shared-Resource Hypothesis	Cross-domain consistency, security, and usability concerns — including ownership governance — for a shared platform resource	H07, H08

Table 5.6: *Emerging findings identified in the case study: the four Adoption-Condition Candidates (A1–A4) and the five Structural Hypothesis Patterns (B1–B5). Both families describe problem and context only; solution-shaped guidance for each item is deferred to the future research agenda (Section 6.3).*

RQ	Principal analysis	Integrated interpretation
RQ1	Section 5.2 (tool use and adoption dynamics).	Section 5.6.1 (Adoption-Condition Candidates; read together with RQ3).
RQ2	Section 5.3 (typological account of the data).	Section 5.6.2 (Structural Hypothesis Patterns).
RQ3	Section 5.4 (benefits and challenges).	Section 5.6.1 (Adoption-Condition Candidates; read together with RQ1).
RQ4	Section 5.5 (nine refinement categories, — observation-led).	—

Table 5.7: *Correspondence between the research questions (Chapter 1) and the organization of evidence in this chapter.*

RQ1 and RQ3 converge on an organizational, not a tool, problem: Both questions produce evidence that Hypo-Stage lowered surface friction, facilitating context-switching, filtering, traceability. Participants recognized the intellectual value of hypothesis management. Still, the same two questions also show that authorship stayed concentrated in senior engineers, who registered hypotheses clustered around protected discussion time rather than distributing across sprint work, and that hypothesis writing under the study’s negative-formulation convention remained cognitively costly. The Adoption-Condition Candidates (A1–A4) name the conditions under which those findings hold: adoption breadth is limited by the organization structure including senior curation, protected cognitive time, statement-level support, and ceremonial touchpoints. These are conditions that a tool alone cannot supply.

RQ2 and the Structural Hypothesis Patterns reveal a content gap in the existing ArchHypo pattern language: The prior pattern literature (SILVA, ADOLFO, *et al.*, 2024; SILVA, ADOLFO, *et al.*, 2024) names mechanics of assessment, deferral, and uncertainty management, but does not name the *types* of hypotheses engineers actually register. When applied in a compliance-heavy, platform-dominant environment, the thirteen registered hypotheses concentrate on regulatory triggers, shared-resource governance, context-agnostic restructuring, and multi-dimensional platform concerns, which are categories the existing patterns do not index. The Structural Hypothesis Patterns therefore complement the mechanics-focused patterns rather than competing with them, and provide the first empirically grounded vocabulary for the content of hypotheses that emerge in practice.

RQ4 closes the formative-evaluation loop: The nine refinement categories are not a separate implementation narrative. Each maps to a specific observed friction surfaced by RQ1, RQ2, and RQ3 questions: the form-level and ceremony-related refinements track A3 and A4; the component-scoped tab tracks B1-shaped concerns about multi-dimensional platform uncertainty; the prioritization dashboard and API filtering respond to the adoption dynamics observed under RQ1. Their derivation from the same observational data is itself evidence that the design-science loop of observation, refinement, and re-observation operated on the Hypo-Stage tool during the study.

Taken together, the chapter supports a bounded claim: tool-mediated hypothesis

engineering succeeds at the integration and traceability layers and produces a useful vocabulary for the content of architectural hypotheses, but its adoption ceiling lies at the organizational layer. The Adoption-Condition Candidates are the natural hand-off to Chapter 6, where they are read as open empirical problems rather than as actionable guidelines.

Chapter 6

Conclusion

The central objective motivating this present work was whether embedding the ArchHypo technique into a tool deployed within a team’s existing workflow could lower the barriers that prior evaluations had identified as threats to its adoption. Both prior studies of ArchHypo — at the Catch Solve startup (SILVA, MELEGATI, WANG, *et al.*, 2024) and at the PropSEP mission-critical project (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025) — demonstrated the technique’s conceptual value, yet neither involved a supporting tool or a multi-team professional software development setting, and both documented a consistent limitation: without tool support and expert facilitation, hypothesis engineering remained a fragile, senior-dependent practice. This work takes up that limitation directly. By deploying Hypo-Stage as a Backstage plugin within a large-scale financial services organization and observing two teams over approximately nine calendar weeks of real sprint work, the study produced empirical data of ArchHypo in a tool-mediated setting. This contribution extended the base of the research line whose prior publications have called explicitly for this kind of evidence (SILVA, ADOLFO, *et al.*, 2024).

The results both confirmed and revealed the expectations established by prior work. On the side of confirmation: hypothesis registration remained predominantly senior-led across both teams, formulating hypotheses under the study’s negative-formulation convention continued to be described as cognitively demanding under sprint pressure, and the absence of team-visible ceremonies for reviewing the hypothesis backlog emerged as a structural obstacle. These observations are consistent with findings from the PropSEP evaluation (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). On the side of novelty: the Hypo-Stage tool integrated in the Backstage proved to be a genuinely meaningful enabler rather than merely a convenience. Visualizing and managing hypotheses in the context of a developer portal that is already adopted by the team reduced context-switching costs in a way that participants named explicitly as a facilitating factor over standalone tools. This suggests that the tool’s usage matters as much as the hypothesis engineering mechanics themselves.

6.1 Synthesis of Contributions

The four research questions of this work produced four contributions that, taken together, described what a tool-mediated instantiation of ArchHypo can deliver in a professional team setting. The contributions were empirical (RQ1, RQ2, RQ3) and an engineering/design contribution (RQ4)—a worked instance of formative evaluation for hypothesis-engineering tools; each was supported by the thirteen-hypothesis dataset and the field-study evidence reported in Chapter 5.

RQ1 — An empirical characterization of senior-led, event-anchored tool use: When ArchHypo was delivered through a tool embedded in the team’s daily workflow, hypothesis registration did not distribute across continuous sprint work but instead concentrated in senior technical leaders and clustered around discrete organizational events such as protected discussion time, training sessions, and dedicated technical meetings such as the Platform Team’s Kafka failure-handling session. The term *event-anchored* refers to the fact that each cluster of registrations was tied to an identifiable team occasion that dedicated time for sustained architectural reasoning, rather than to ad hoc moments within ordinary sprint execution. Both technical leaders treated the backlog as an externalization instrument for tacit concerns and used those occasions as structured prompts for architectural reasoning. The contribution was the empirical observation that even with tool support in place, the usage depended on the persistent efforts of the senior engineers.

RQ2 — A first empirically grounded vocabulary for the content of architectural hypotheses: The five Structural Hypothesis Patterns — *Hypothesis Cluster*, *Regulatory Trigger Hypothesis*, *Service Decomposition Hypothesis*, *Monitored Quality-Attribute Hypothesis*, and *Shared-Resource Hypothesis* — named recurring shapes that architectural hypotheses took when ArchHypo was applied in a regulated, platform-heavy professional context. They complemented the mechanics-focused patterns of the existing ArchHypo pattern language (SILVA, ADOLFO, *et al.*, 2024): the existing patterns addressed *how* hypotheses were assessed and deferred; the Structural Hypothesis Patterns addressed *what kind* of hypothesis was being assessed. The contribution was nominative and was explicitly bounded by the dataset size, with no frequency claims being made.

RQ3 — A benefit-cost characterization of tool-mediated hypothesis engineering, condensed into the Adoption-Condition Candidates: Participants recognized hypothesis management as filling a gap that ADRs leave open, capturing the pre-decision state of uncertainty rather than the post-decision rationale; yet this perceived value was consistently held hostage to sprint delivery pressure and to the cognitive demand of the study’s negative-formulation convention. The four Adoption-Condition Candidates (A1–A4) condensed this tension into named organizational conditions on which sustainable hypothesis engineering depended: senior curation, protected cognitive time, statement-level scaffolding, and ceremonial touchpoints. The contribution was the explicit observation that these conditions are systematically outside what a tool alone could supply, making them the proper starting point for follow-up intervention research.

RQ4 — A worked instance of formative evaluation for hypothesis-engineering tools: Nine refinement categories were derived during the observation period, each grounded in observed friction rather than in design foresight, and several were redeployed and re-observed within the same study. This operationalized the formative-evaluation model advocated in flexible-design guidance (ROBSON and McCARTAN, 2016) for artifacts whose value proposition is tightly coupled to workflow fit. The most structurally consequential refinement — the evolution chart that tracked uncertainty and impact ratings across successive assessments — directly addressed the assessment-history gap of the initial Hypo-Stage version by CORRÊA (2025) and gave teams a mechanism for judging whether their technical-plan actions had the intended effect.

Overarching synthesis. The four contributions converged on a single bounded claim: tool-mediated ArchHypo succeeded at the integration and traceability layer (RQ1, RQ4) and produced a usable vocabulary for the content of architectural hypotheses (RQ2), while the limitations on its use are dictated by the organization itself (RQ3). The Adoption-Condition Candidates therefore functioned as this work’s principal hand-off to the future-work agenda (Section 6.3), where they are read as open empirical problems rather than as actionable guidelines.

6.2 Implications for Practice and Research

The findings of this study carry implications at two levels: for practitioners who are considering adopting hypothesis-engineering tools in their organizations, and for researchers who study tool-mediated architectural practices.

For **practitioners**, the most actionable implication concerns the integration surface of the tool. The study demonstrates that embedding Hypo-Stage within Backstage removed a class of adoption friction that a standalone tool would have created. This is consistent with the broader platform engineering literature on Internal Developer Portals (CORRÊA, 2025) and with the Team Topologies argument that tooling should follow, rather than reshape, team cognitive pathways (SKELTON and PAIS, 2019). A general principle suggested by these findings is that hypothesis-engineering tools are more likely to achieve sustainable adoption when they are delivered as extensions to an existing workflow rather than as a new workflow to be added to an already busy sprint cadence.

A second implication for practitioners concerns the organizational conditions for broad participation. Tool availability is a necessary but not sufficient condition for distributing architectural decision-making beyond senior engineers. The evidence from this study — that twelve of thirteen hypotheses were authored by senior engineers even with tool support in place — suggests that organizations intending to use ArchHypo to operationalize the socio-technical architecture paradigm described by BETTS *et al.* (2024) should combine tool deployment with three complementary investments: protected time for architectural sessions separate from sprint ceremonies (adoption-condition candidate A3); applied guided templates to the hypothesis registration to minimize the difficulty of creating falsifiable statements (adoption-condition candidate A4, actualized in the Hypo-Stage form improvements); and at least one recurring team-visible ceremony where the hypothesis

backlog is reviewed collectively. Without these conditions, the tool risks reinforcing the senior dependency it was designed to dissolve.

For **researchers**, the study offers three methodological observations that may benefit future work in tool-mediated empirical software engineering. First, the formative evaluation design provided richer insight into the causal relationship between tool design choices and adoption behavior than a retrospective evaluation would. This design is especially appropriate for artifacts like Hypo-Stage, whose value proposition is tightly coupled to usability and workflow integration. Second, participant observation combined with artifact analysis of the tool’s own data (the hypothesis records) proved a productive combination: field notes captured the social nuances of the team’s usage, while the tool’s hypothesis records provided a verifiable benchmark to ground observational interpretations. Third, we demonstrate that even brief observation windows in naturalistic settings can identify meaningful structural patterns. This supports the view that flexible case-study designs can produce theoretically generative findings even when the dataset is limited in scale (ROBSON and McCARTAN, 2016).

More broadly, this study contributed to ongoing research on what it means to support distributed architectural decision-making with tools. The vision of a socio-technical architecture (BETTS *et al.*, 2024; HUNTER and STOJMILOVA, 2025; TIBBITTS, 2025) in which architecture work is shared across teams rather than centralized in a dedicated function depends on practitioners at all levels having both the vocabulary and the tooling to surface, assess, and communicate architectural uncertainties. Hypo-Stage, as evaluated here, represents one operational instantiation of that vision. The findings suggest that the tool’s contribution is genuine but partial: it succeeds in lowering context-switching cost and in making hypothesis management available within the daily workflow, but the deeper challenge remains an organizational and pedagogical problem that a tool, by itself, can not solve.

6.3 Limitations and Future Work

Any interpretation of the findings presented in this work must be read against its methodological boundaries. Acknowledging these boundaries is not a disclaimer but a research responsibility: it makes explicit the conditions under which the findings hold and identifies the empirical work required to extend them.

6.3.1 Scope and Transferability

The study was conducted within a single technology organization that already operated a mature Backstage deployment and maintained a scaled service product in Latin America. The enabling condition of Backstage availability is non-trivial: it means that findings about adoption friction may not transfer directly to organizations without an Internal Developer Portal. Similarly, the dominance of Reliability and Auditability as quality attributes in the hypothesis dataset reflects the specific regulatory environment of the study site; organizations in other domains may exhibit very different hypothesis profiles. The organizational culture, the team size differences between the Product and Platform teams, and the presence of a researcher who was also the tool’s supervising author all constitute

contextual specifics whose influence cannot be fully separated from the findings. These characteristics do not invalidate the results; they define the population to which the results most directly apply and orient the replication agenda described below.

6.3.2 Dataset Scope and Hypothesis Lifecycle

Thirteen hypotheses produced by two teams over approximately nine calendar weeks is a dataset appropriate for initial conceptual characterization but insufficient for statistical or frequency claims. The Adoption-Condition Candidates and Structural Hypothesis Patterns should be read as emerging propositions rather than as consolidated findings. Their value at this stage is nominative: they provide a vocabulary for recognition that can be tested and refined in future studies. That observation window was also too short to observe full hypothesis lifecycle completion: eleven of the thirteen hypotheses remained in Open status at the study's end, and the tool's ability to support teams through validation, invalidation, and retrospective learning after hypothesis closure remains empirically unexamined.

6.3.3 Researcher Role and Potential Bias

The researcher's dual position as organizational employee and academic investigator introduces a known potential for bias. Although strategies were employed to maintain role separation, the researcher's status as the tool's supervising author and as a domain expert in ArchHypo may have influenced how participants described their experience. Social desirability effects cannot be ruled out, and the absence of a comparison condition means that it is not possible to distinguish the effect of the tool from the effect of the structured technique itself, or from the effect of the researcher's presence.

A related methodological limitation follows from the active-listening data-collection choice (Section 4.6.3): the sessions were not audio-recorded or transcribed verbatim, so the RQ3 attributions rest on the researcher's dated field notes and post-session written recollection. That absence of a verbatim record is a threat to confirmability—an independent auditor cannot re-inspect the original utterances—and future work should address it with recorded and transcribed sessions (or equivalently auditable contemporaneous notes reviewed by participants).

6.3.4 Deviation from Prescribed Hypothesis Formulation

A construct-validity limitation concerns how hypotheses were formulated during the study. The teams and the researcher adopted a local convention of registering hypotheses in a negative formulation—phrased to describe what would prove each hypothesis false—rather than following ArchHypo's prescribed practice of writing simple, direct, falsifiable statements of the team's current belief (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025). This convention is a plausible cause of the Negative-Formulation Barrier (pattern A4): the cognitive cost that participants reported may stem from the study's operationalization rather than from the technique as designed. Future replications should adopt the prescribed free-form falsifiable statements—optionally aided by interrogative templates such as QUEST (MELEGATI and WANG, 2019)—to test whether the barrier is intrinsic to ArchHypo or an artifact of the convention used here.

6.3.5 Future Work Agenda

The most immediate priority is **replication in additional organizational contexts** with teams of varying size, domain, and tenure profile, and ideally with a researcher who does not hold a pre-existing relationship with the tool. Replication studies could allow the Adoption-Condition Candidates and Structural Hypothesis Patterns to be tested against new hypothesis datasets, providing the basis for either consolidating them into the ArchHypo pattern language (SILVA, ADOLFO, *et al.*, 2024) or revising them in light of disconfirming evidence.

A **retrospective classification of prior ArchHypo datasets** would test the reach and coverage of the Structural Hypothesis Patterns beyond the thirteen-hypothesis corpus of this study. The four hypotheses from the Catch Solve start-up evaluation (SILVA, MELEGATI, WANG, *et al.*, 2024) and the ten from the PropSEP mission-critical project (SILVA, MELEGATI, SILVEIRA, *et al.*, 2025) constitute a natural extension corpus: applying the Structural Hypothesis Patterns vocabulary to these published hypotheses would reveal whether the archetypes generalize across organizational contexts and whether the vocabulary captures structural shapes not represented in the present dataset.

A **longitudinal design** spanning multiple development cycles within the same organization would address the lifecycle limitation directly. Such a study could allow observation of hypothesis closure events and of whether the tool adequately supports retrospective learning once the hypothesis is closed. It would also permit examination of the longer-term trajectory of senior-anchor dependency: whether, over time and with repeated practice, hypothesis registration becomes more broadly distributed across experience levels, or whether the senior-anchored pattern is structurally stable.

The **integration of Hypo-Stage with Architecture Decision Record tools** represents a natural architectural extension of the research. Hypo-Stage manages the pre-decision lifecycle of architectural uncertainty; ADR tools manage the post-decision record. Connecting the two would create a single traceable thread from hypothesis registration through technical plan execution to final architectural decision, addressing a gap in the current tool landscape identified in Chapter 3. Empirical investigation of whether this integration changes how teams reason about the relationship between uncertainty management and decision documentation would constitute a meaningful contribution to the architectural knowledge management literature.

Finally, the **question of broadening participation** beyond senior engineers remains the most consequential open problem in the ArchHypo adoption research agenda. The findings of this study, and of the two prior ArchHypo evaluations, consistently point to a pattern in which the technique's conceptual value is recognized but its practice remains concentrated in the hands of those who already possess architectural fluency. Intervention studies that test whether guided hypothesis templates, changes to the team's review processes, or explicit architectural hygiene ceremonies can measurably increase registration breadth. This would provide the empirical basis for translating the Adoption-Condition Candidates from emerging propositions into actionable adoption guidelines.

Beyond the five strategic directions above, the nine emerging findings summarized in Table 5.6 open a concrete, empirically anchored research agenda. Chapter 5 deliberately

leaves each item in descriptive form, without prescribing a solution; the studies outlined below sketch the empirical work that would be required before any such prescription could be justified.

For the Adoption-Condition Candidates (team-behavior conditions for sustainable hypothesis engineering):

- **A1 — Senior-Curated Hypothesis Backlog** (*mid- and long-horizon uncertainties stay visible mainly because senior engineers author them*): investigate whether review models can shift authorship distribution across all sorts of experience levels without degrading hypothesis quality.
- **A2 — Long-Horizon Registration Priority** (*registration effort is best concentrated on slow, high-consequence uncertainties; fast-resolved micro-decisions rarely repay the cost and are not captured as ADRs either*): design and evaluate lightweight triage cues that help teams distinguish long-horizon uncertainties worth registering from transient micro-decisions that are not.
- **A3 — Registration as a Slack-Dependent Practice** (*well-formulated registration depends on protected time, not on sprint slack*): compare sustained registration rate and hypothesis quality under (i) a periodic architectural-review cadence with protected cognitive time versus (ii) in-sprint ad hoc registration.
- **A4 — Negative-Formulation Barrier** (*the study's negative-formulation convention, not ArchHypo itself, made hypothesis writing persistently cognitively costly*): design and test whether ArchHypo's prescribed free-form falsifiable statements—optionally aided by templates such as QUEST—lower the barrier relative to the negative-formulation convention used here.

For the Structural Hypothesis Patterns (structural archetypes of architectural hypotheses):

- **B1 — Hypothesis Cluster** (*a multi-dimensional platform concern decomposing into a set of linked falsifiable hypotheses*): investigate decomposition strategies and cluster-level review practices for such concerns, and evaluate how cluster-aware tooling affects traceability and decision velocity.
- **B2 — Regulatory Trigger Hypothesis** (*an external compliance obligation surfacing as tracked architectural uncertainty*): evaluate whether modeling compliance obligations as tracked hypotheses, with explicit falsification criteria against the current architecture, improves regulated-domain outcomes compared with implicit compliance assumptions.
- **B3 — Service Decomposition Hypothesis** (*uncertainty about whether and how to decompose a service into microservices along domain boundaries*): investigate whether requiring a discovery phase before a team commits to a decomposition produces measurably better outcomes than committing directly to the refactor.
- **B4 — Monitored Quality-Attribute Hypothesis** (*a quality attribute empirically validated under load but continuously monitored while implementation is still open*): define and empirically test lifecycle states that distinguish “empirically validated

but not yet implemented” from “Open with unresolved uncertainty,” and evaluate reviewer comprehension across the two.

- **B5 — Shared-Resource Hypothesis** (*cross-domain consistency, security, and usability concerns — including ownership governance — for a shared platform resource*): investigate how such concerns (consistency, security, usability, and governance facets such as ownership, cross-domain read and write, and transfer) can be expressed and tracked as architectural hypotheses, and measure the downstream effect on consistency violations.

Appendix A

Research Instruments

This appendix reproduces the instruments and communications used during data collection. Section A.1 contains the open-ended active listening guide used in the end-of-period sessions. Section A.2 reproduces the informed consent framework communicated to participants. Section A.3 reproduces the training communication sent to both teams at the start of the observation period.

A.1 Open-Ended Active Listening Guide

Note on instrument design. The study did not employ a fixed questionnaire. As described in Section 4.6.3, data collection followed an open-ended active listening format consistent with the semi-structured interview approach recommended for flexible designs (ROBSON and McCARTAN, 2016). The guide below provided a thematic scaffold for conversations; the actual order and phrasing of questions varied with the natural flow of each session. Sessions were conducted individually with the senior technical leader of each team at the end of the observation period (one session per team, two in total; approximately 20–30 minutes each). They were not audio-recorded or transcribed; see Section 4.6.3.

Thematic Areas and Guiding Questions

Area 1 — Overall experience of Hypo-Stage.

- Can you walk me through the moments when you actually opened Hypo-Stage during the observation period?
- What made you decide to register a hypothesis at those particular moments?
- Were there situations where you thought about registering a hypothesis but ended up not doing so? What stopped you?

Area 2 — Perceived usefulness of the technique and the tool.

- Has using Hypo-Stage changed the way you think about or discuss architectural decisions with your team?

- Are there specific decisions you handled differently because you had a hypothesis registered?
- What aspect of the tool did you find most useful in practice?
- What aspect of the tool was most frustrating or felt like friction?

Area 3 — Learning curve and technique understanding.

- How confident do you feel formulating a hypothesis statement that is falsifiable? Did that feel natural or forced?
- How did your team members react to the concept of an architectural hypothesis? Do you think the concept would be accessible to a mid-level engineer without facilitation?
- What would have made the onboarding experience easier?

Area 4 — Workflow integration.

- How well did Hypo-Stage fit into your existing sprint ceremonies and day-to-day tooling?
- Did the Backstage integration help or was it irrelevant to adoption?
- If you were to continue using the tool after this study, what would need to change — either in the tool or in your team's process — for it to be sustainable?

Area 5 — Future outlook.

- If you had to name one architectural uncertainty that you are currently managing tacitly that you *wish* you had registered in Hypo-Stage, what would it be?
- Would you recommend Hypo-Stage to another team in the organization? Under what conditions?

A.2 Informed Consent Framework

No formal written consent form was required for this study, as the research involved no sensitive personal data and participation was entirely voluntary (Section 4.9). The following points were communicated verbally and in writing to all participants prior to the observation period and were confirmed by engineering leadership in the organization:

1. **Voluntary participation.** All engineers were free to use Hypo-Stage as much or as little as they chose. Participation in the active listening sessions was optional and had no bearing on the researcher's organizational role.
2. **Anonymization.** The name of the organization, the names of all participants, and the names of internal software products and components would be anonymized in all research outputs. Backstage entity references in hypothesis data would be replaced with generic labels (see Appendix B, Table B.1).

3. **Data use.** Observations, artifact records, and session notes would be used exclusively for academic research purposes and would not be shared with organizational management in any individually identifiable form.
4. **Right to withdraw.** Any participant could withdraw previously provided data or discontinue participation at any point without consequence.
5. **Research knowledge and consent of leadership.** The study was conducted with the knowledge and informal consent of the relevant engineering leadership in the organization, who confirmed that participation was compatible with the teams' work obligations.

A.3 Training Communication

The following text reproduces, in anonymized and lightly edited form, the written communication sent to both teams prior to the training session that inaugurated the observation period. The purpose of this communication was to set expectations, introduce the research context, and invite voluntary participation.

Subject: Architectural Hypothesis Management with Hypo-Stage — Invitation to Participate

Hi Team,

I am reaching out to invite you to participate in a research initiative I am conducting as part of my Master's degree at IME-USP.

What is it about? We will be exploring a technique called Architectural Hypothesis Engineering (ArchHypo) and a supporting tool called Hypo-Stage, which I have integrated into our Backstage portal. The idea is simple: instead of leaving architectural uncertainties as tacit knowledge held by a few engineers, we write them down as structured hypotheses, assess their urgency, and track the actions we are taking to resolve them.

What does participation involve? I will run a short training session (approximately 60–90 minutes) to introduce the technique and the tool. After that, you are invited to use Hypo-Stage whenever you encounter architectural uncertainties in your normal work over the next four weeks. At the end of the period, I would like to have a brief conversation (20–30 minutes) with interested team members about their experience.

Is it mandatory? No. Participation is entirely voluntary. Using the tool as much or as little as you find valuable is up to you.

What happens to the data? All records will be anonymized before being reported in my thesis. The names of our teams, our products, and our systems will not appear in any research output. Observations are used only for academic research.

Please reply if you have any questions, or let me know if you would like to discuss the initiative before the training session.

Looking forward to working with you on this!

[Researcher — role and name anonymized]

Editorial note on observation length. The invitation above references a four-week horizon. In practice, teams continued to use Hypo-Stage voluntarily beyond that window. Hypothesis registration timestamps in Appendix B span 11 February 2026 to 12 April 2026 (61 calendar days, approximately nine weeks), corresponding to five consecutive two-week sprint periods from the earliest to the latest registration in the organization's cadence.

Appendix B

Registered Hypotheses Data

This appendix presents the complete set of thirteen architectural hypotheses registered by the two participating teams in Hypo-Stage during the observation period. Registration timestamps span 11 February 2026 through 12 April 2026 (61 calendar days, approximately nine weeks), covering five consecutive two-week sprint periods from the earliest to the latest creation date in the organization’s cadence. All entity references, service names, and internal artifact links have been anonymized in accordance with the ethical agreement described in Section 4.9. The anonymization legend is provided in Table B.1. Technical planning entries are reproduced in full in Section B.3.

B.1 Anonymization Legend

Hypotheses in Hypo-Stage are linked to components in the organization’s Backstage catalog through `entityRef` strings (human-readable labels that identify a service, worker, platform capability, or template entry). To protect confidentiality under the ethical agreement summarized in Section 4.9, this appendix does not reproduce real catalog identifiers. Instead, each hypothesis table in Section B.2 lists *anonymized* entity references in its **Entities** row; Table B.1 explains what each anonymized string denotes at a functional level, without recovering the original catalog paths or service names.

The table is intentionally typeset here as *non-floating* material so it remains with this section even though it occupies most of a page: the goal is to keep the legend immediately available when reading the hypothesis records, rather than deferring it past the large tables in Section B.2.

Anonymized entity reference	Description
catalog: Kotlin + JOOQ + Spring service template	Internal catalog entry for the JVM service scaffold using Kotlin, JOOQ, and Spring — intended to replace a legacy functional-language stack in the regulated financial domain.
internal-framework: JVM platform stack with Kafka integration	The organization’s shared application framework for JVM services, covering Kafka exposure conventions, request logging, and platform integrations.
service: secure blob storage – API	HTTP API for storing and governing binary documents with security and lifecycle rules.
service: secure blob storage – worker	Background worker for the secure blob storage domain (async processing and maintenance).
service: confidential records & credit-bureau integration – API	API for identity, credit-bureau, and confidential document flows, currently coupled to a single business vertical.
service: confidential records & credit-bureau integration – worker	Worker for the same confidential records and bureau integration context.
worker: high-volume billing aggregation	Long-running worker that aggregates large billing datasets.
platform: shared streaming infrastructure	Shared Kafka streaming platform and consumption standards, modeled as a catalog component.

Table B.1: *Anonymized Backstage entity references and their descriptions.*

B.2 Hypothesis Records

Each hypothesis is typeset as a single record table: fixed metadata rows, a full-width statement row, then author notes and technical-plan references (see Section B.3). The bold header on each record repeats the hypothesis identifier and short label exactly as in Table 5.2 of Chapter 5.

H01 — Kotlin/JOOQ/Spring template readiness

Team	Product
Registered	2026-03-18
Lifecycle status	Open
Source type	Requirements
Uncertainty / Impact	Very High / Very High
Quality attributes	Performance Reliability Maintainability
Entities	catalog: Kotlin + JOOQ + Spring service template

Statement

“The Kotlin + JOOQ + Spring service template is ready to scale for use in regulated financial product services, replacing a legacy functional-language stack.”

Notes	The organization intends to deprecate the legacy stack in that domain; the Kotlin + JOOQ + Spring template is the leading candidate to replace it.
Technical plans	None registered at end of observation period.

H02 — Kafka direct exposure via platform framework

Team	Platform
Registered	2026-02-11
Lifecycle status	Open
Source type	Requirements
Uncertainty / Impact	Very Low / Medium
Quality attributes	Configurability Flexibility Extensibility Modifiability
Entities	internal-framework: JVM platform stack with Kafka integration

Statement

“Kafka should be exposed to application layers for direct use via the internal platform framework.”

Notes	Subject on hold; technical plans expected for the following quarter.
Technical plans	None registered at end of observation period.

H03 — Decouple eng. logs from security/compliance audit

Team	Platform
Registered	2026-02-19
Lifecycle status	Open
Source type	Requirements
Uncertainty / Impact	Medium / Very High
Quality attributes	Observability Auditability
Entities	internal-framework: JVM platform stack with Kafka integration

Statement

“Decouple engineering/observability logs from security/compliance auditing by evolving database logging toward a unified Request Logs capability on the internal platform framework.”

Notes	The platform framework must let applications derive full log context including database logs. It is expected to supersede a legacy audit library and satisfy organization-wide security audit requirements.
Technical plans	TP-03-A, TP-03-B, TP-03-C (see Section B.3).

H04 — Blob storage GDPR right-to-erasure

Team	Product
Registered	2026-03-18
Lifecycle status	Open
Source type	Requirements
Uncertainty / Impact	Medium / Very Low
Quality attributes	Auditability Usability
Entities	service: secure blob storage – API service: secure blob storage – worker

Statement

“The secure blob storage service must retain a reference to the data subject in its domain data so all sensitive blobs can be deleted when the user exercises their right to erasure under applicable privacy regulation.”

Notes	Concern raised in an internal RFC about fully erasing user sensitive data, including binary objects.
Technical plans	None registered at end of observation period.

H05 — Credit-bureau context-agnostic restructuring

Team	Product
Registered	2026-03-18
Lifecycle status	In Review
Source type	Solution
Uncertainty / Impact	High / Very High
Quality attributes	Interoperability Extensibility Reusability
Entities	service: confidential records & credit-bureau integration – API ... worker

Statement

“Credit-bureau integrations can be decoupled from vertical-specific domain entities into context-agnostic middleware that accepts only each provider’s minimum required data (e.g., national ID), reusable across domains without duplicated cost.”

Notes	Discovery focused on: (1) dependency matrix — current vs. real dependency per integration; (2) agnostic interface — contract without domain concepts, only minimal identifiers and bureau outputs; (3) migration plan — which flows can isolate immediately vs. requiring deeper business-logic refactors.
Technical plans	TP-05-A (see Section B.3).

H06 — Billing aggregation chunked keyset pagination

Team	Product
Registered	2026-02-24
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Very Low / Very High
Quality attributes	Performance Reliability Scalability
Entities	worker: high-volume billing aggregation

Statement

“Chunked keyset pagination with isolated repository-level transactions keeps memory strictly bounded by chunk size and prevents out-of-memory failures when aggregating millions of billing records.”

Notes	Validated with a capacity/volume test at 5× production scale (45M rows total, 1.5M rows/month). Profiling showed JPA streams can cause first-level cache growth; chunking scopes transactions (≈10k rows) so the garbage collector can reclaim memory (saw-tooth heap pattern). Composite index on (billing_year_month, status, id).
Technical plans	TP-06-A (see Section B.3).

H07 — Blob ownership transfer across service domains

Team	Product
Registered	2026-04-07
Lifecycle status	Validated
Source type	Requirements
Uncertainty / Impact	Very Low / Very High
Quality attributes	Scalability Usability Configurability Flexibility
Entities	service: secure blob storage – API service: secure blob storage – worker

Statement

“Secure blob storage must support transferring blob ownership across service domains so blobs can move from one owning service to another.”

Notes	Validated through a technical discovery spike that defined the transfer capability and produced a prioritized implementation backlog.
Technical plans	TP-07-A (see Section B.3).

H08 — Blob cross-client sharing governance

Team	Product
Registered	2026-03-18
Lifecycle status	Validated
Source type	Requirements
Uncertainty / Impact	Very Low / Very High
Quality attributes	Security Usability Interoperability Extensibility
Entities	service: secure blob storage – API

Statement

“Secure blob storage must support sharing a blob across multiple clients so the same object can be read from different domain services.”

Notes	Example: a customer-facing channel receives documents later routed to support tooling; support needs read access in chat history without breaking vault invariants.
Technical plans	TP-08-A (see Section B.3).

H09 — DLT for audit vs. offset-based replay

Team	Platform
Registered	2026-04-12
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Medium / Very High
Quality attributes	Auditability Reliability
Entities	platform: shared streaming infrastructure

Statement

“Using the Dead-Letter Topic (DLT) only for auditing and diagnosis, and reprocessing from the original topic’s offset, preserves consistency and auditability better than re-enqueueing messages from the DLT.”

Notes	How to falsify: controlled replay from the DLT without mutating the event must be proven safer, simpler, and more auditable than replay by offset.
Technical plans	TP-09-A, TP-09-B (see Section B.3).

H10 — Synchronous retry for transient consumer failures

Team	Platform
Registered	2026-04-12
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Medium / High
Quality attributes	Interoperability Performance
Entities	platform: shared streaming infrastructure

Statement

“For short-lived transient failures, limited synchronous retry on the consumer thread preserves per-partition ordering better and with lower overall complexity than asynchronous retry on a separate queue.”

Notes	How to falsify: a per-partition retry queue preserves order with less blocking time and acceptable operational cost.
Technical plans	TP-10-A (see Section B.3).

H11 — Failure state inside vs. outside Kafka flow

Team	Platform
Registered	2026-04-12
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Medium / High
Quality attributes	Reliability Modifiability Interoperability
Entities	platform: shared streaming infrastructure

Statement

“Keeping failure state inside the Kafka message flow increases coordination complexity across threads and complicates failure handling, compared to handling failures in the execution layer and keeping consumption non-blocking.”

Notes	How to falsify: modelling failure state in Kafka shows lower operational complexity and equal or better recovery than handling failures in the execution layer.
Technical plans	TP-11-A (see Section B.3).

H12 — Error classification before retry

Team	Platform
Registered	2026-04-12
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Medium / High
Quality attributes	Performance Reliability
Entities	platform: shared streaming infrastructure

Statement

“Classifying errors as recoverable vs. non-recoverable before applying retry reduces consumer lag and computational cost without worsening the successful processing rate.”

Notes	How to falsify: indiscriminate retries improve final success rate without meaningful latency/capacity cost.
Technical plans	TP-12-A, TP-12-B (see Section B.3).

H13 — Payload immutability before re-enqueueing

Team	Platform
Registered	2026-04-12
Lifecycle status	Open
Source type	Solution
Uncertainty / Impact	Very Low / Very High
Quality attributes	Auditability
Entities	platform: shared streaming infrastructure

Statement

“Allowing fixes by mutating the payload before re-enqueueing increases consistency violations and makes traceability and auditing harder.”

Notes	How to falsify: a formal correction mechanism fully preserves audit trail, causality, and reproducibility of the event.
--------------	---

Technical plans	TP-13-A (see Section B.3).
------------------------	---

B.3 Technical Plans Catalog

The following entries reproduce all technical plans registered in Hypo-Stage during the observation period. Each entry is identified by a plan code (e.g., TP-03-A), the parent hypothesis, the action type, the target date, a description of the planned action, and the expected outcome. Internal artifact links (tickets, pull requests, design documents) have been redacted.

TP-03-A (H03 — Decouple eng. logs from security/compliance audit) *Action type:* Architectural Spike *Target date:* 2026-03-31

Description: Discover and implement default mapping of database transaction IDs into the Request Log structure within the internal framework so every request and application log entry carries database transaction IDs for correlation and audit trail without per-application integration.

Expected outcome: Request logs consistently include database transaction ID fields; applications using the framework obtain audit-ready context without custom wiring.

TP-03-B (H03) *Action type:* Modularization / Flexibility *Target date:* 2026-03-31

Description: Extend framework logging to worker modules so critical database transactions in workers log with the same context (transaction IDs, correlation) as request paths.

Expected outcome: Workers emit logs with database transaction IDs aligned with the Request Log infrastructure; audit coverage for asynchronous/background work.

TP-03-C (H03) *Action type:* Modularization / Flexibility *Target date:* 2026-03-31

Description: Stream database transaction properties to the same log buckets and pipeline as Request and Application logs for a single audit/observability backbone.

Expected outcome: Database transaction properties follow the same ingestion and retention path as Request logs; no separate audit-only pipeline.

TP-05-A (H05 — Credit-bureau context-agnostic restructuring) *Action type:* Discovery
Target date: ongoing

Description: Structured discovery covering: (1) dependency matrix (current vs. real dependencies per integration); (2) agnostic interface proposal (contract free of domain concepts, accepting only minimal provider identifiers and bureau outputs); (3) migration sequencing (which flows can isolate immediately vs. requiring deeper business-logic refactors).

Expected outcome: A prioritized plan for context-agnostic restructuring and a draft interface contract.

TP-06-A (H06 — Billing aggregation chunked keyset pagination) *Action type:* Implementation (Other) *Target date:* 2026-03-13

Description: Implement the billing aggregation worker with chunked keyset pagination using a `LIMIT` loop with `id > :lastProcessedId`; avoid broad `@Transactional` so the Hibernate first-level cache clears per chunk; add composite index on `(billing_year_month, status, id)`.

Expected outcome: Process 300k+ rows/month safely without out-of-memory failures; stable sawtooth heap profile; healthy database CPU/I/O via index scans.

TP-07-A (H07 — Blob ownership transfer across service domains) *Action type:* Architectural Spike *Target date:* 2026-04-30

Description: Technical discovery to define the transfer capability on the blob storage API.

Expected outcome: A backlog for implementing the ownership transfer feature.

TP-08-A (H08 — Blob cross-client sharing governance) *Action type:* Guideline Implementation *Target date:* 2026-04-06

Description: Alignment session on cross-domain sharing and transfers so that long-lived access (e.g., messaging history) does not require definitive ownership by every consumer.

Expected outcome: A governance decision enabling prioritization of sharing and transfer implementation work.

TP-09-A (H09 — DLT for audit vs. offset-based replay) *Action type:* Guideline Implementation *Target date:* 2026-06-30

Description: Produce a runbook for reprocessing by offset with minimal tooling for controlled replay.

Expected outcome: Documentation of how to reprocess by offset with minimal tooling.

TP-09-B (H09) *Action type:* Experiment *Target date:* 2026-06-30

Description: Operational exercise with a simulated incident.

Expected outcome: Measure MTTR, duplicate rate, and divergence between the original event and the final state.

TP-10-A (H10 — Synchronous retry for transient failures) *Action type:* Tracer Bullet
Target date: 2026-06-30

Description: Tracer bullet on a critical consumer path; inject transient failures with windows of 1s, 5s, and 30s.

Expected outcome: Measure out-of-order delivery, consumer lag, throughput, and retry success rate.

TP-11-A (H11 — Failure state inside vs. outside Kafka flow) *Action type:* Architectural Spike
Target date: 2026-06-30

Description: Spike comparing both failure-state approaches on a real consumer.

Expected outcome: Measure consumer lag, blocked-thread time, states/transitions, and operational recovery steps for each approach.

TP-12-A (H12 — Error classification before retry) *Action type:* Guideline Implementation
Target date: 2026-06-30

Description: Define an error taxonomy and mandatory classification step before retry is applied.

Expected outcome: Practical directives for classifying errors before retry, consumable by all teams using the shared streaming platform.

TP-12-B (H12) *Action type:* Software Analytics (Architectural Trigger) *Target date:* 2026-06-30

Description: Instrument analytics by error class on a real consumer: success rate, average attempts, extra time per message, and discards to DLT.

Expected outcome: A metrics dataset sufficient to falsify or support the hypothesis.

TP-13-A (H13 — Payload immutability before re-enqueueing) *Action type:* Experiment
Target date: 2026-06-30

Description: Enforce an immutable-payload guideline; add automated validation in replay tooling; require mandatory architectural review for exceptions.

Expected outcome: Experiment results that validate or refute the hypothesis and establish a baseline for payload governance in the shared streaming platform.

References

- [ANJUM 2026] Mateen Ali ANJUM. “Platform engineering and internal developer portals: a multivocal literature review”. *Frontiers in Computer Science* 8 (2026), p. 1814498. DOI: [10.3389/fcomp.2026.1814498](https://doi.org/10.3389/fcomp.2026.1814498) (cit. on pp. 16, 21).
- [L. BASS *et al.* 2021] Len BASS, Paul CLEMENTS, and Rick KAZMAN. *Software Architecture in Practice*. 4th ed. Addison-Wesley, 2021 (cit. on p. 9).
- [BETTS *et al.* 2024] Thomas BETTS, Blanca García GIL, Eran STILLER, Daniel BRYANT, and Rafal GANCARZ. *InfoQ Software Architecture and Design Trends Report — 2024*. <https://www.infoq.com/articles/architecture-trends-2024/>. Grey literature — practitioner report. Accessed: 23 March 2026. 2024 (cit. on pp. 1, 2, 4, 6, 25, 69, 70).
- [BOROWA *et al.* 2023] Klara BOROWA, Rafał LEWANCZYK, Klaudia STPICZYNSKA, Patryk STRADOMSKI, and Andrzej ZALEWSKI. “What rationales drive architectural decisions? an empirical inquiry”. In: *Software Architecture: 17th European Conference on Software Architecture, ECSA 2023, Istanbul, Turkey, September 18–22, 2023, Proceedings*. Vol. 14212. Lecture Notes in Computer Science. Springer, 2023, pp. 303–318. DOI: [10.1007/978-3-031-42592-9_21](https://doi.org/10.1007/978-3-031-42592-9_21) (cit. on p. 24).
- [CONWAY 1968] Melvin E. CONWAY. “How do committees invent?” *Datamation* 14.4 (1968), pp. 28–31 (cit. on p. 24).
- [CORRÊA 2025] Pedro Henrique Mariano CORRÊA. *ArchHypo Tooling: Enabling Practical Hypothesis Engineering for Software Architecture*. Capstone Project Report (Bachelor), Institute of Mathematics and Statistics, University of São Paulo. Supervised by José G. Lima Neto; co-supervised by Paulo Meirelles. 2025 (cit. on pp. xvi, xvii, 3–7, 17–21, 37, 53, 69).
- [DA SILVEIRA NETO *et al.* 2025] Manoel Valério DA SILVEIRA NETO, Andreia MALUCELLI, and Sheila REINEHR. “Agile software architecture: perceptions on quality and architectural technical debt management”. In: *Proceedings of the XIX Brazilian Symposium on Software Components, Architectures, and Reuse (SBCARS)*. SBC, 2025. DOI: [10.5753/sbcars.2025.11890](https://doi.org/10.5753/sbcars.2025.11890) (cit. on p. 24).

- [DASANAYAKE *et al.* 2016] Sandun DASANAYAKE, Jouni MARKKULA, Sanja AARAMAA, and Markku OIVO. “An empirical study on collaborative architecture decision making in software teams”. In: *Software Architecture: 10th European Conference, ECSA 2016, Copenhagen, Denmark, November 28–December 2, 2016, Proceedings*. Vol. 9839. Lecture Notes in Computer Science. Springer, 2016. DOI: [10.1007/978-3-319-48992-6_18](https://doi.org/10.1007/978-3-319-48992-6_18) (cit. on p. 24).
- [FAGERHOLM *et al.* 2017] Fabian FAGERHOLM, Alejandro Sanchez GUINEA, Hanna MÄENPÄÄ, and Jürgen MÜNCH. “The RIGHT model for continuous experimentation”. *Journal of Systems and Software* 123 (2017), pp. 292–305. DOI: [10.1016/j.jss.2016.03.034](https://doi.org/10.1016/j.jss.2016.03.034) (cit. on p. 2).
- [FOOTE and YODER 1999] Brian FOOTE and Joseph YODER. “Big ball of mud”. *Pattern Languages of Program Design* 4 (1999), pp. 654–692. URL: <https://www.laputan.org/mud/mud.html> (cit. on p. 10).
- [HERBSLEB and GRINTER 1999] James D. HERBSLEB and Rebecca E. GRINTER. “Splitting the organization and integrating the code: Conway’s law revisited”. In: *Proceedings of the 21st International Conference on Software Engineering (ICSE)*. ACM/IEEE, 1999, pp. 85–95. URL: https://www.cs.cmu.edu/afs/cs.cmu.edu/Web/People/jdh/collaboratory/research_papers/ICSE99.pdf (cit. on pp. 1, 24).
- [HUNTER and STOJMILOVA 2025] Peter HUNTER and Elena STOJMILOVA. “Empowering teams: decentralizing architectural decision-making”. *InfoQ* (Nov. 2025). Grey literature — practitioner article. URL: <https://www.infoq.com/articles/empowering-decentralizing-architectural-decision-making/> (cit. on pp. 25, 70).
- [KAZMAN *et al.* 1998] Rick KAZMAN *et al.* “The architecture tradeoff analysis method”. In: *Proceedings of the Fourth International Conference on Engineering of Complex Computer Systems (ICECCS)*. IEEE, 1998, pp. 68–78. DOI: [10.1109/ICECCS.1998.706657](https://doi.org/10.1109/ICECCS.1998.706657) (cit. on p. 26).
- [MELEGATI and WANG 2019] Jorge MELEGATI and Xiaofeng WANG. “QUEST: new practices to represent hypotheses in experiment-driven software development”. In: *Proceedings of the 2nd ACM SIGSOFT International Workshop on Software-Intensive Business: Start-ups, Platforms, and Ecosystems (IWSiB)*. 2019, pp. 13–18. DOI: [10.1145/3340481.3342732](https://doi.org/10.1145/3340481.3342732) (cit. on pp. 10, 60, 71).
- [MELEGATI, WANG, and ABRAHAMSSON 2019] Jorge MELEGATI, Xiaofeng WANG, and Pekka ABRAHAMSSON. “Hypotheses engineering: first essential steps of experiment-driven software development”. In: *IEEE/ACM Joint 4th International Workshop on Rapid Continuous Software Engineering and 1st International Workshop on Data-Driven Decisions, Experimentation and Evolution (RCoSE/DDrEE)*. 2019, pp. 16–19. DOI: [10.1109/RCoSE/DDrEE.2019.00011](https://doi.org/10.1109/RCoSE/DDrEE.2019.00011) (cit. on pp. 2, 10, 21).

REFERENCES

- [NYGARD 2011] Michael NYGARD. *Documenting Architecture Decisions*. <https://cognitect.com/blog/2011/11/15/documenting-architecture-decisions>. Grey literature — practitioner blog post introducing Architecture Decision Records. Nov. 2011 (cit. on p. 26).
- [OLIVEIRA DE DIANA *et al.* 2011] Mauricio José de OLIVEIRA DE DIANA, Fabio KON, and Marco Aurélio GEROSA. “Conducting an architecture group in a multi-team agile environment”. In: *Proceedings of the 5th Brazilian Symposium on Software Components, Architectures and Reuse (SBCARS)*. SBC, 2011. URL: https://www.ime.usp.br/~kon/papers/arch_group.pdf (cit. on pp. 23, 24, 26).
- [OLSSON and BOSCH 2014] Helena Holmström OLSSON and Jan BOSCH. “The HYPEX model: from opinions to data-driven software development”. In: *Continuous Software Engineering*. Ed. by Jan BOSCH. Cham: Springer, 2014, pp. 155–164. DOI: [10.1007/978-3-319-11283-1_13](https://doi.org/10.1007/978-3-319-11283-1_13) (cit. on p. 2).
- [PERRY and WOLF 1992] Dewayne E. PERRY and Alexander L. WOLF. “Foundations for the study of software architecture”. *ACM SIGSOFT Software Engineering Notes* 17.4 (1992), pp. 40–52. DOI: [10.1145/141874.141884](https://doi.org/10.1145/141874.141884) (cit. on p. 1).
- [PRYCE 2024] Nat PRYCE. *adr-tools: Command-Line Tools for Working with Architecture Decision Records*. <https://github.com/npryce/adr-tools>. Grey literature — open-source project documentation. 2024 (cit. on p. 26).
- [RAZAVIAN *et al.* 2019] Maryam RAZAVIAN, Barbara PAECH, and Antony TANG. “Empirical research for software architecture decision making: an analysis”. *Journal of Systems and Software* 149 (2019), pp. 360–381. DOI: [10.1016/j.jss.2018.12.003](https://doi.org/10.1016/j.jss.2018.12.003) (cit. on p. 24).
- [RICHARDS and FORD 2020] Mark RICHARDS and Neal FORD. *Fundamentals of Software Architecture: An Engineering Approach*. O’Reilly Media, 2020. ISBN: 9781492043454 (cit. on p. 9).
- [RICHARDSON 2018] Chris RICHARDSON. *Microservices Patterns: With Examples in Java*. Manning Publications, 2018. ISBN: 9781617294549 (cit. on p. 62).
- [ROBSON and McCARTAN 2016] Colin ROBSON and Kieran McCARTAN. *Real World Research: A Resource for Users of Social Research Methods in Applied Settings*. 4th. John Wiley & Sons, 2016. ISBN: 978-1-118-74523-6 (cit. on pp. xi, xiii, xvii, 29–35, 37–42, 46, 69, 70, 75).
- [SALAMEH and J. M. BASS 2021] Abdallah SALAMEH and Julian M. BASS. “An architecture governance approach for Agile development by tailoring the Spotify model”. *AI & Society* 37 (2021), pp. 1225–1243. DOI: [10.1007/s00146-021-01240-x](https://doi.org/10.1007/s00146-021-01240-x) (cit. on pp. 23, 26).

- [SILVA, ADOLFO, *et al.* 2024] Kelson SILVA, Luciane Baratto ADOLFO, *et al.* “Patterns for using hypothesis engineering to manage architectural uncertainty”. In: *Proceedings of the 29th European Conference on Pattern Languages of Programs, People, and Practices (EuroPLoP 2024)*. Irsee, Germany: ACM, 2024. DOI: [10.1145/3698322.3698333](https://doi.org/10.1145/3698322.3698333) (cit. on pp. [xvi](#), [6](#), [10](#), [50](#), [55](#), [58](#), [60](#), [65](#), [67](#), [68](#), [72](#)).
- [SILVA, MELEGATI, SILVEIRA, *et al.* 2025] Kelson SILVA, Jorge MELEGATI, Fábio SILVEIRA, *et al.* “ArchHypo: managing software architecture uncertainty using hypotheses engineering”. *IEEE Transactions on Software Engineering* 51.2 (2025), pp. 481–499. DOI: [10.1109/TSE.2024.3520477](https://doi.org/10.1109/TSE.2024.3520477) (cit. on pp. [xvi](#), [xvii](#), [1–4](#), [10–16](#), [21](#), [26](#), [32](#), [36](#), [42](#), [47](#), [50](#), [52–54](#), [60](#), [67](#), [71](#), [72](#)).
- [SILVA, MELEGATI, WANG, *et al.* 2024] Kelson SILVA, Jorge MELEGATI, Xiaofeng WANG, Maurício FERREIRA, and Eduardo GUERRA. “Using hypotheses to manage technical uncertainty and architecture evolution in a software start-up”. *IEEE Software* 41.4 (2024), pp. 7–13. DOI: [10.1109/MS.2024.3383628](https://doi.org/10.1109/MS.2024.3383628). URL: <https://dl.acm.org/doi/abs/10.1109/MS.2024.3383628> (cit. on pp. [xvii](#), [1–4](#), [10](#), [12–14](#), [16](#), [21](#), [26](#), [47](#), [53](#), [60](#), [67](#), [72](#)).
- [SKELTON and PAIS 2019] Matthew SKELTON and Manuel PAIS. *Team Topologies: Organizing Business and Technology Teams for Fast Flow*. IT Revolution Press, 2019 (cit. on pp. [9](#), [10](#), [16](#), [24](#), [25](#), [69](#)).
- [SOUZA *et al.* 2019] Eric SOUZA, Ana MOREIRA, and Miguel GOULÃO. “Deriving architectural models from requirements specifications: a systematic mapping study”. *Information and Software Technology* 109 (2019), pp. 26–39. DOI: [10.1016/j.infsof.2019.01.004](https://doi.org/10.1016/j.infsof.2019.01.004) (cit. on pp. [1–4](#), [6](#), [7](#), [9](#), [20](#), [21](#), [24](#)).
- [SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION 2024] SPOTIFY and CLOUD NATIVE COMPUTING FOUNDATION. *Backstage: An Open Source Framework for Building Developer Portals*. <https://backstage.io>. Open-source project; over 3,000 companies adopted as of 2025. CNCF Incubating Project. 2024 (cit. on pp. [xvii](#), [17](#), [18](#), [21](#)).
- [TANG *et al.* 2017] Antony TANG, Maryam RAZAVIAN, Barbara PAECH, and Tom-Michael HESSE. “Human aspects in software architecture decision making: a literature review”. In: *2017 IEEE International Conference on Software Architecture (ICSA)*. IEEE, 2017, pp. 107–116. DOI: [10.1109/ICSA.2017.15](https://doi.org/10.1109/ICSA.2017.15) (cit. on p. [24](#)).
- [TIBBITTS 2025] Lindsey TIBBITTS. “Decentralized architecture needs more than autonomy”. *InfoQ* (June 2025). Grey literature — practitioner article. URL: <https://www.infoq.com/articles/decentralized-architecture-advice-process/> (cit. on pp. [25](#), [70](#)).
- [VAILLANT 2024] Thomas VAILLANT. *Log4brains: Architecture Decision Records Management and Publication Tool*. <https://github.com/thomvaill/log4brains>. Grey literature — open-source project documentation. 2024 (cit. on p. [26](#)).
- [YIN 2009] Robert K. YIN. *Case Study Research: Design and Methods*. 4th. Thousand Oaks, CA: SAGE Publications, 2009 (cit. on pp. [29–31](#), [34](#), [37](#), [41](#)).